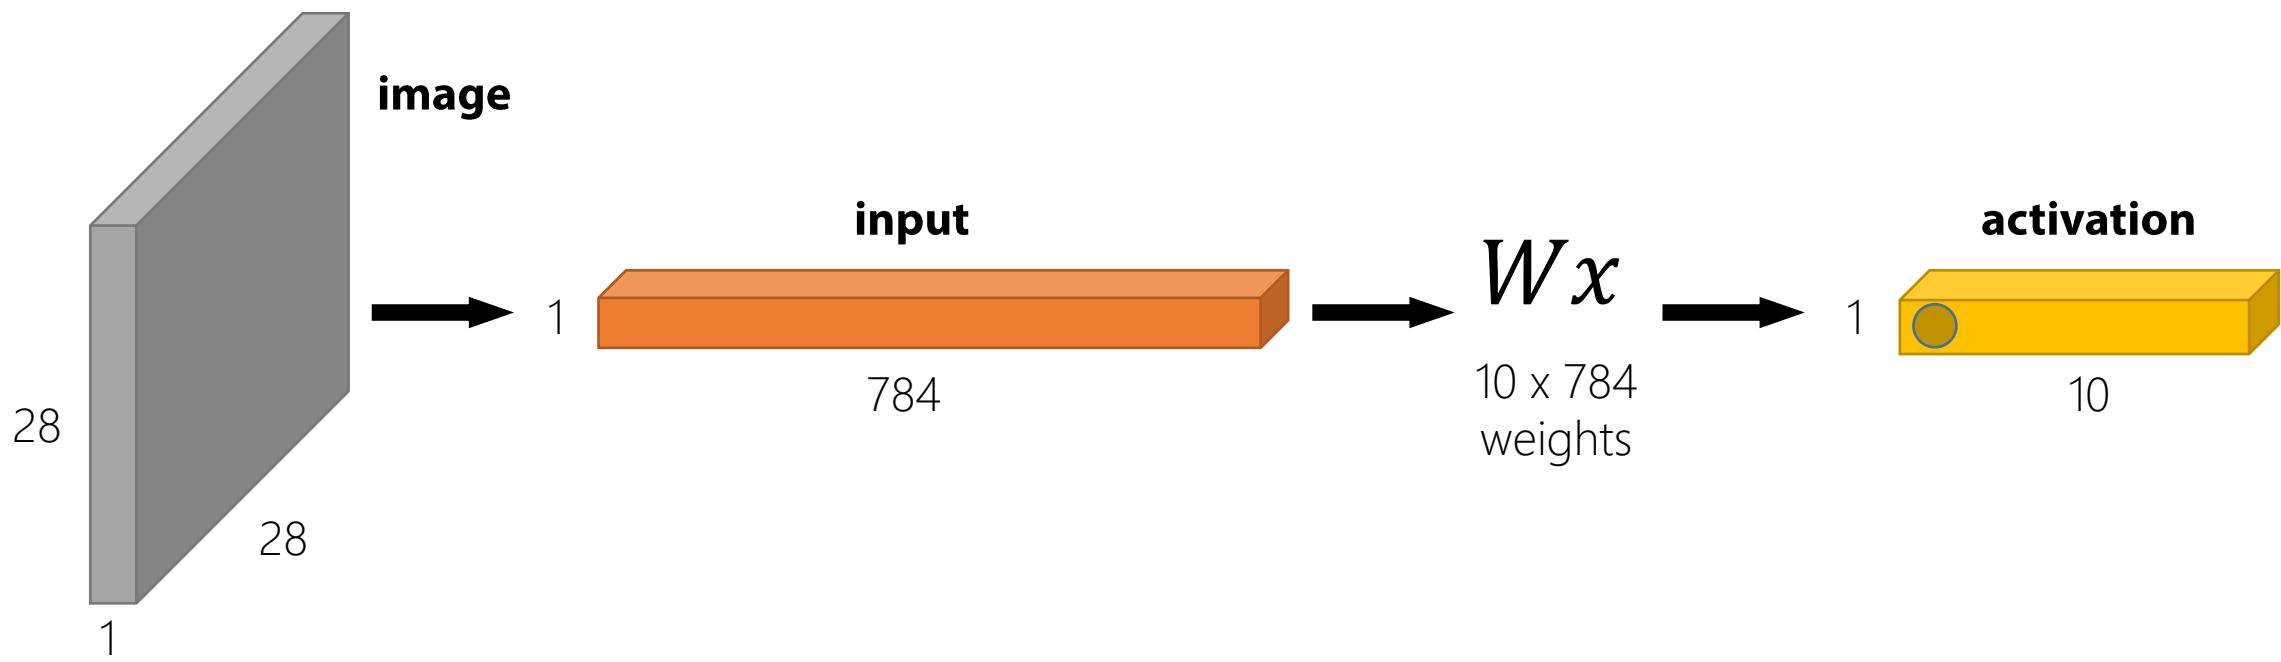


Convolutional Neural Networks

Stephen Baek

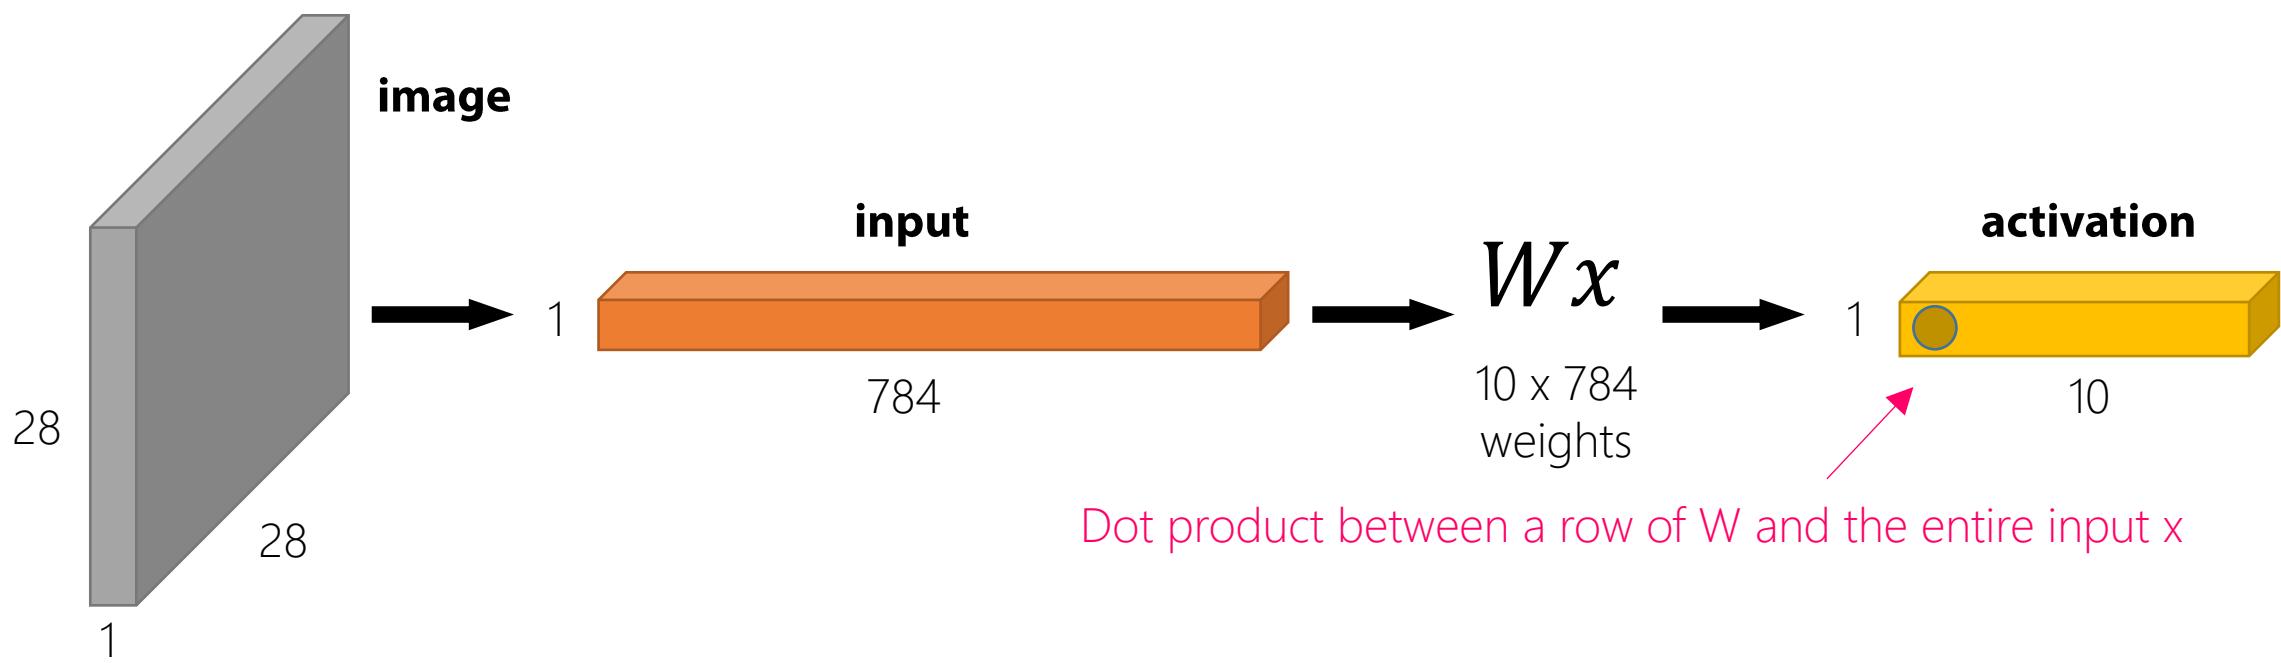
Fully Connected Layer

- 28×28 image → stretch to 784×1
- $64 \times 64 \times 3$ image → stretch to 12288×1
- ...



Fully Connected Layer

- 28×28 image → stretch to 784×1
- $64 \times 64 \times 3$ image → stretch to 12288×1
- ...



Draw your number here



Downsampled drawing:

2

First guess:

2

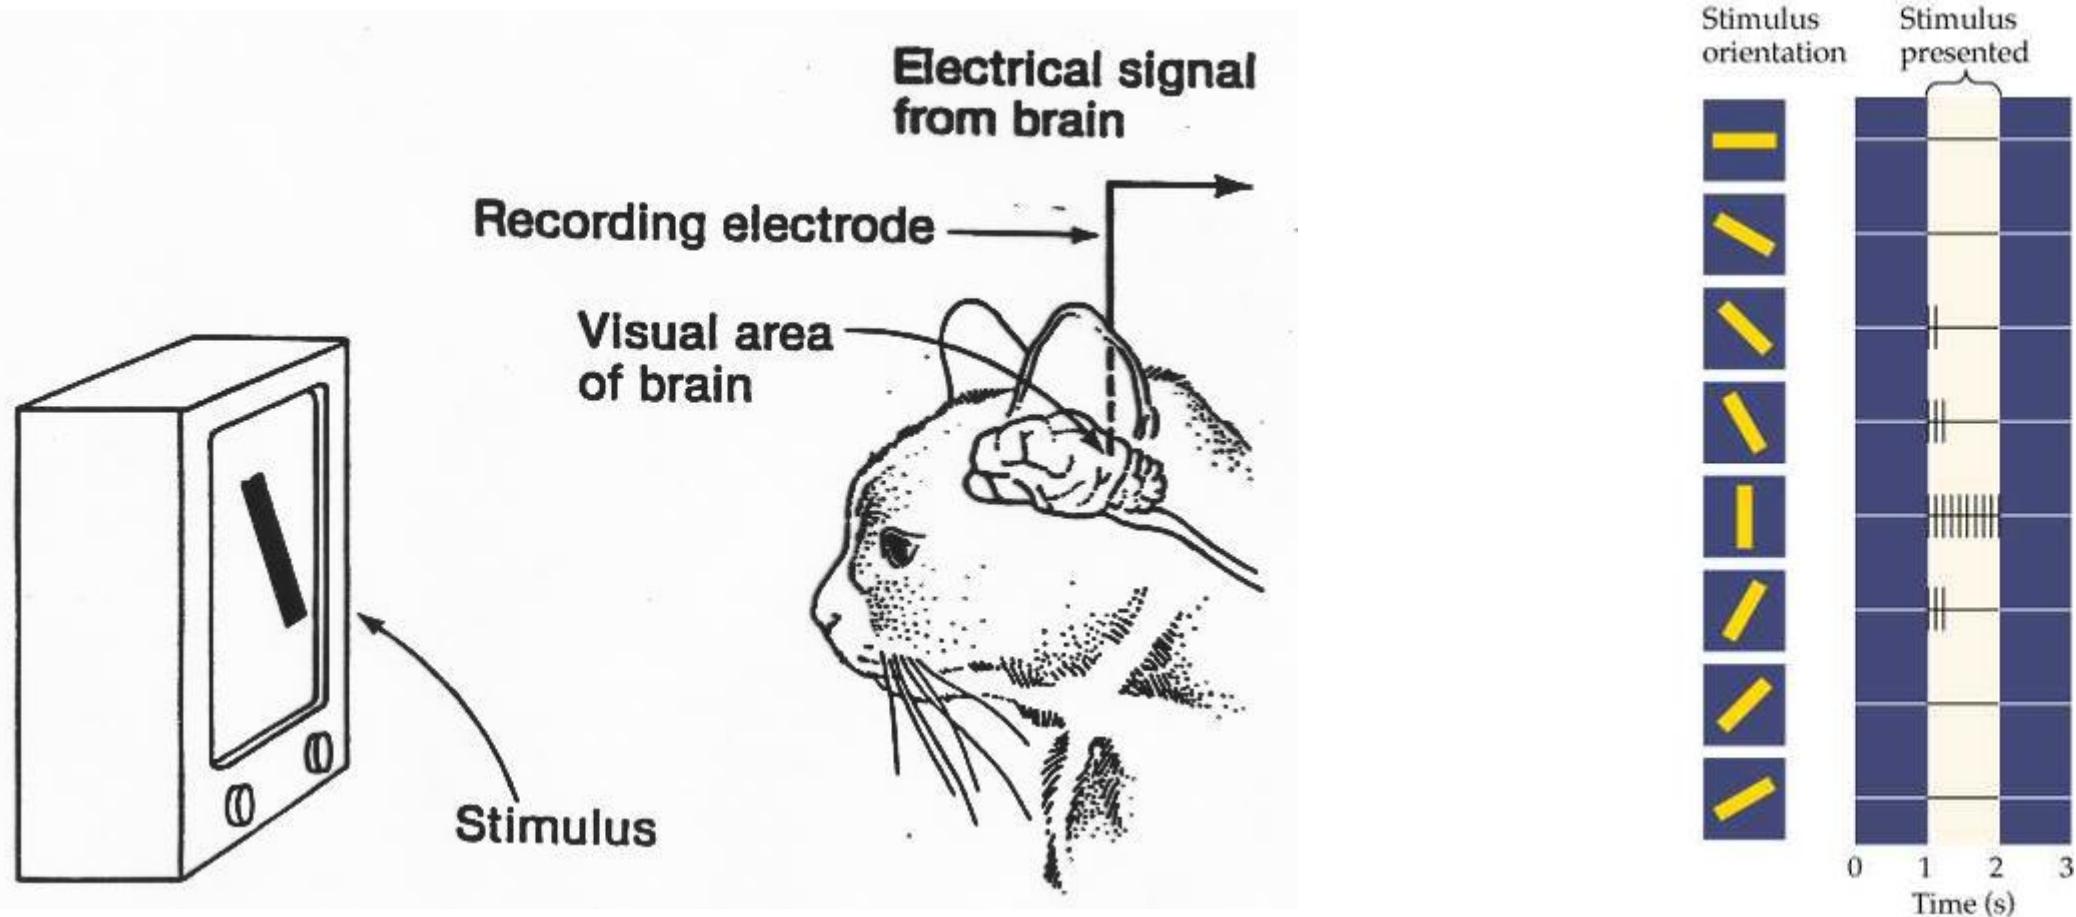
Second guess:

1

0 1 2 3 4 5 6 7 8 9
█ █ █ █ █ █ █ █ █ █



Hubel & Wiesel (1959 ~)



Neurons in the visual cortex respond selectively to oriented edges. Neurons in visual cortex typically respond vigorously to a bar of light oriented at a particular angle and weakly or not at all to other orientations.

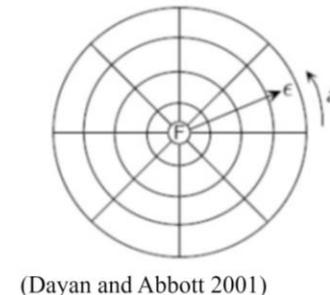
Hubel & Wiesel (1959 ~)

Warning: GRAPHIC CONTENT

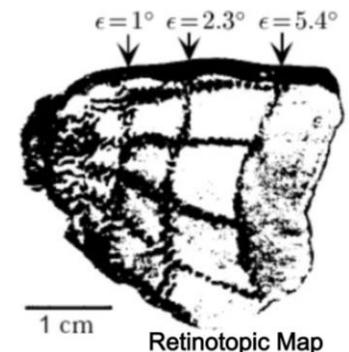


Topographic Maps in Cortex

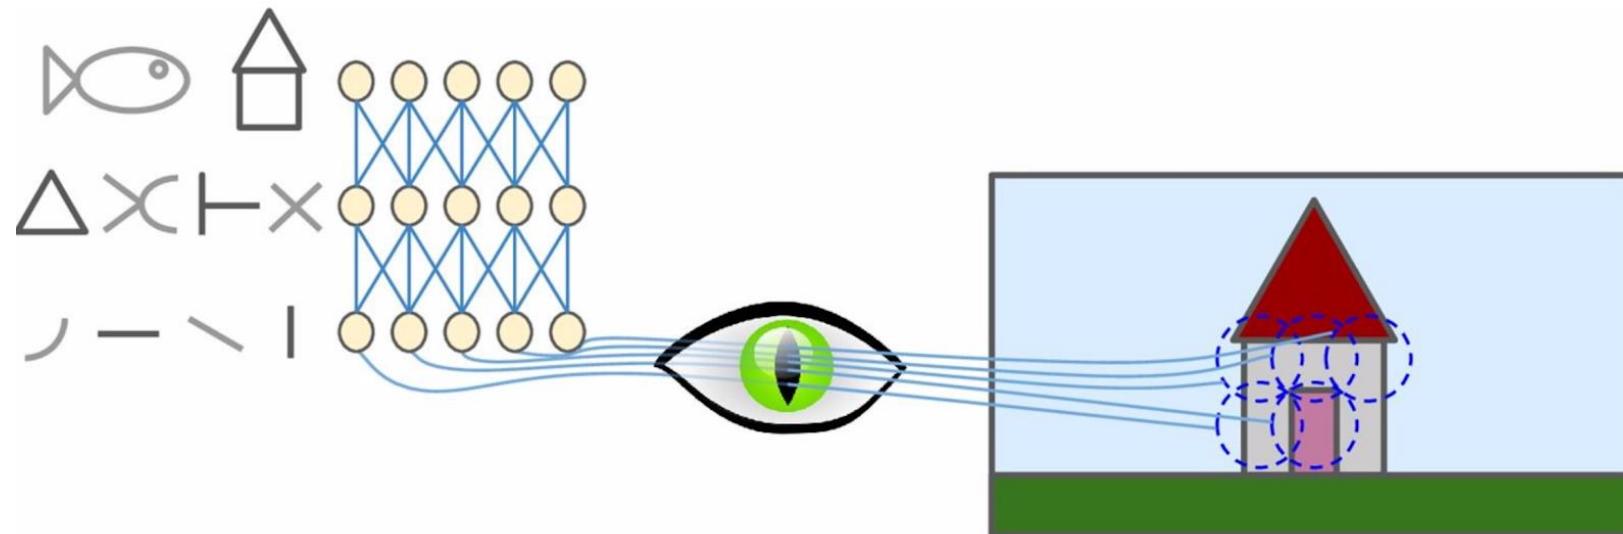
- Each visual sensitive cell only responses to stimuli of a limited region (receptive field)
- Neighboring cells have partially overlapping receptive fields
- Neighboring points in a visual image evoke activity in neighboring regions of visual cortex
- In this manner, the visual system easily maintain the information of the spatial location of stimulus



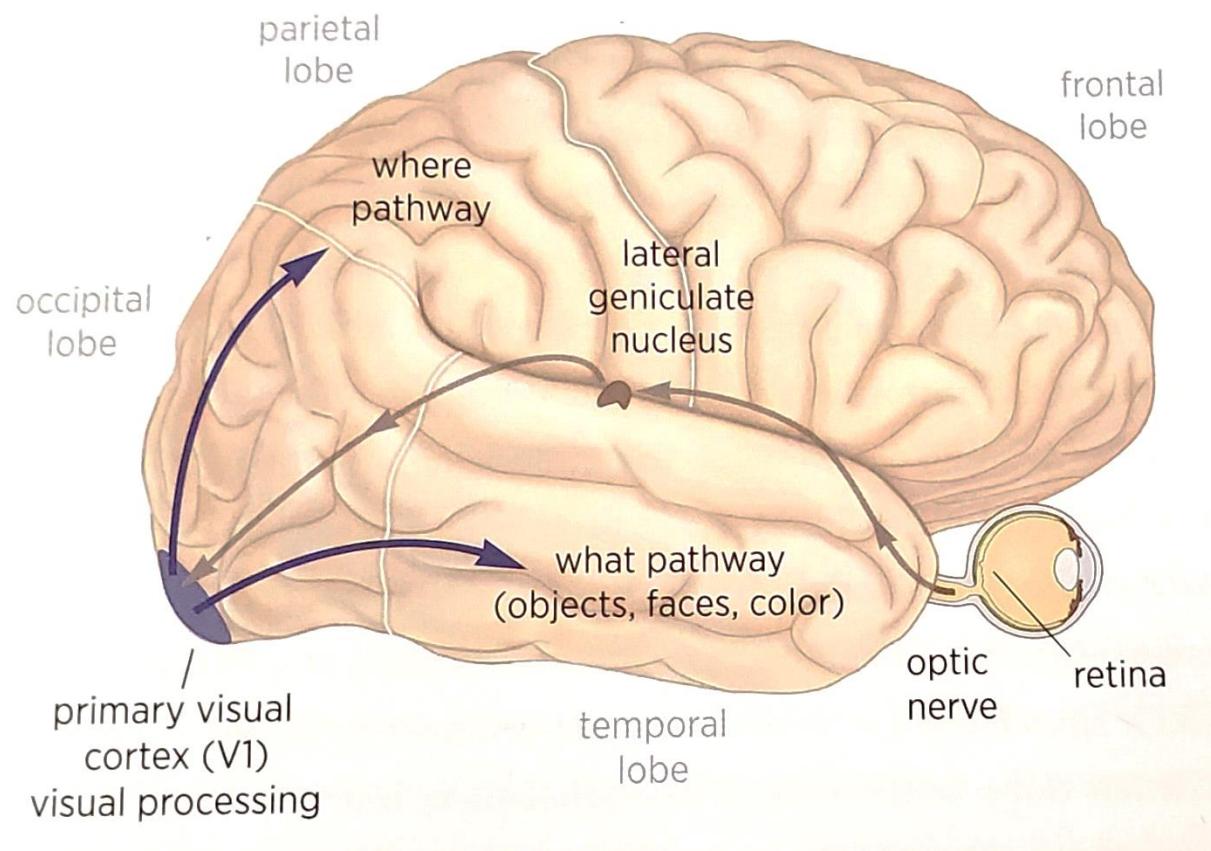
(Dayan and Abbott 2001)



7



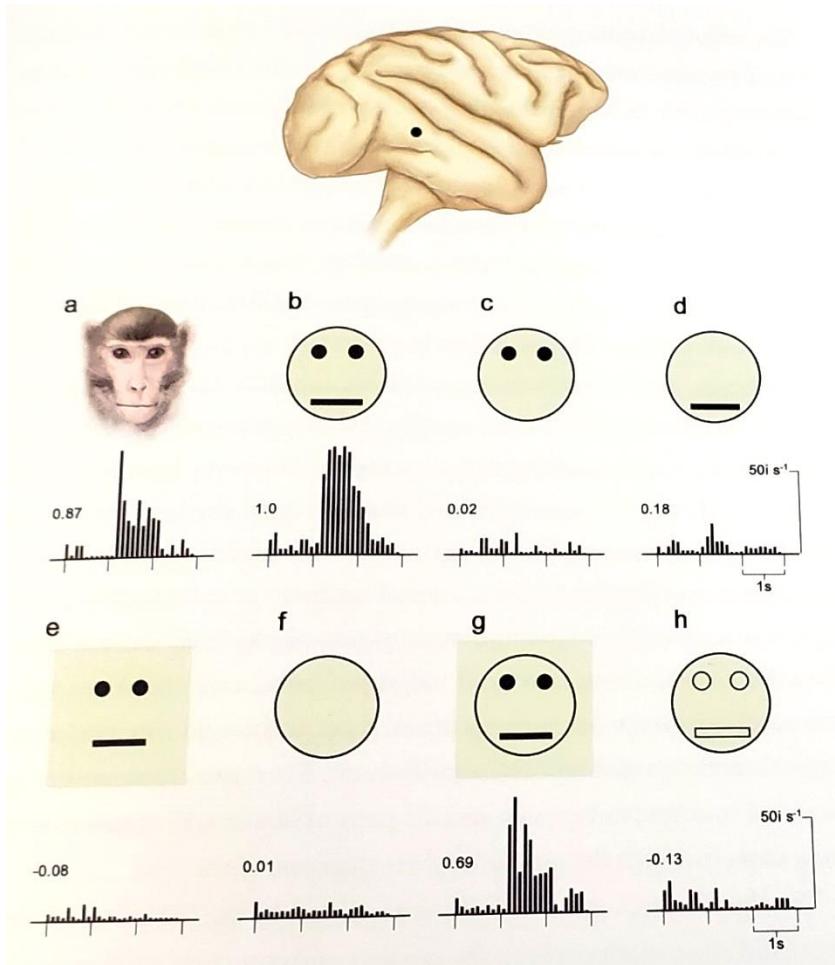
The human visual system



- Retina: visual input
- Retina → Lateral Geniculate Nucleus
 - Visual information flows through the optic nerve
- Lateral Geniculate Nucleus (LGN):
 - A small, ovoid object at the end of the optic tract
 - One on each side of the brain
 - In humans, each LGN has six layers of neurons
 - Sends information to the primary visual cortex (V1)

Image Courtesy: Kandel, "Reductionism in Art and Brain Science," 2016

e.g. Facial Recognition

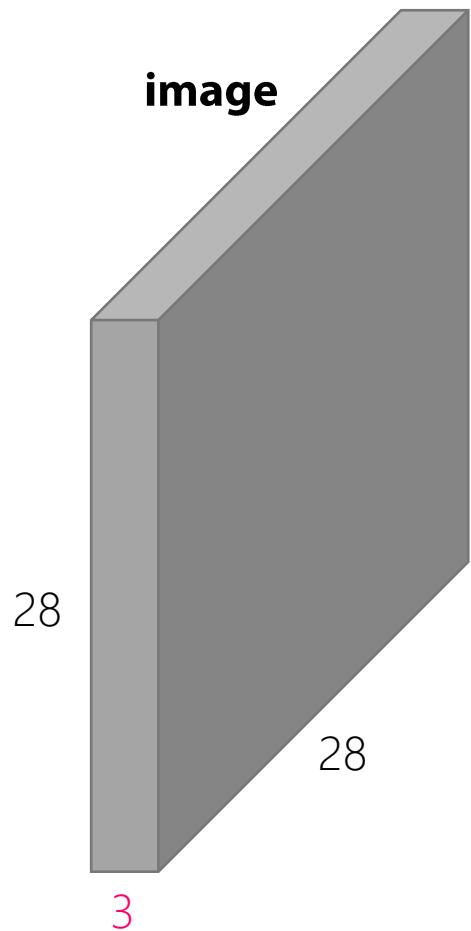


- Holistic face detection
 - “Face cell” in the inferior temporal cortex
 - Fires when there is a face-like object

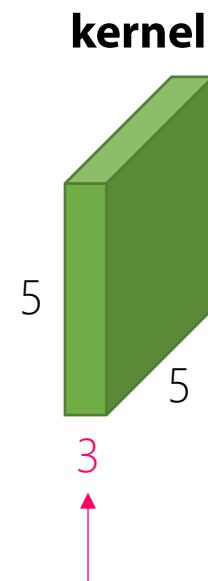
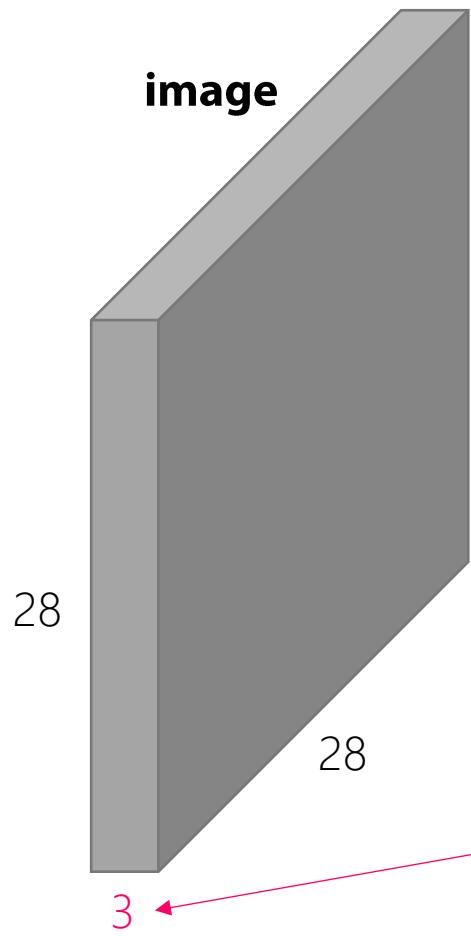
Convolutional Neural Networks

- Key idea:
 - As like how we (humans) understand a visual scene, if neural nets could see small pieces, understand patterns and textures, combine the pieces to see a bigger picture, computers should be able to recognize images.
- A bonus:
 - Typical neural networks are “fully connected”.
 - In an image domain, this means all the pixels are interconnected.
 - However, pixels far apart have no significant meaning...
 - By connecting only the nearing neighbors, computational load could be much lower.

Convolution Layer

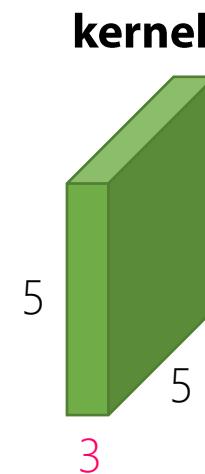
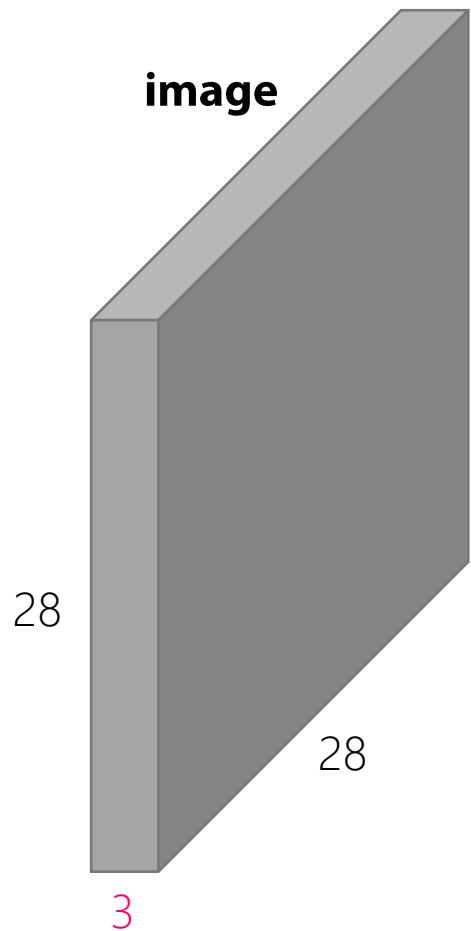


Convolution Layer



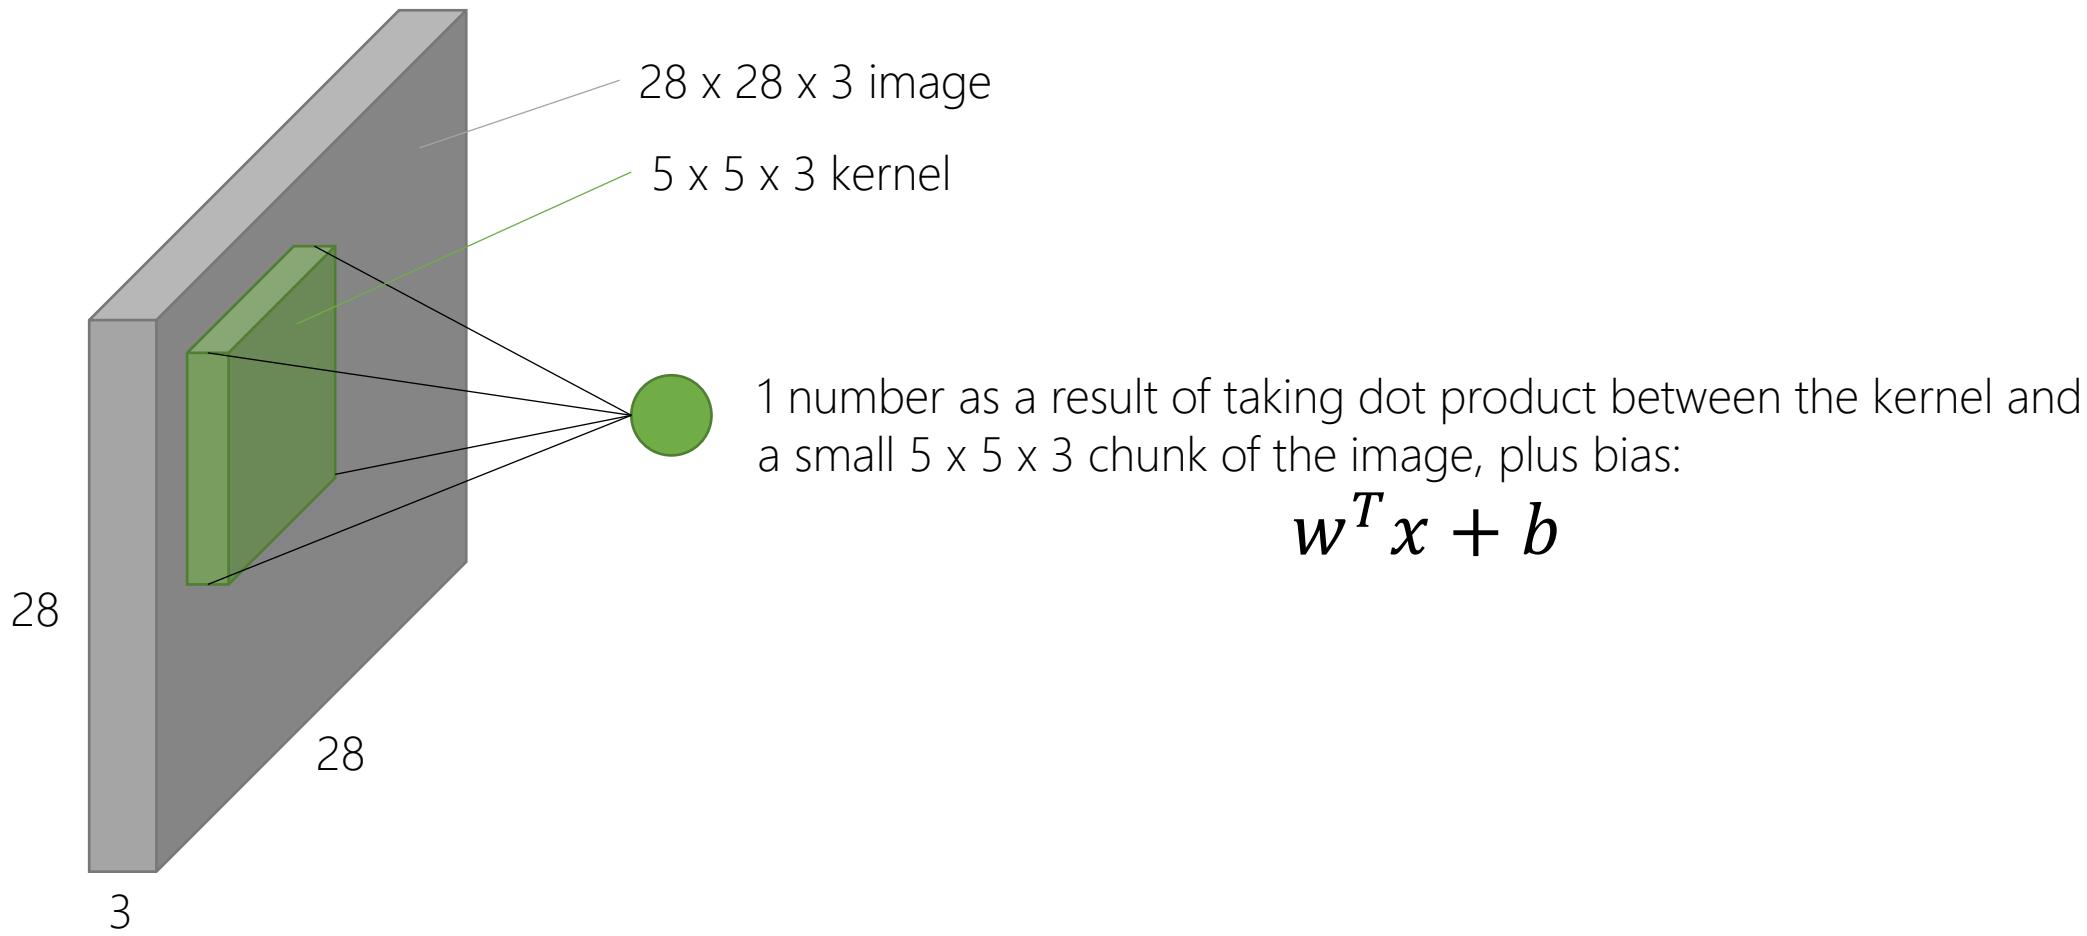
Depth of the kernel must be the same as the depth (i.e. number of channels) of the input image.

Convolution Layer

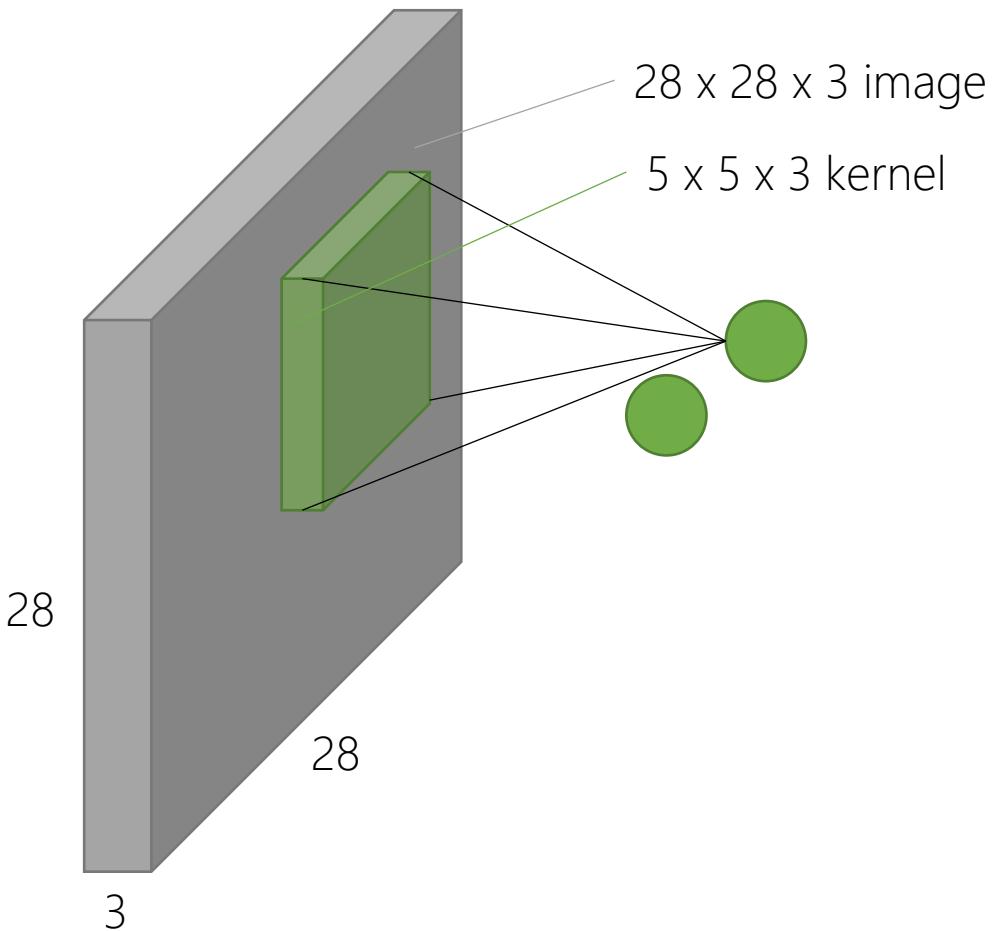


Convolve the kernel with the image!
In other words...
“slide the kernel over the image,
compute dot products each time.”

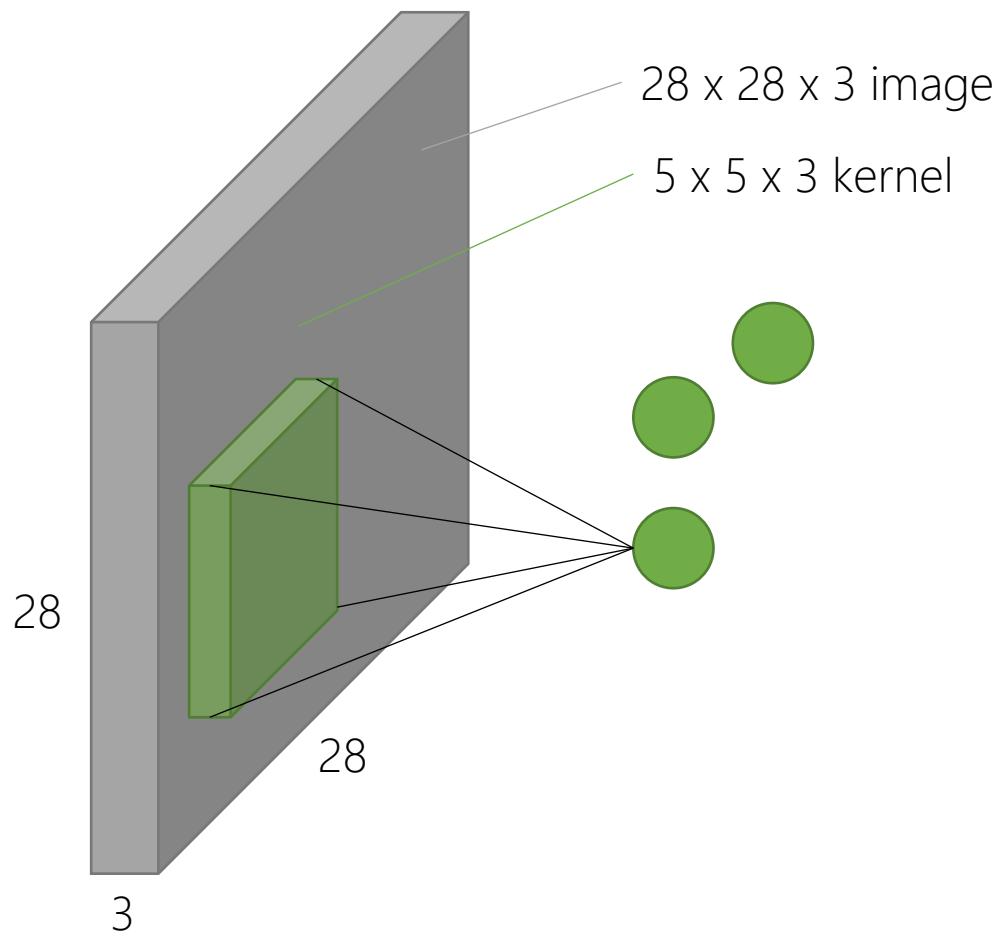
Convolution Layer



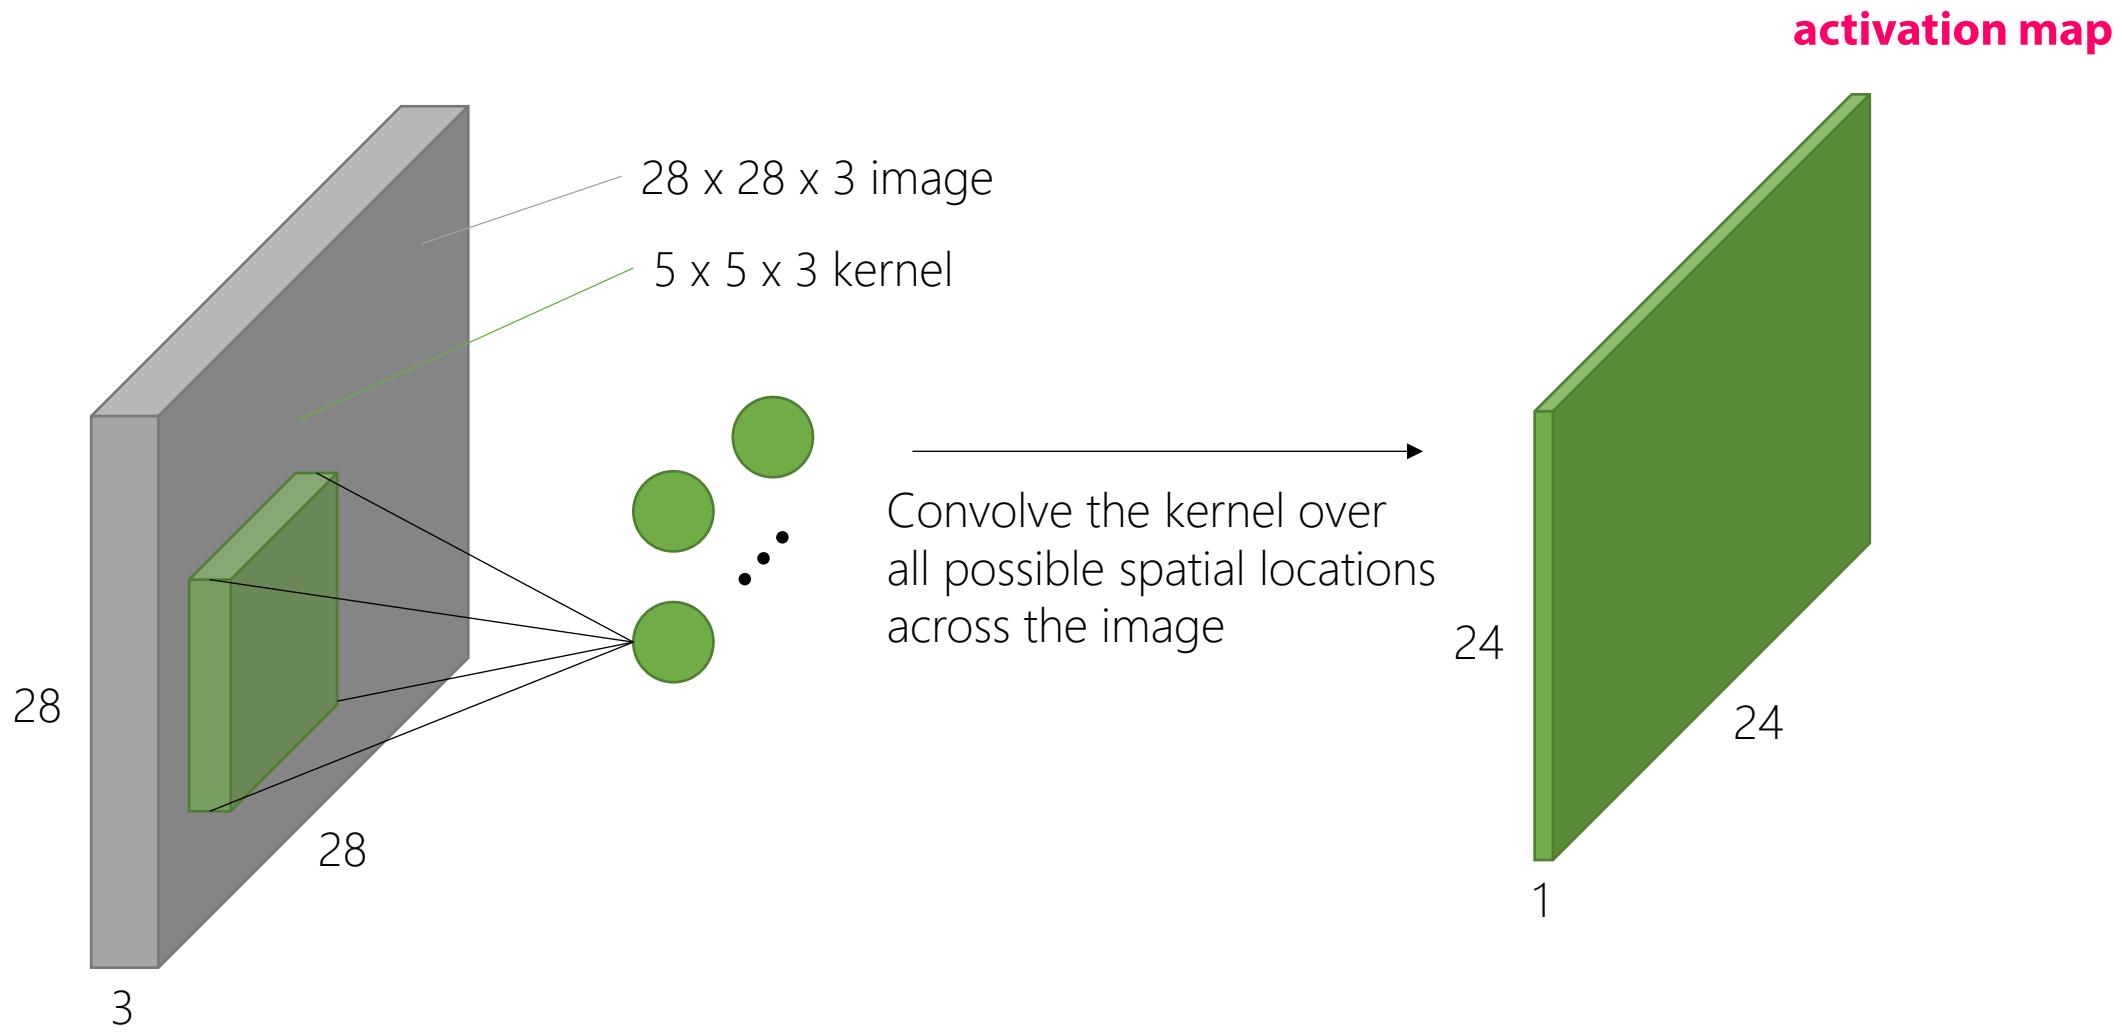
Convolution Layer



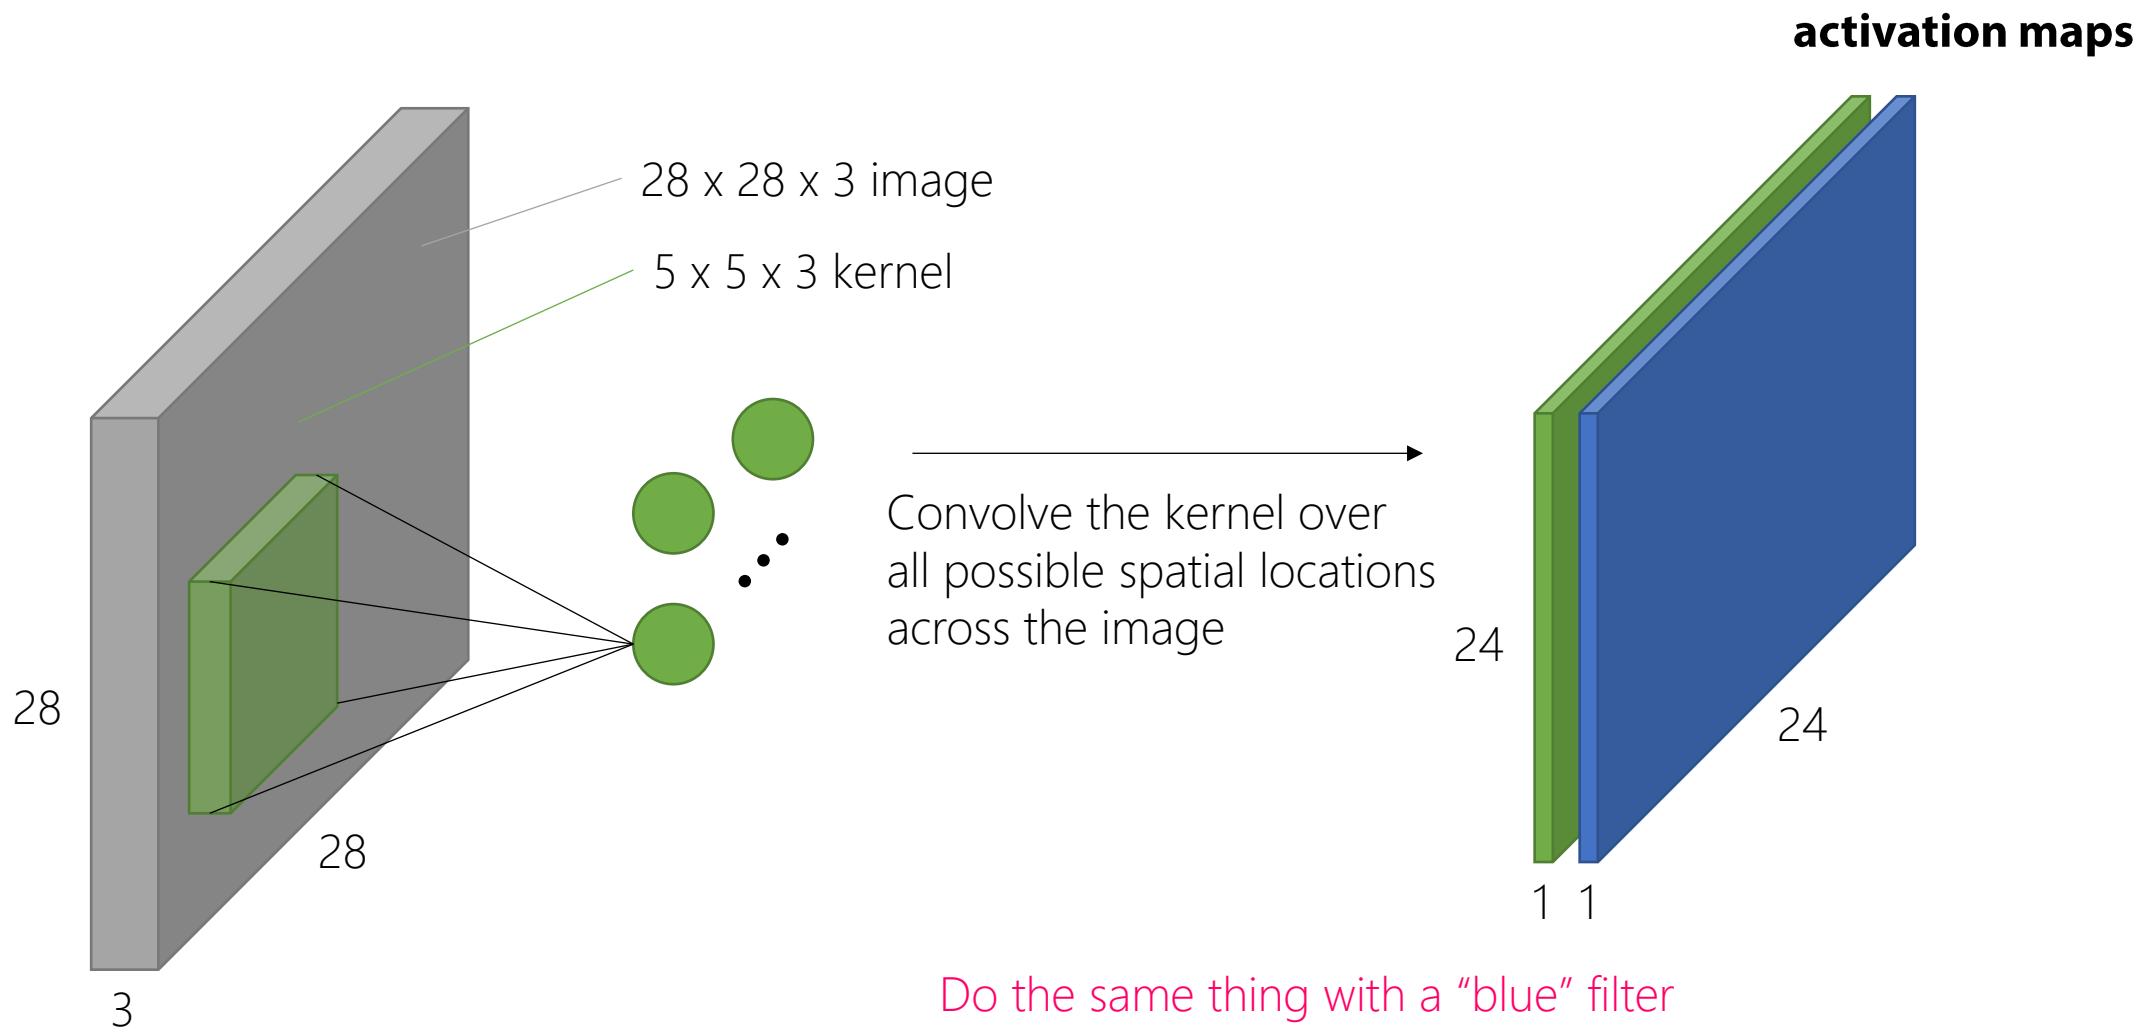
Convolution Layer



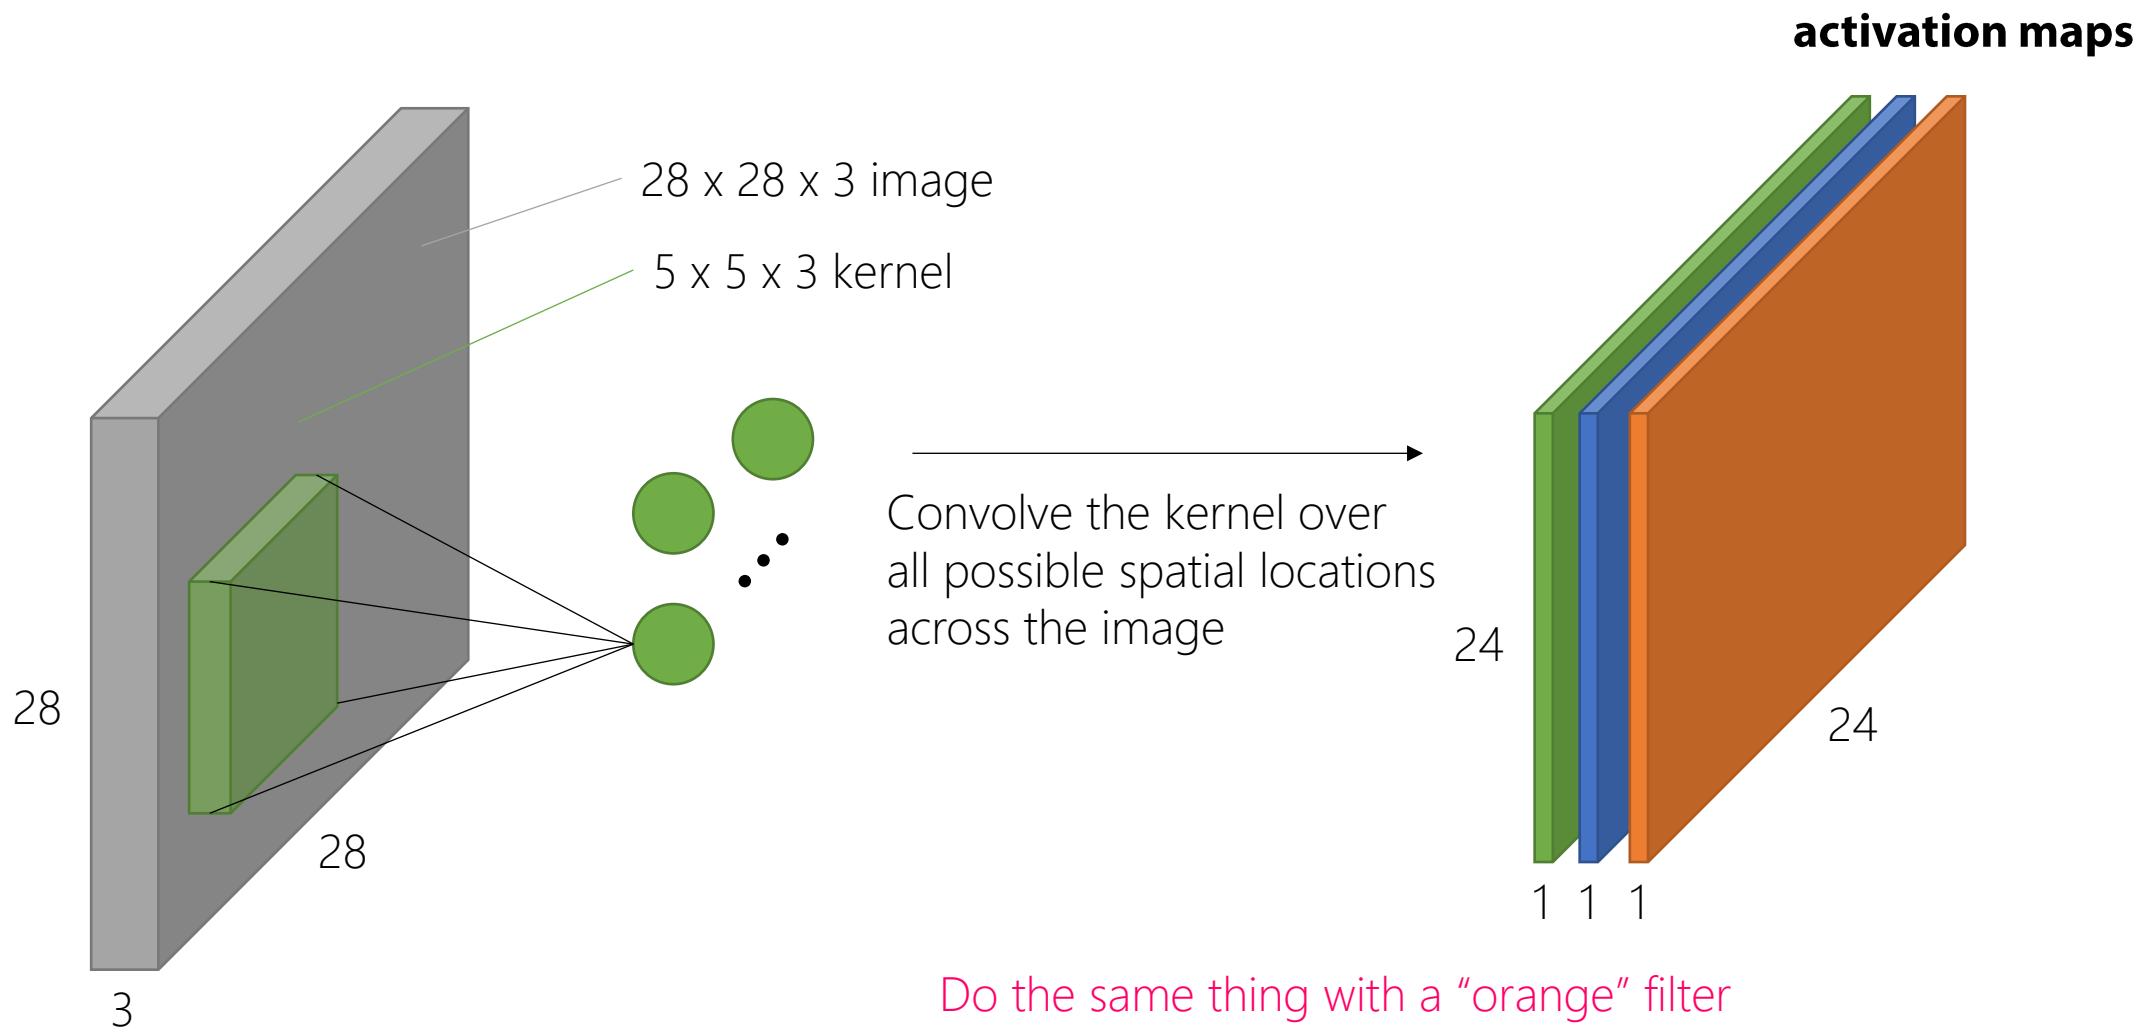
Convolution Layer



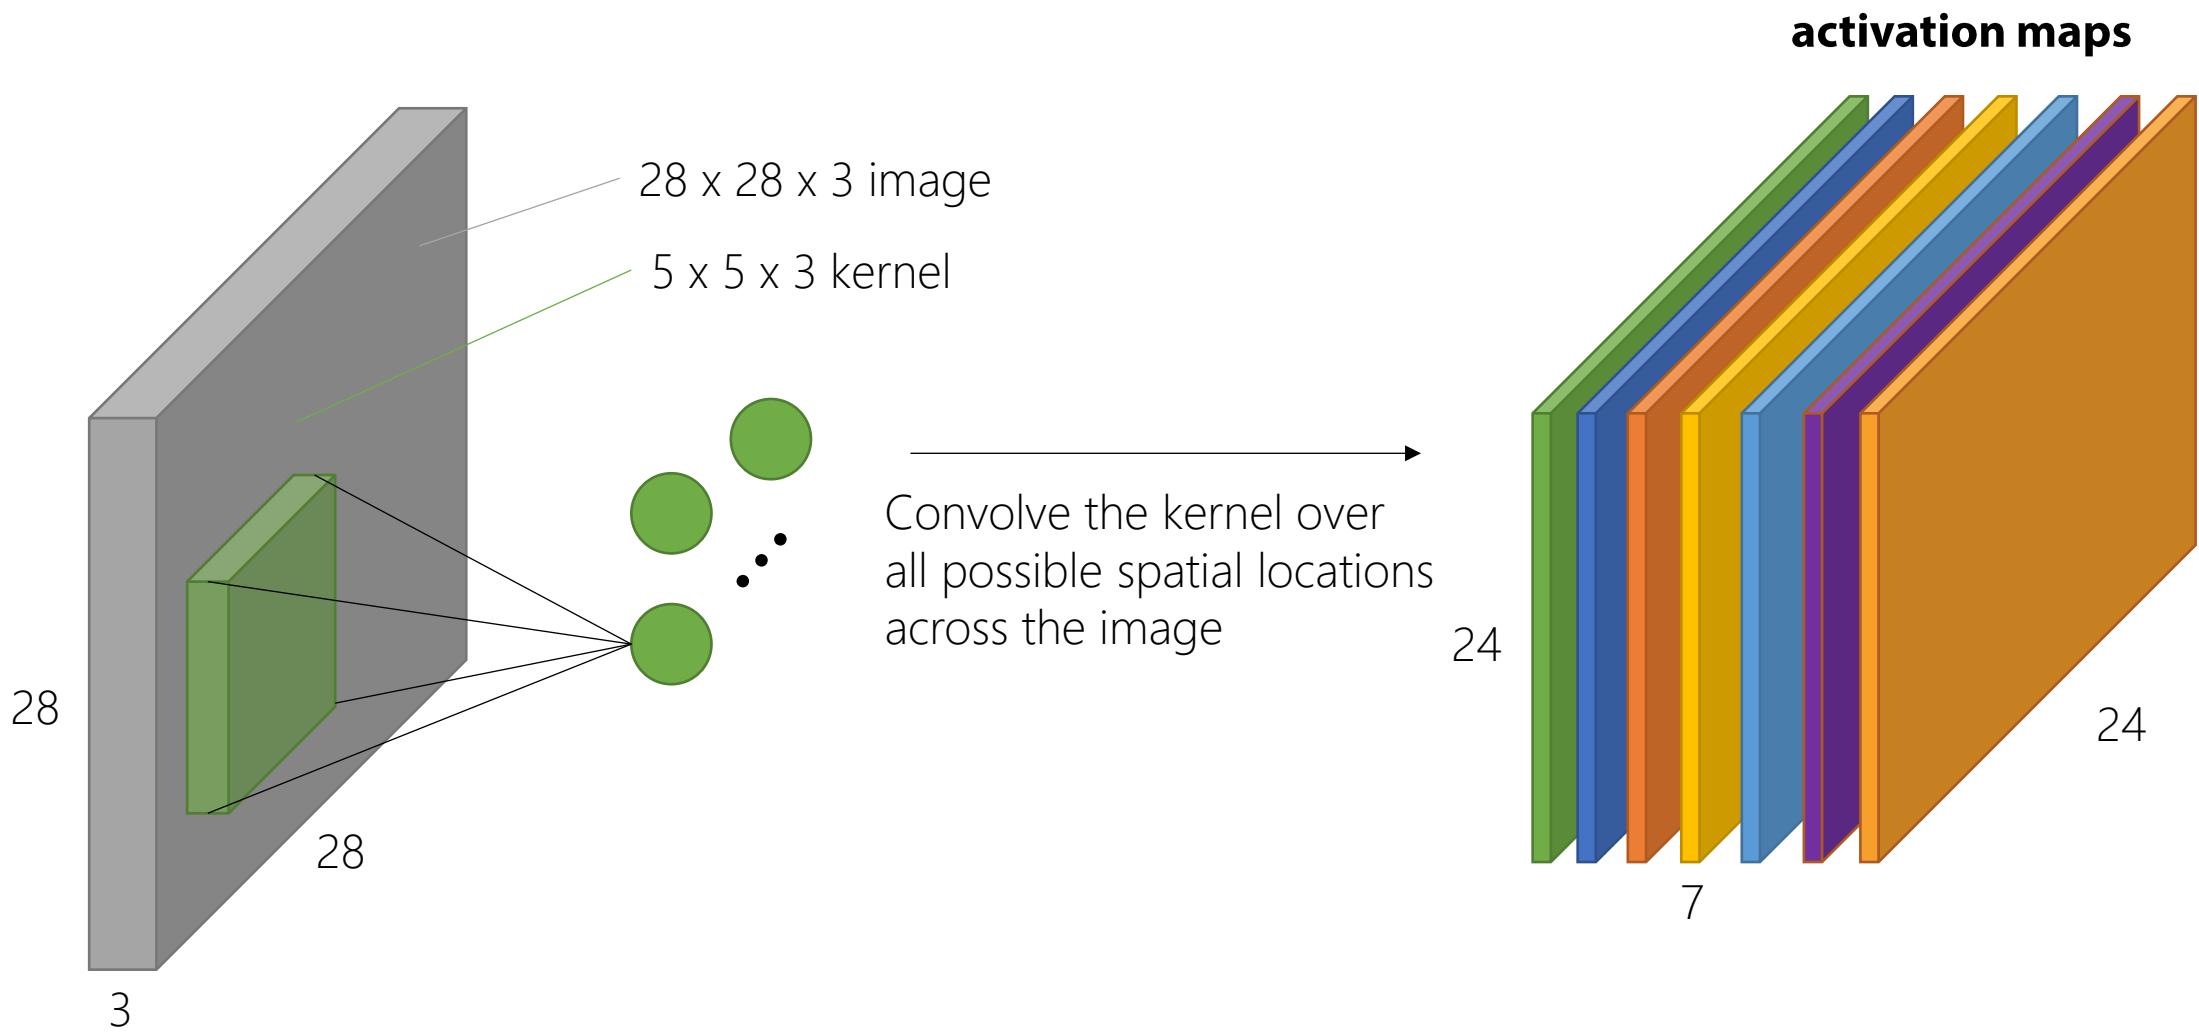
Convolution Layer



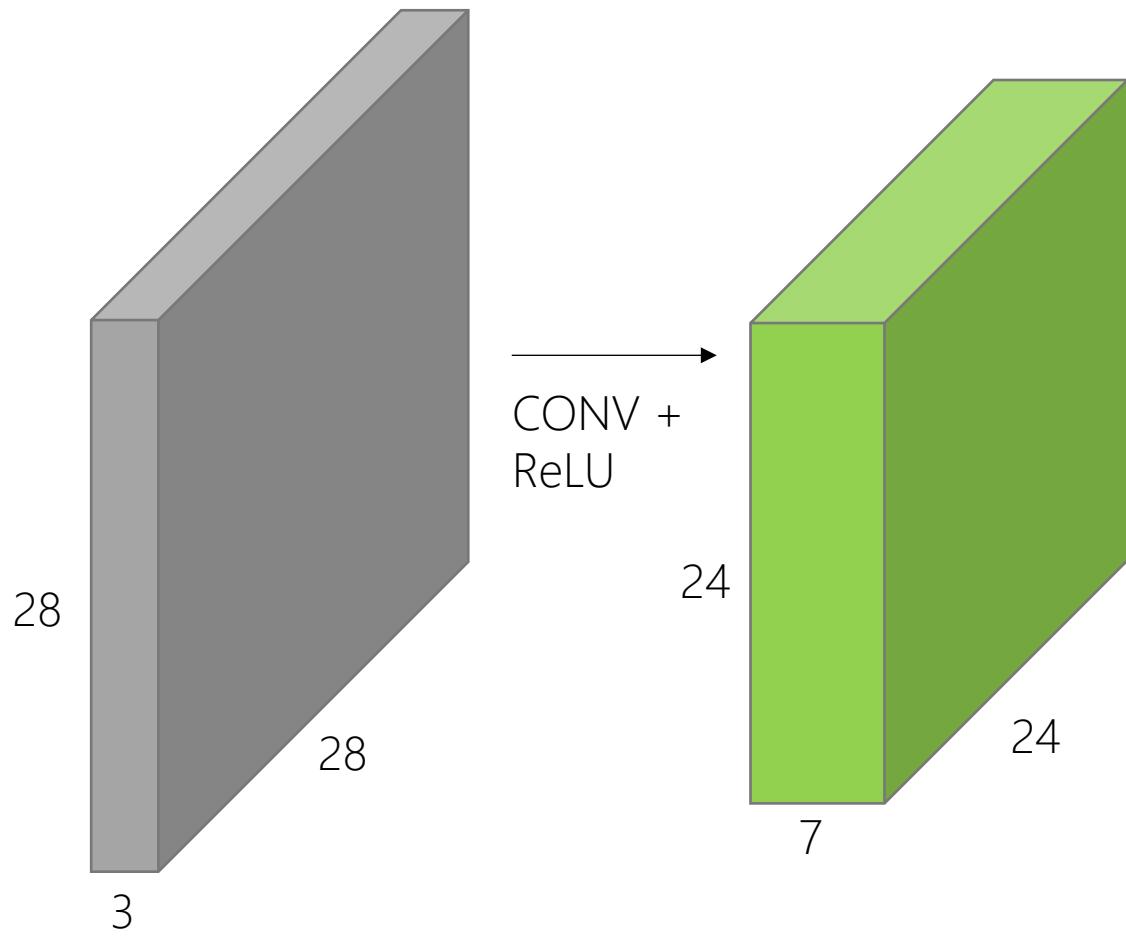
Convolution Layer



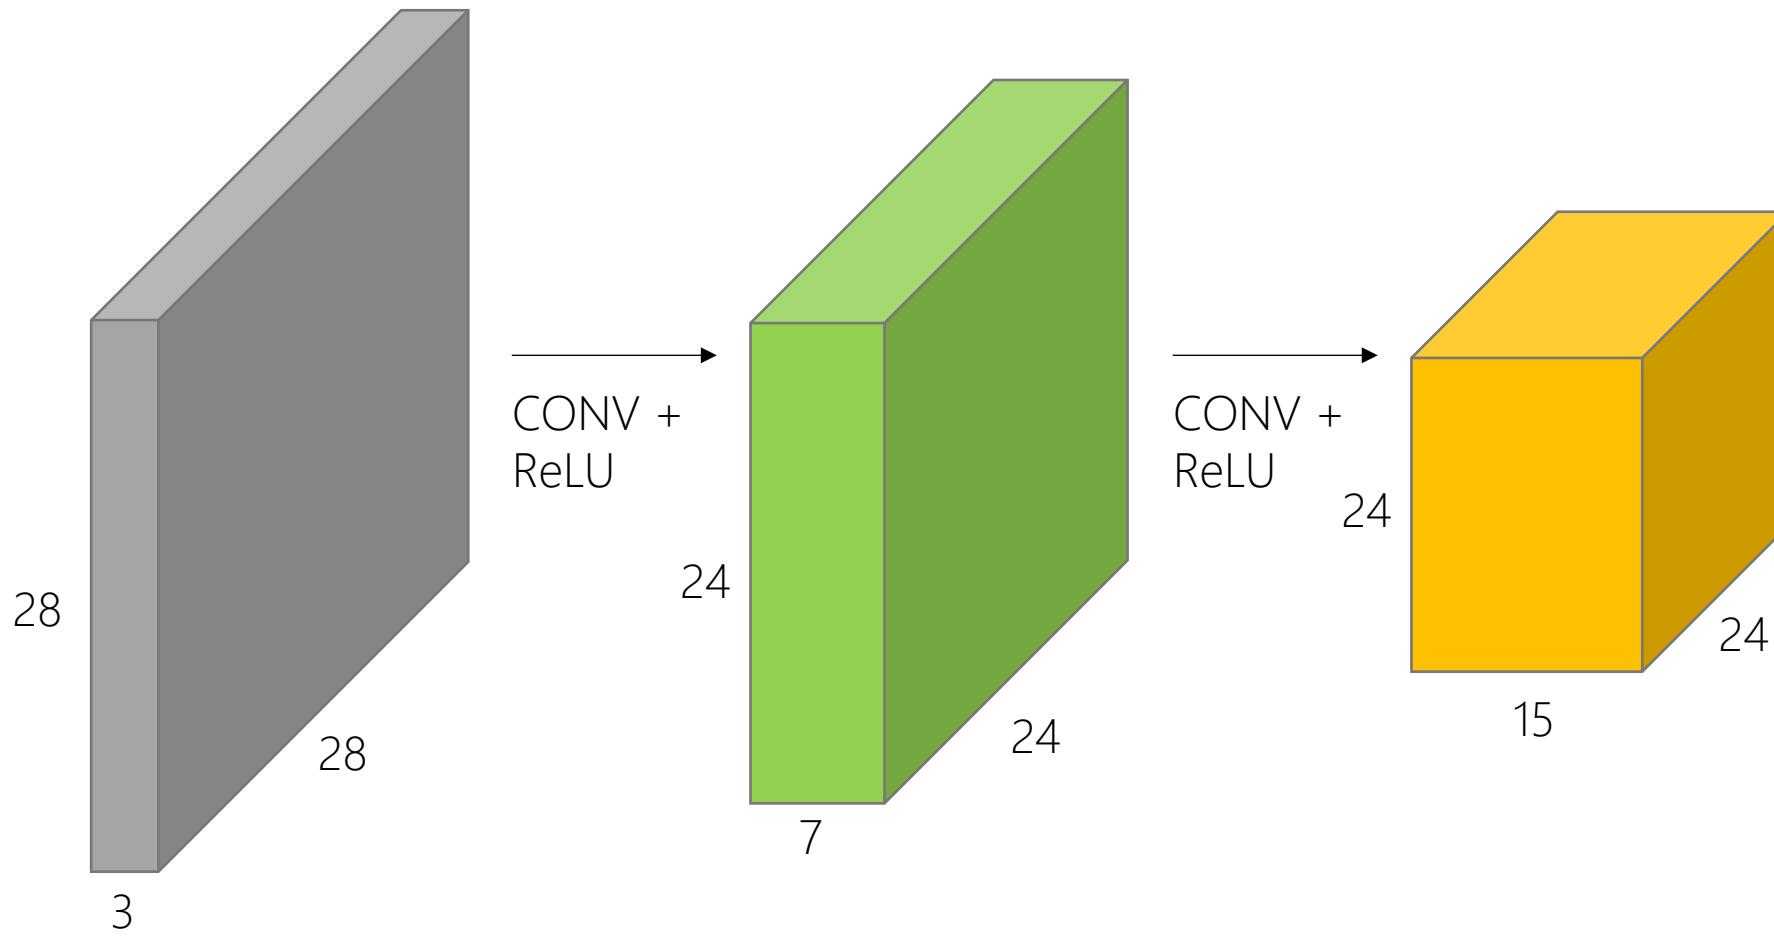
Convolution Layer



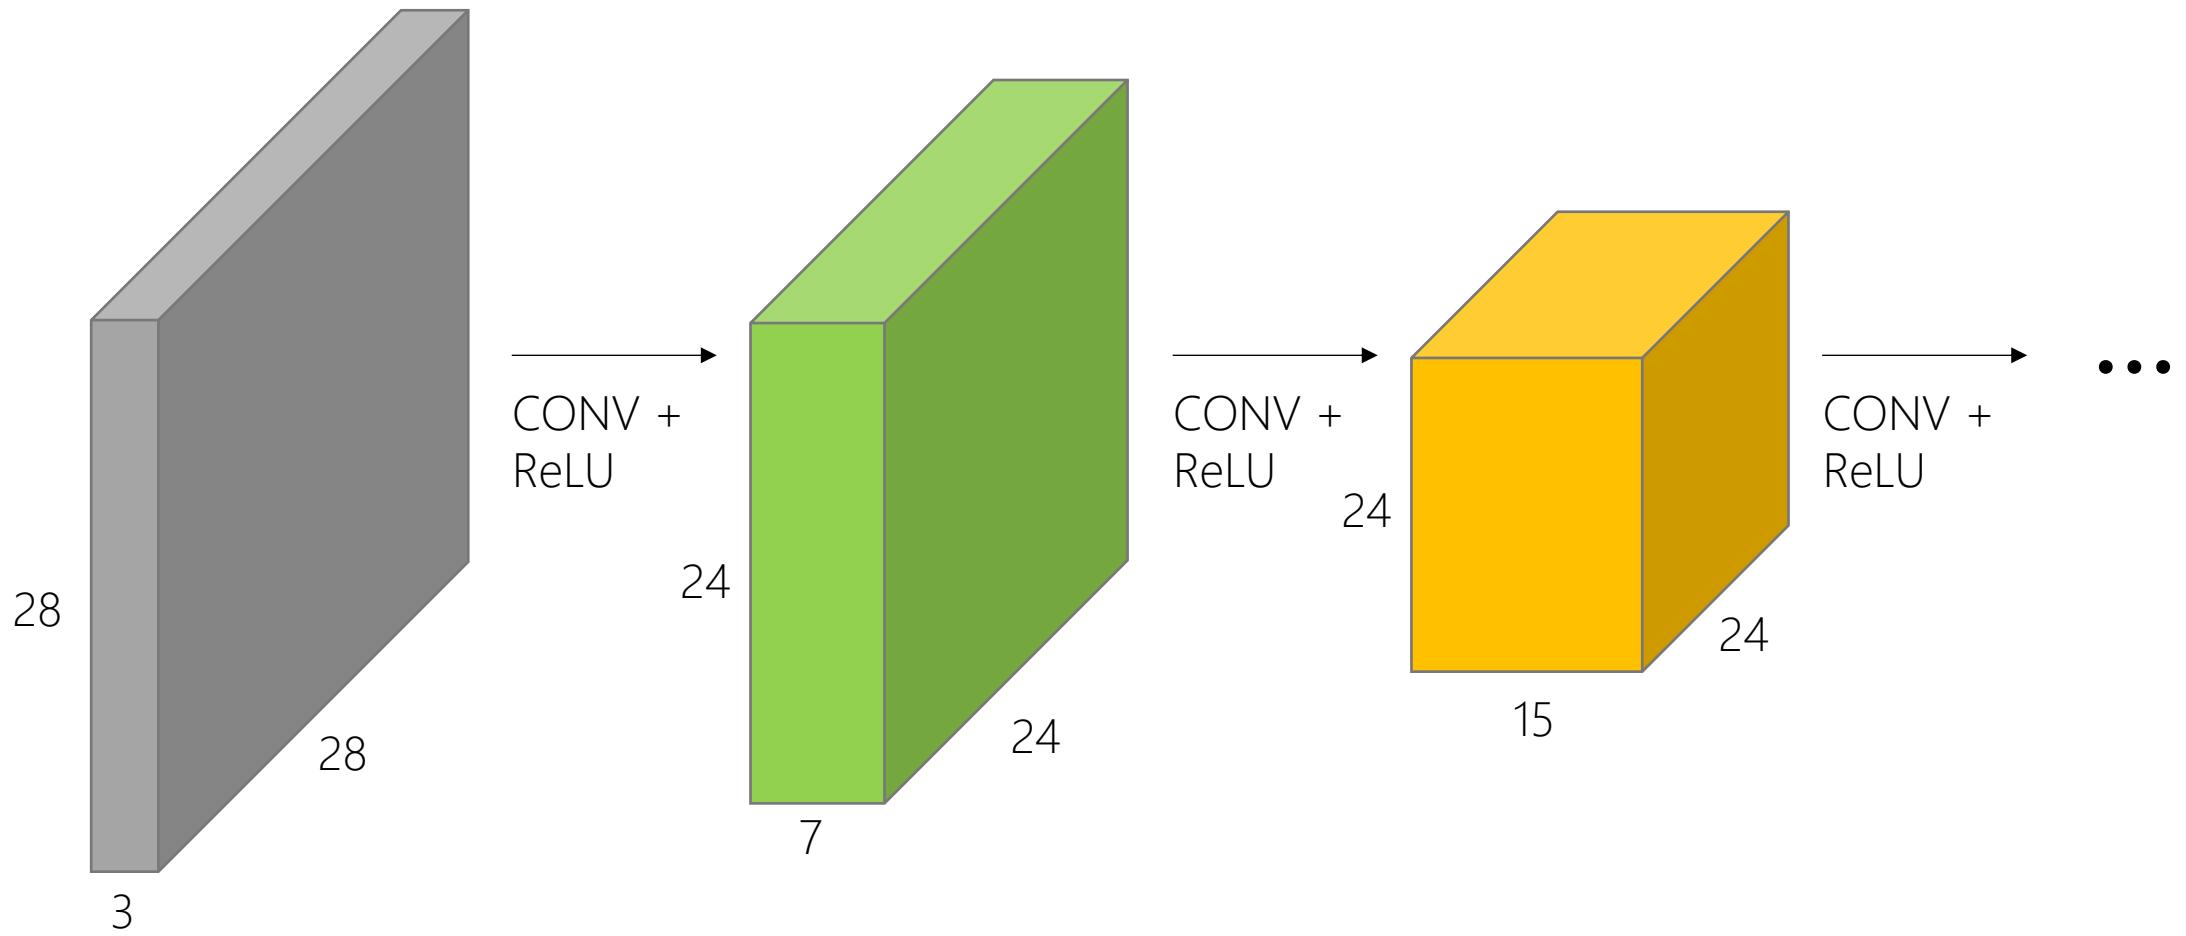
ConvNet is a sequence of convolution layers



ConvNet is a sequence of convolution layers



ConvNet is a sequence of convolution layers



Draw your number here



Downsampled drawing: **2**

First guess: **2**

Second guess: **1**

Layer visibility

Input layer

Show

Convolution layer 1

Show

Downsampling layer 1

Show

Convolution layer 2

Show

Downsampling layer 2

Show

Fully-connected layer 1

Show

Fully-connected layer 2

Show

Output layer

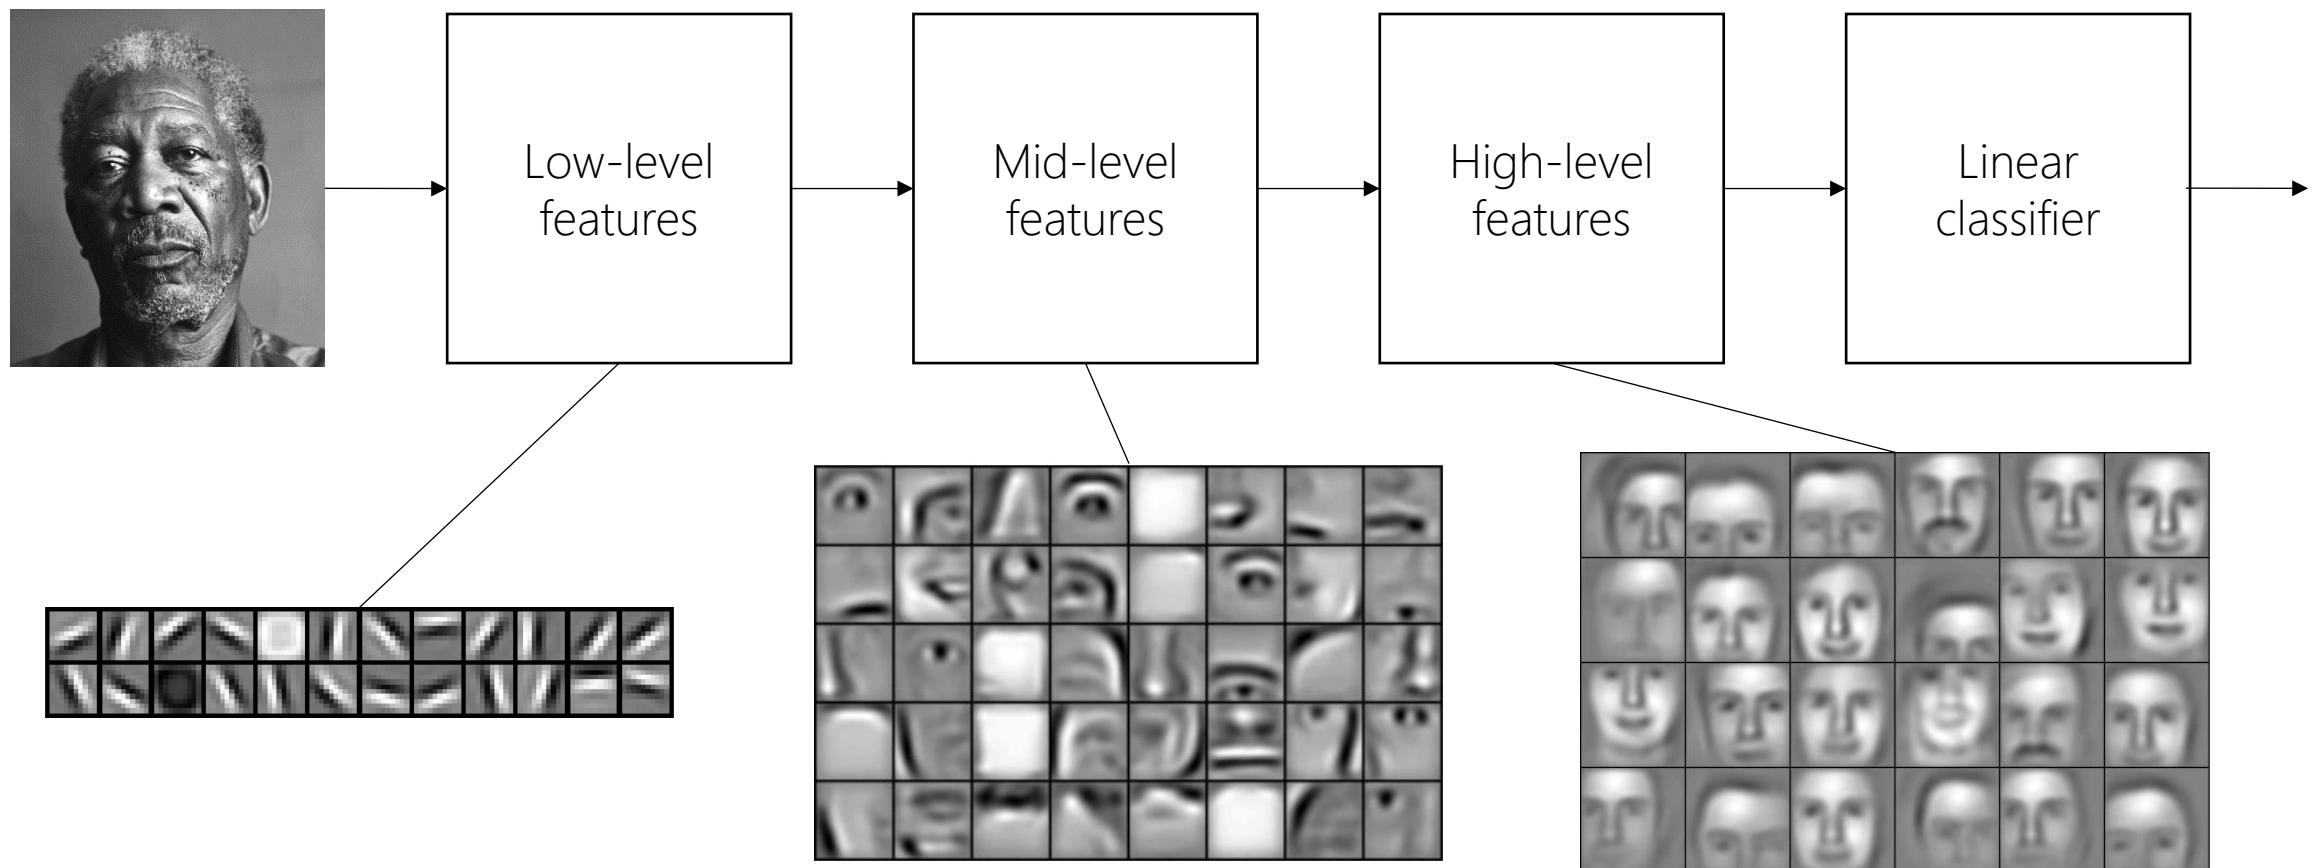
Show

0 1 2 3 4 5 6 7 8 9



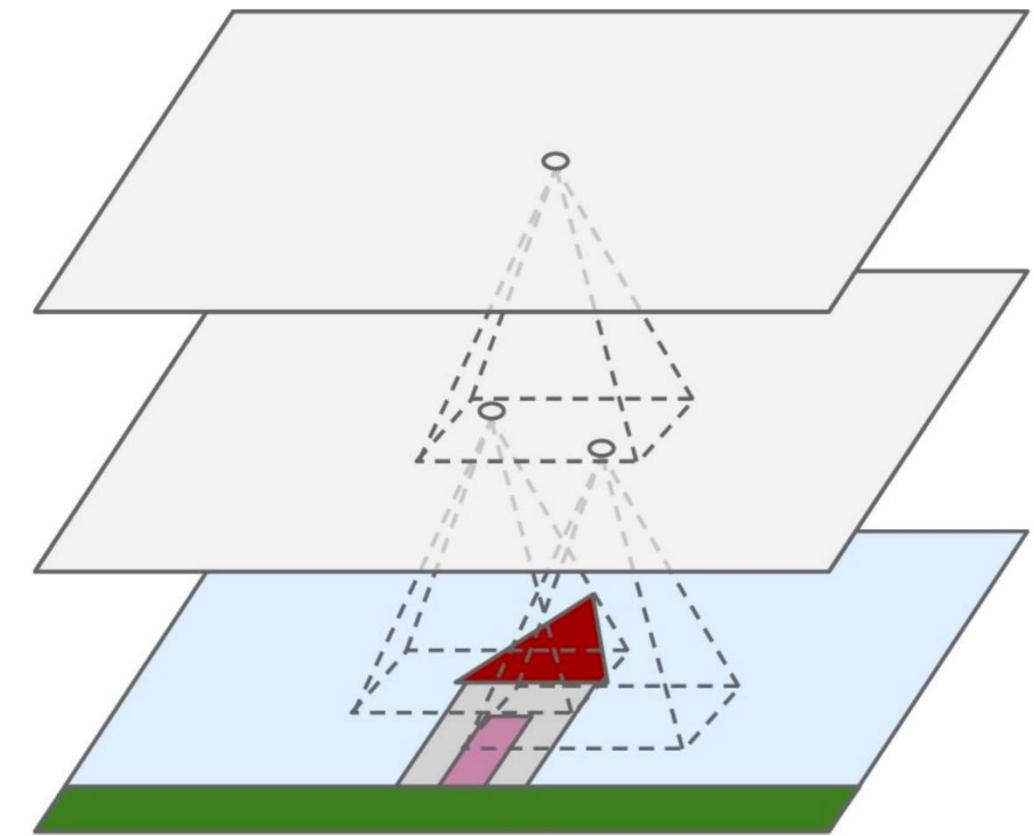
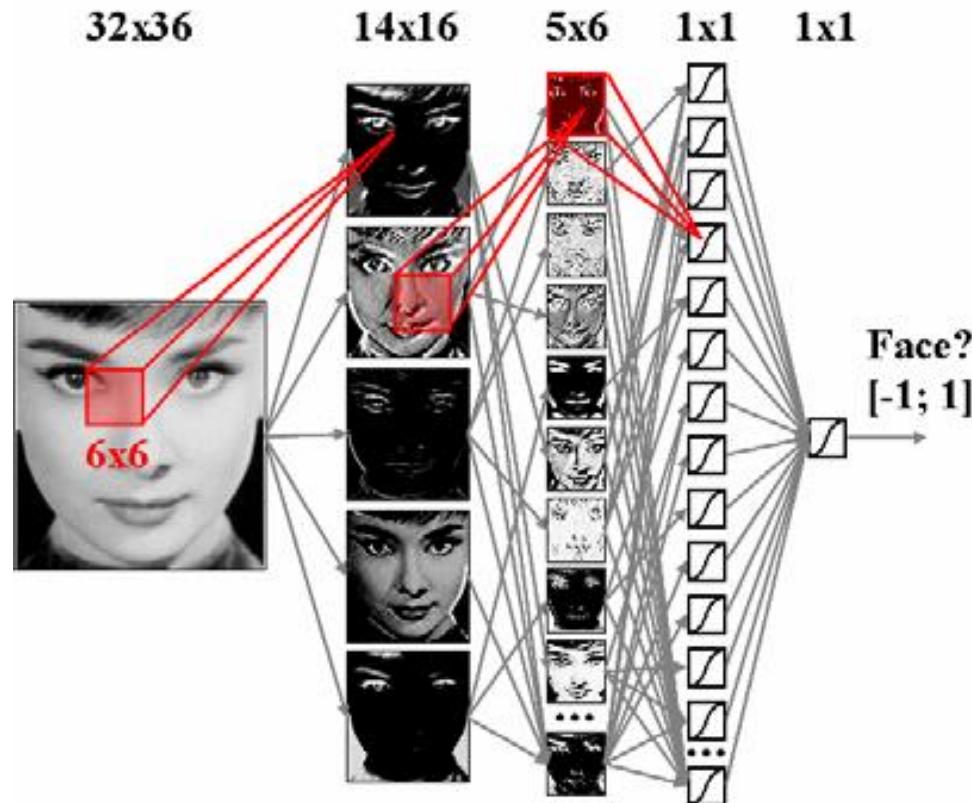
Made by [Adam Harley](#). [Project details](#)

ConvNet is a sequence of convolution layers

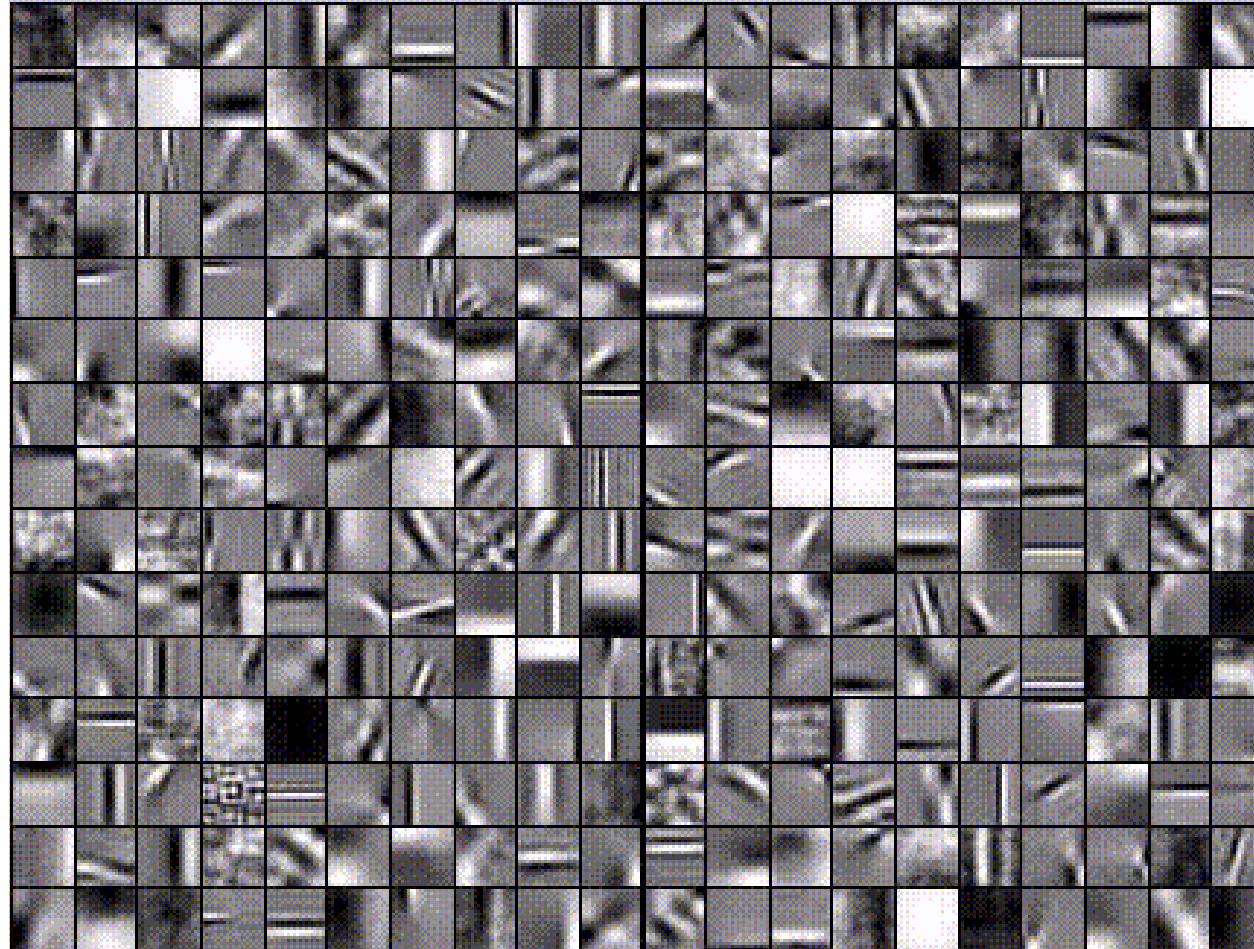


Lee et al. (2009)

ConvNet is a sequence of convolution layers



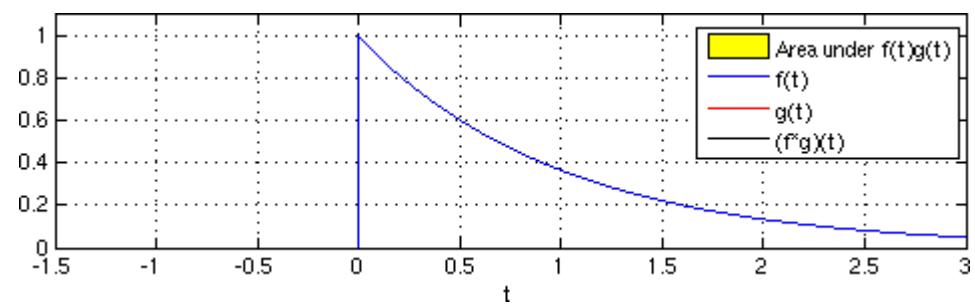
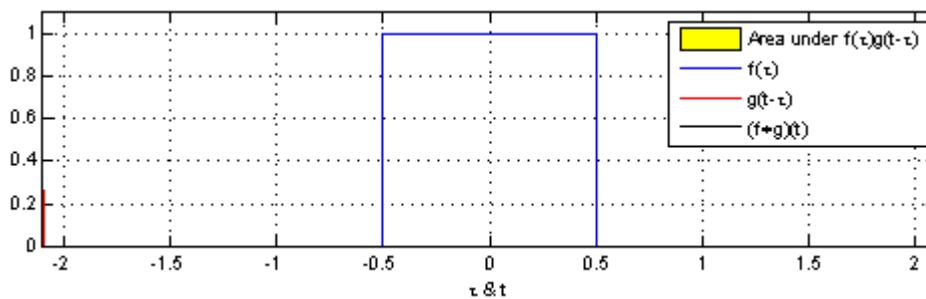
Conv kernels are trainable



A closer look...

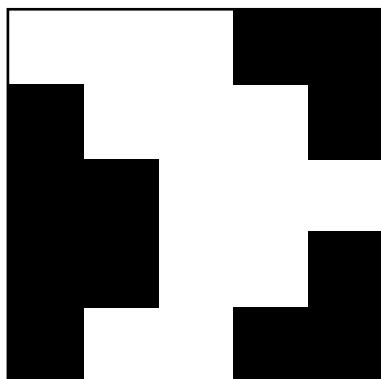
- Convolution

$$(f * g)(t) := \int_{-\infty}^{\infty} f(\tau)g(t - \tau) d\tau$$



A closer look...

- 2D Discrete Convolution



*

$$\begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{bmatrix}$$



1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

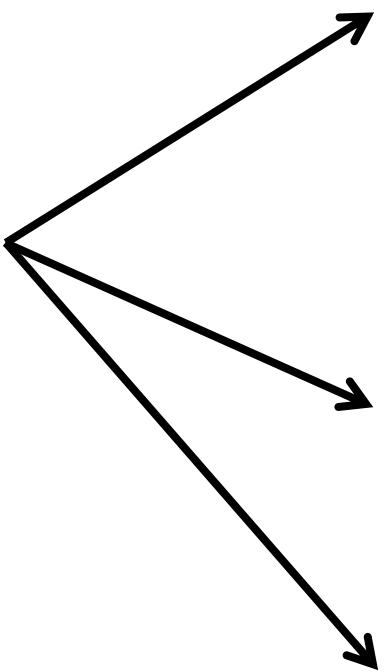
Image

4			

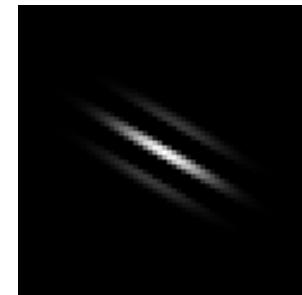
Convolved
Feature

A closer look...

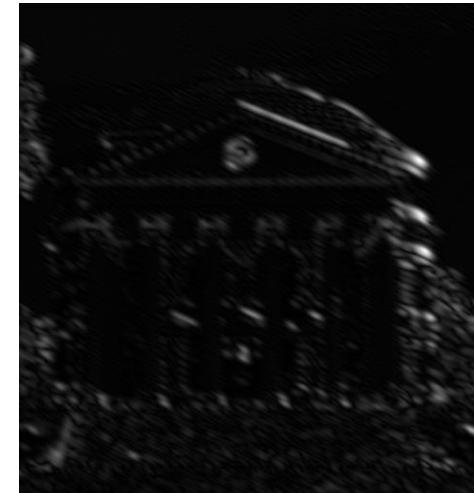
- 2D Discrete Convolution



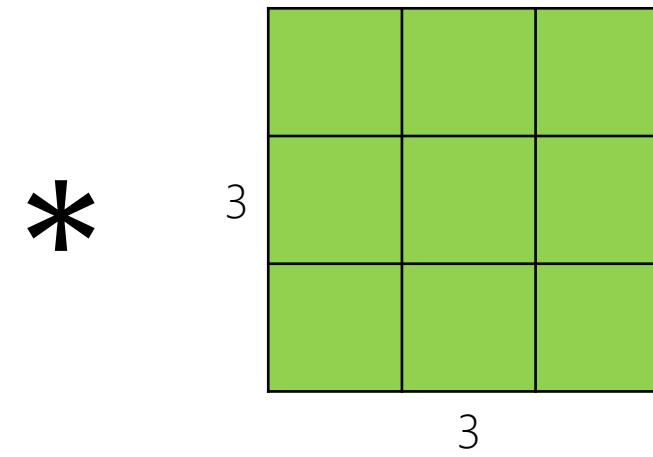
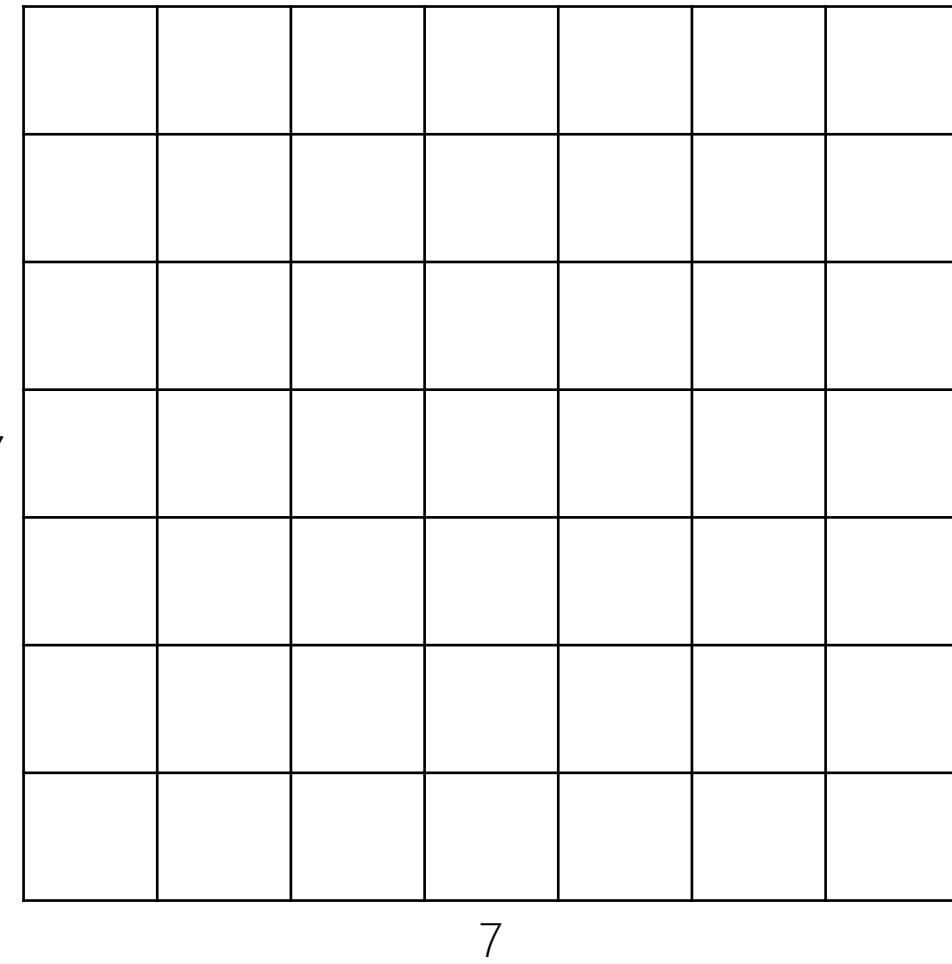
$$\ast \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} =$$



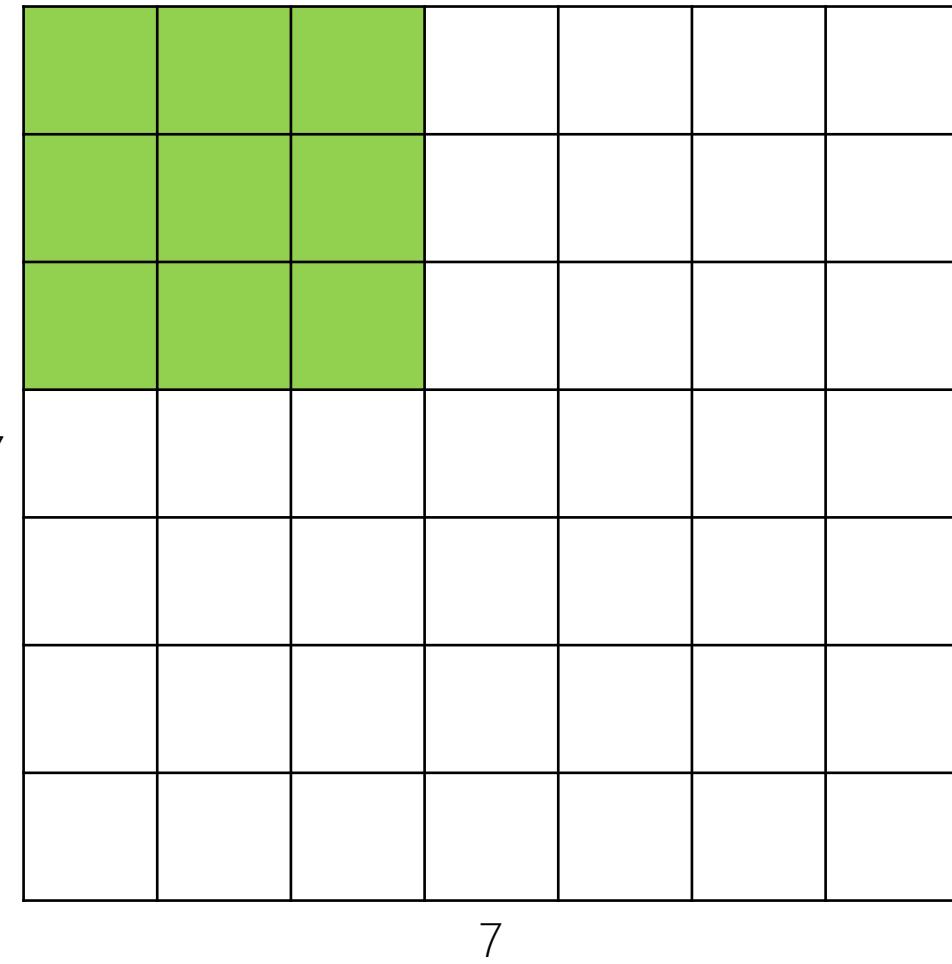
⋮



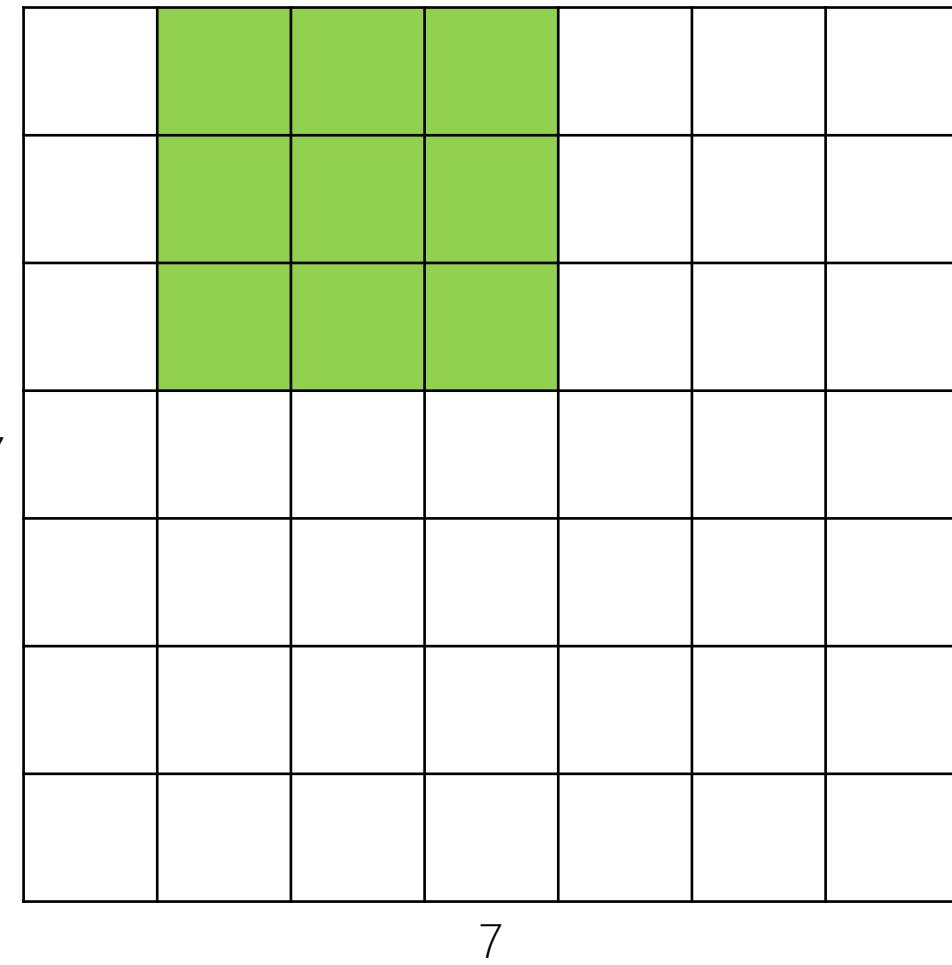
Stride & Padding



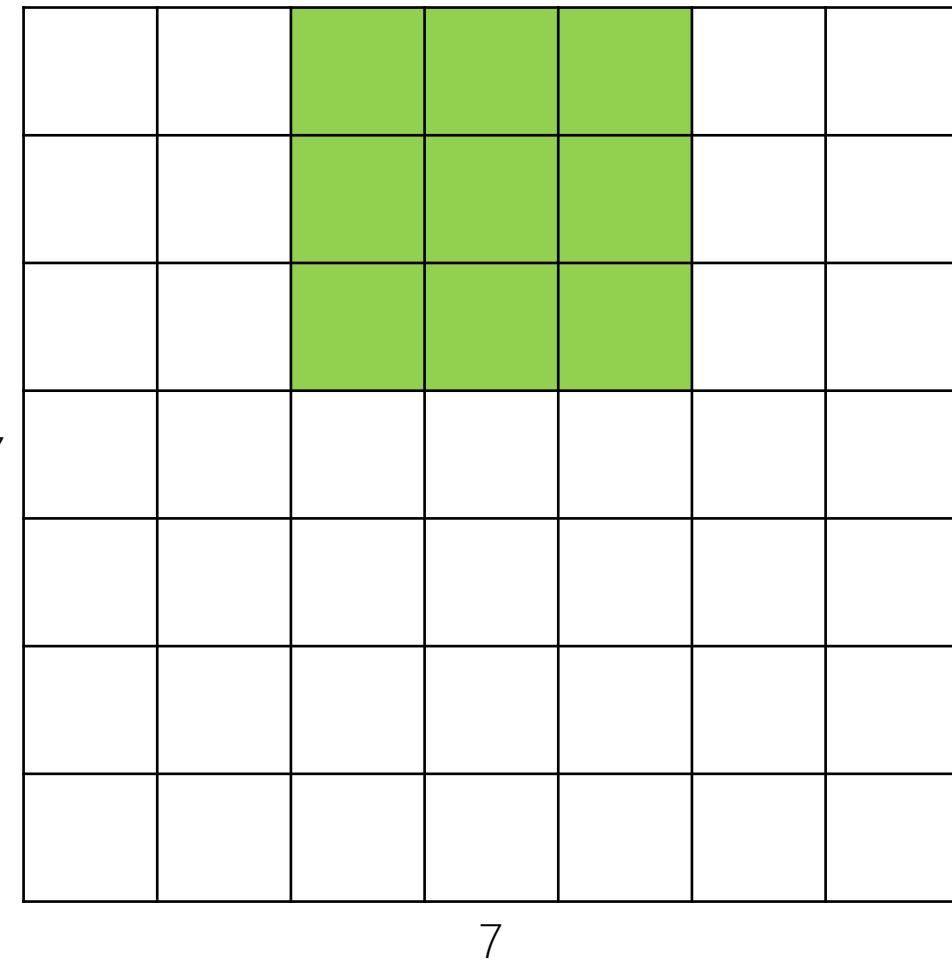
Stride & Padding



Stride & Padding

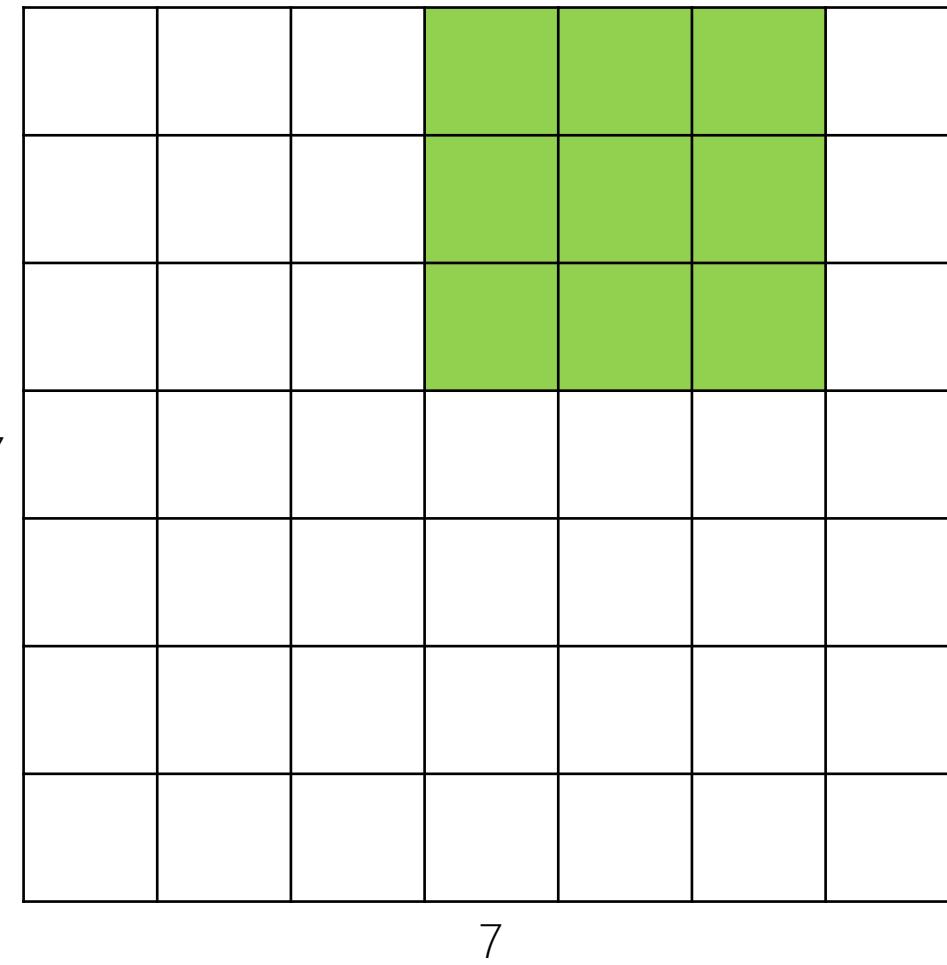


Stride & Padding

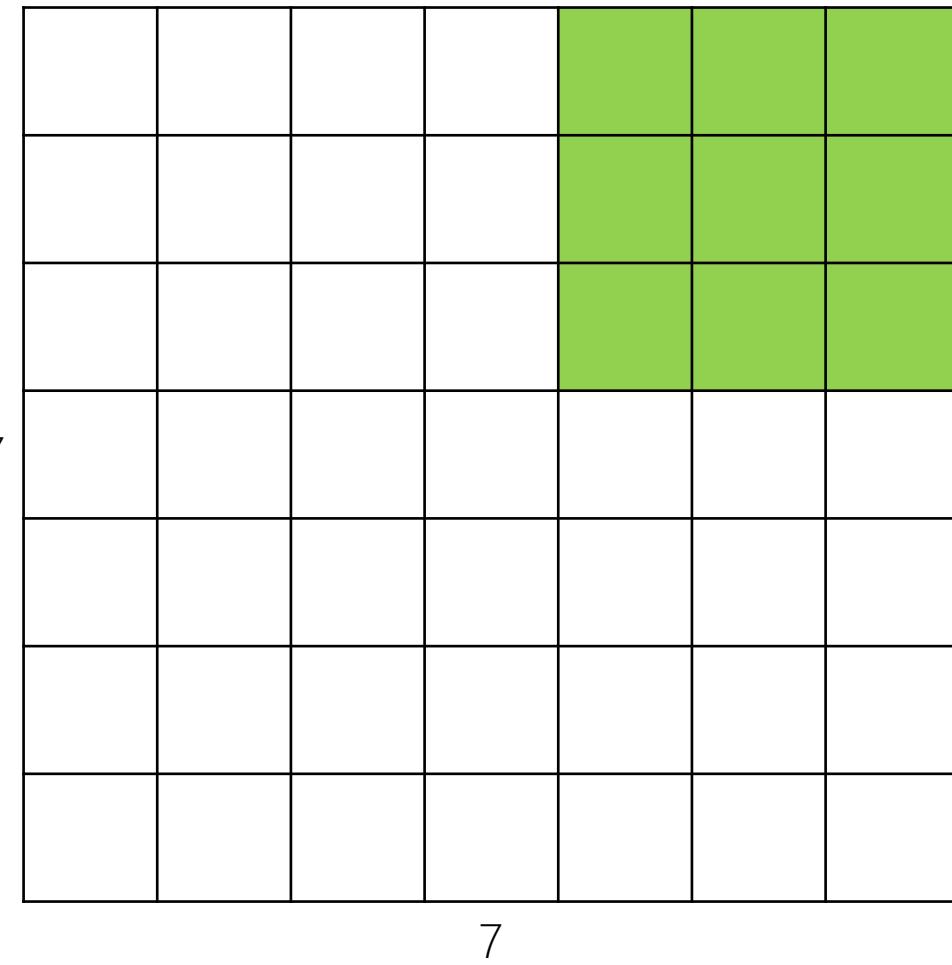


7

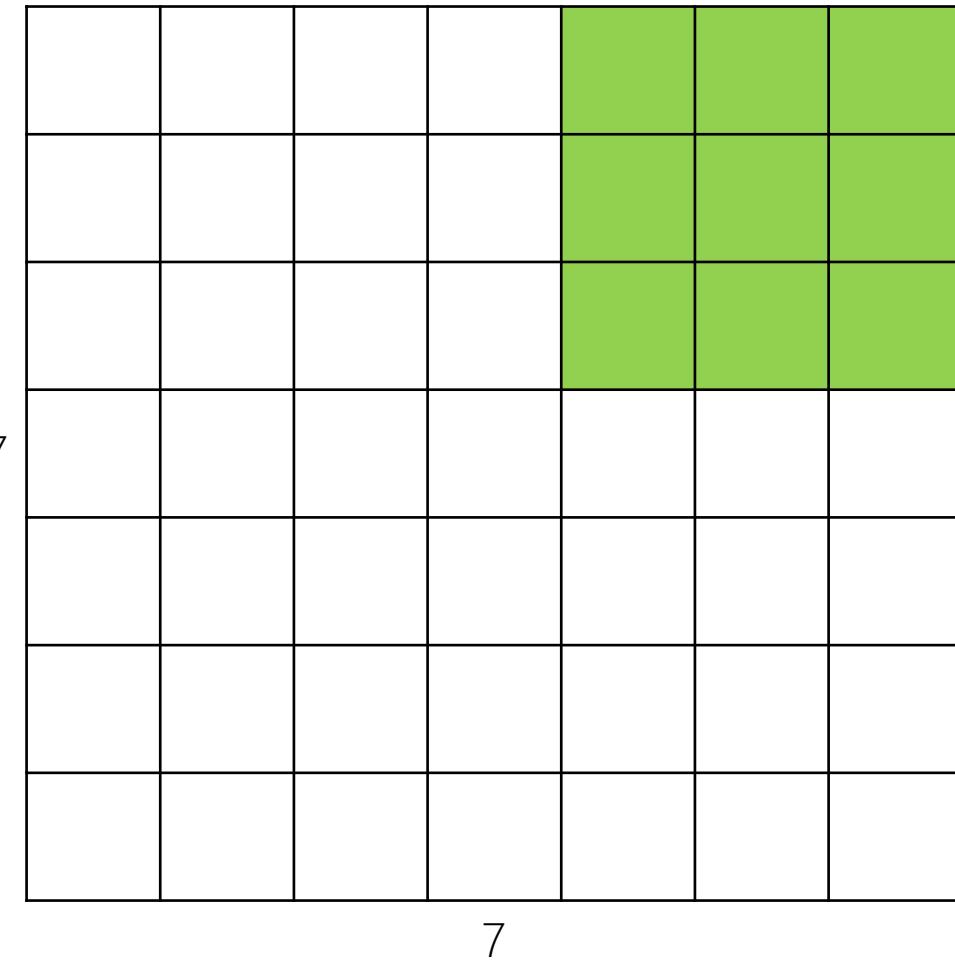
Stride & Padding



Stride & Padding

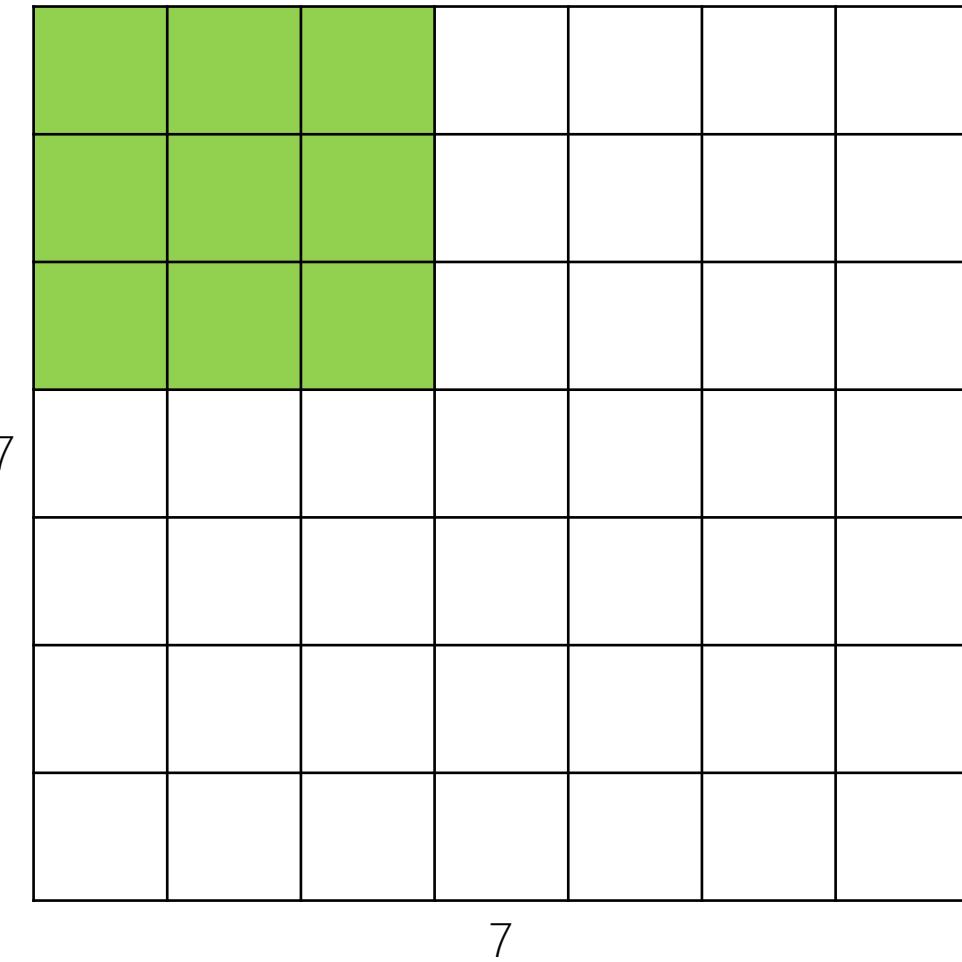


Stride & Padding



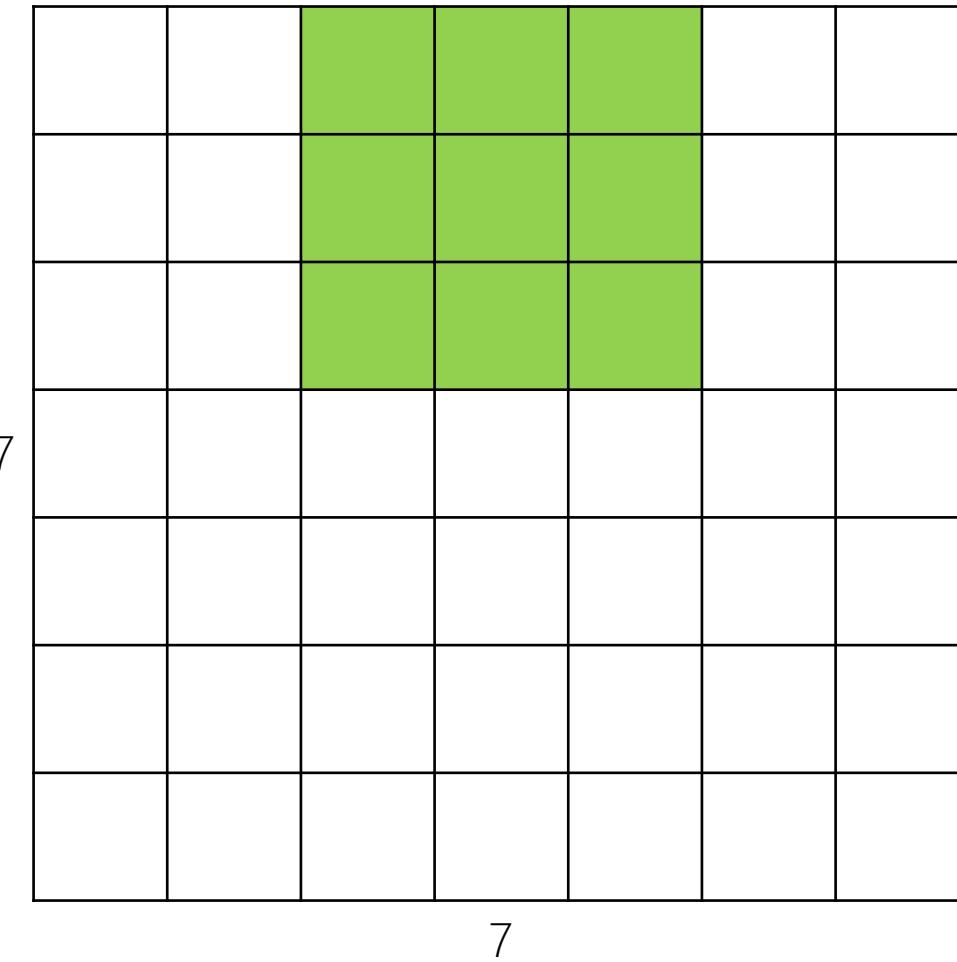
→ 5x5 output

Stride & Padding



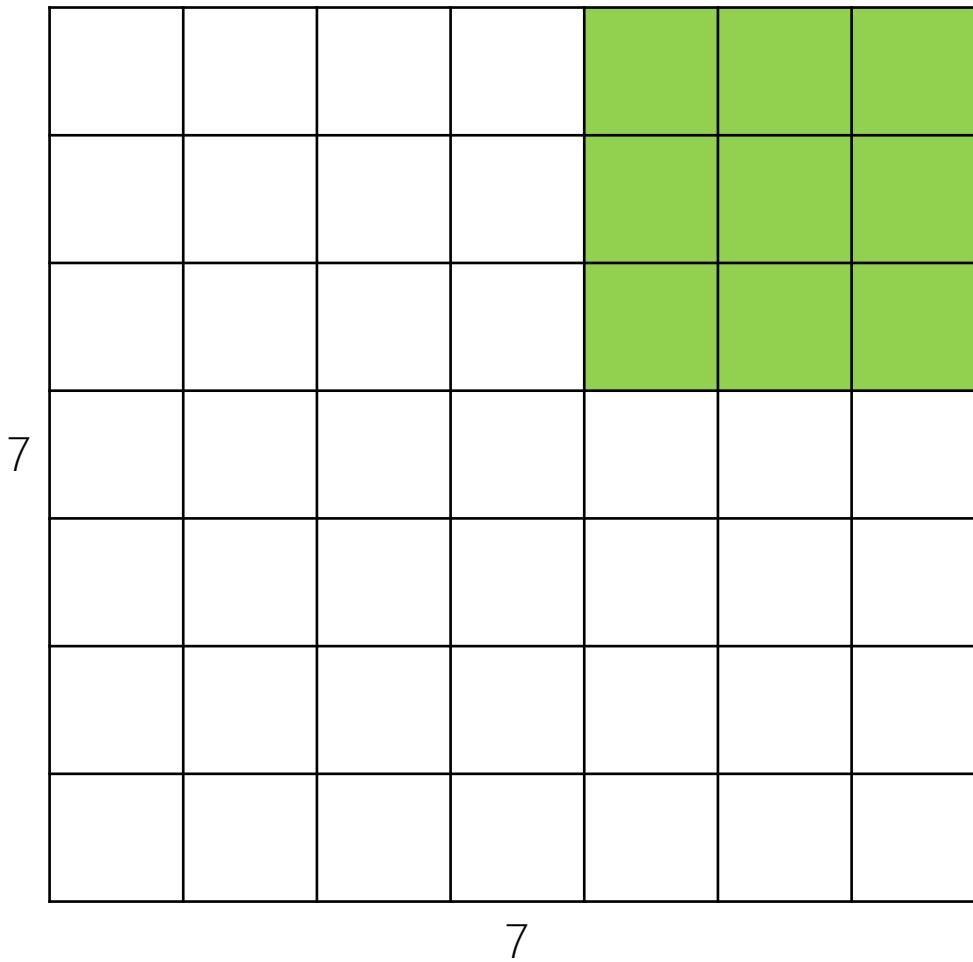
Now with **stride 2**

Stride & Padding



Now with **stride 2**

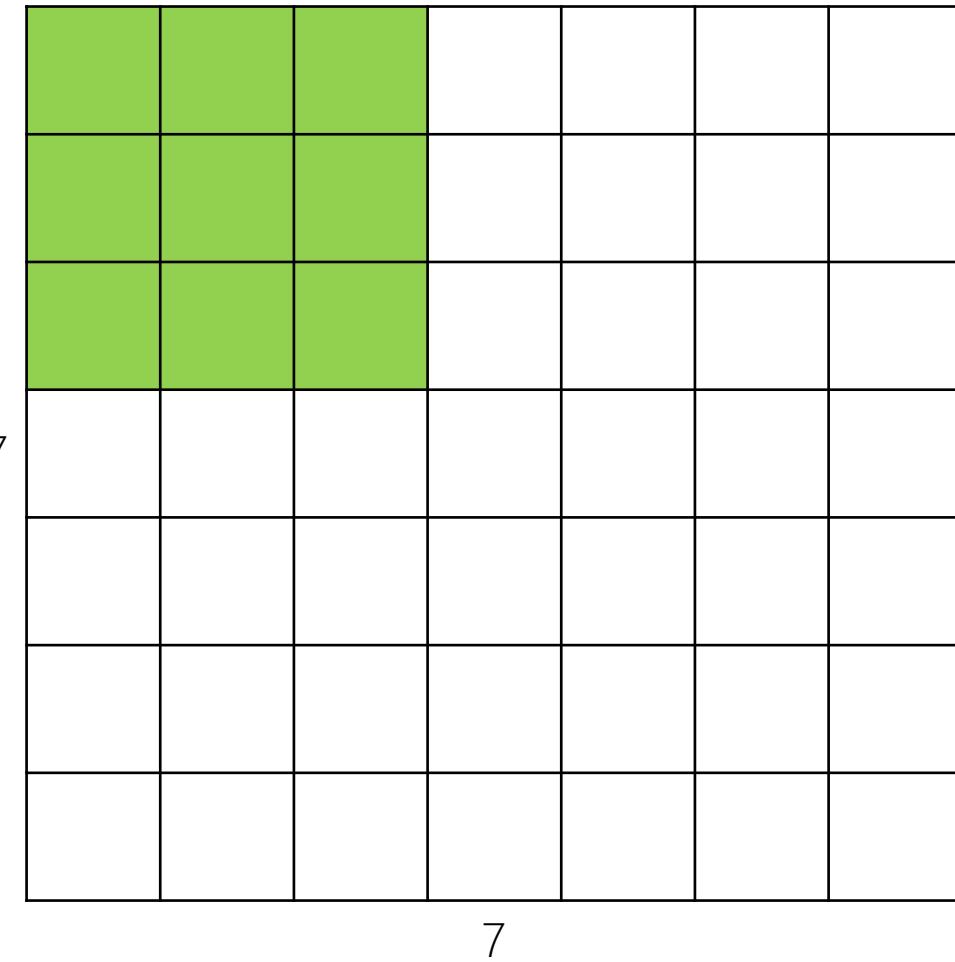
Stride & Padding



Now with **stride 2**

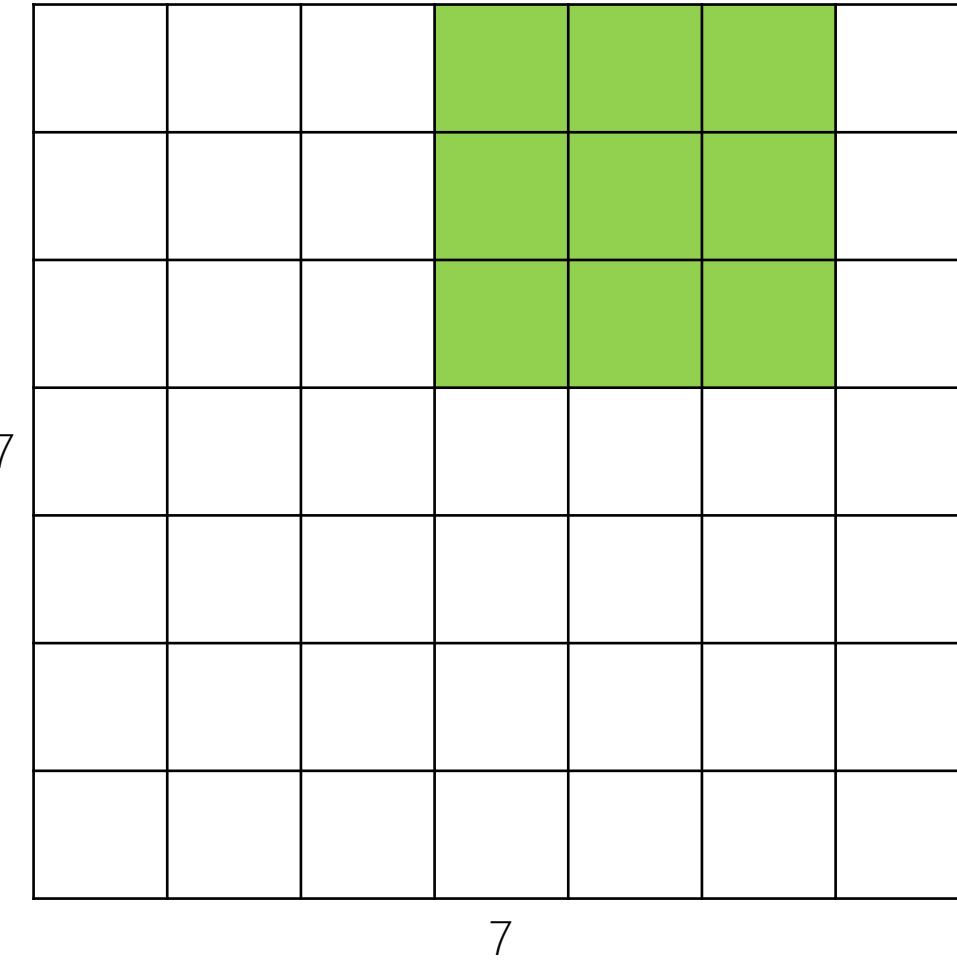
→ 3x3 output

Stride & Padding



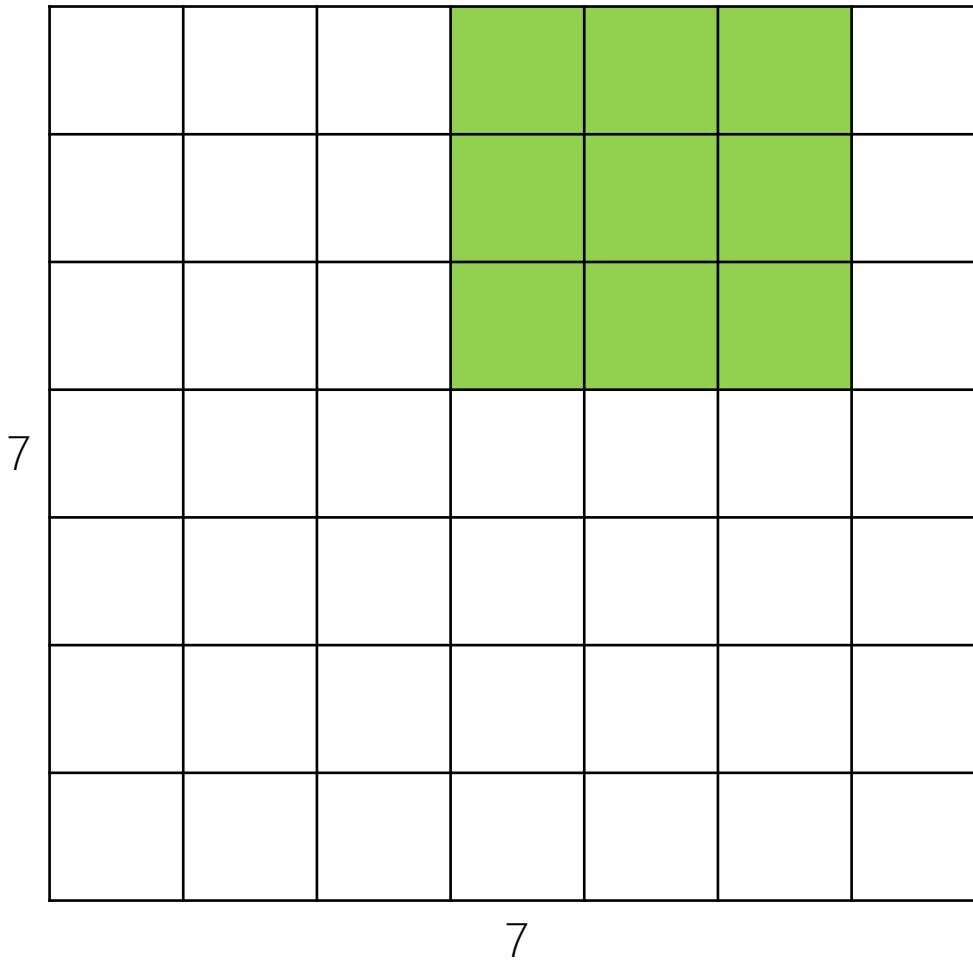
How about with **stride 3**?

Stride & Padding



How about with **stride 3**?

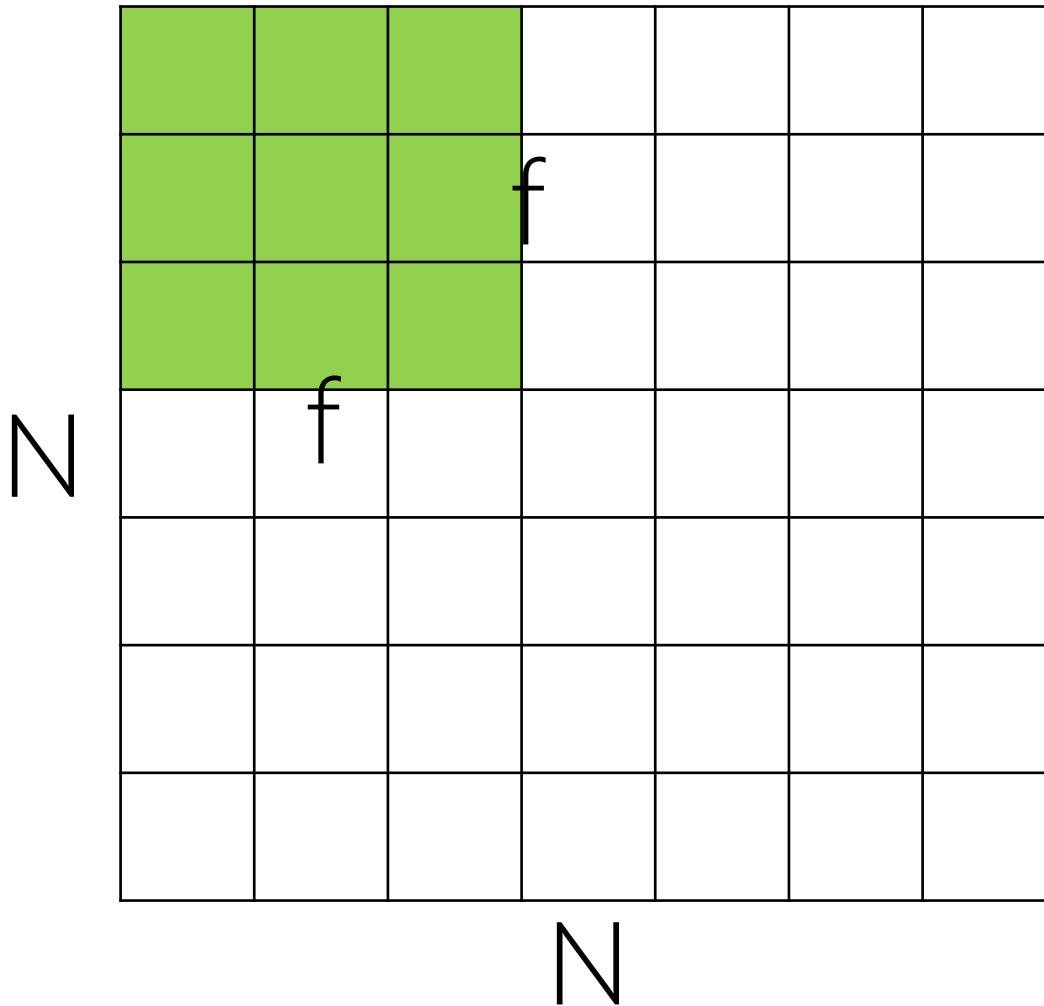
Stride & Padding



How about with **stride 3**?

→ ...?

Stride & Padding



Output size:
 $(N-f)/\text{stride} + 1$

Example, $N=7$, $f=3$:

$$\text{Stride 1: } (7-3)/1 + 1 = 5$$

$$\text{Stride 2: } (7-3)/2 + 1 = 3$$

$$\text{Stride 3: } (7-3)/3 + 1 = 2.3333$$

Stride & Padding

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

Output size:
 $(N-f)/\text{stride} + 1$

Q. 7x7 input, 3x3 kernel, with stride 1,
padded with 1 pixel. Output size?

Stride & Padding

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

Output size:
 $(N-f)/\text{stride} + 1$

Q. 7x7 input, 3x3 kernel, with stride 1,
padded with 1 pixel. Output size?

→ 7x7

Stride & Padding

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

Output size:
 $(N-f)/\text{stride} + 1$

Q. 7x7 input, 3x3 kernel, with stride 1,
padded with 1 pixel. Output size?

→ 7x7

Q. 7x7 input, 3x3 kernel, with stride 3.
You want to make 3x3 output, padding?

Stride & Padding

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

Output size:
 $(N-f)/\text{stride} + 1$

Q. 7x7 input, 3x3 kernel, with stride 1,
padded with 1 pixel. Output size?

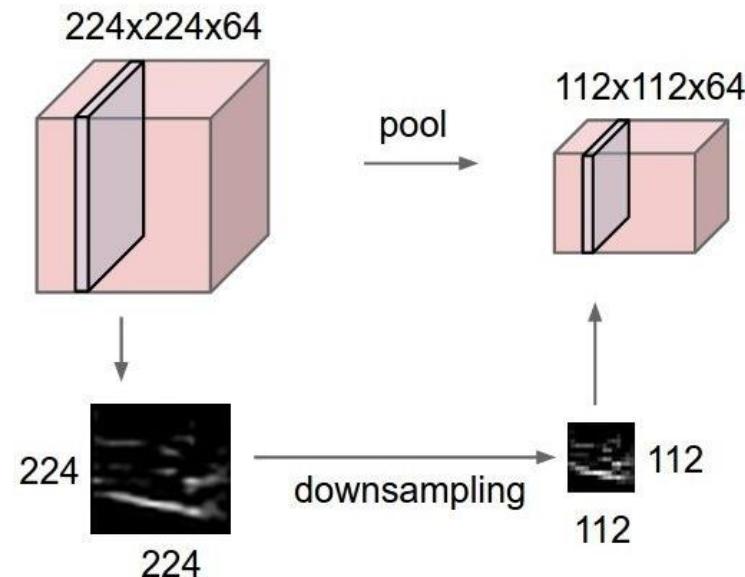
→ 7x7

Q. 7x7 input, 3x3 kernel, with stride 3.
You want to make 3x3 output, padding?

→ 1

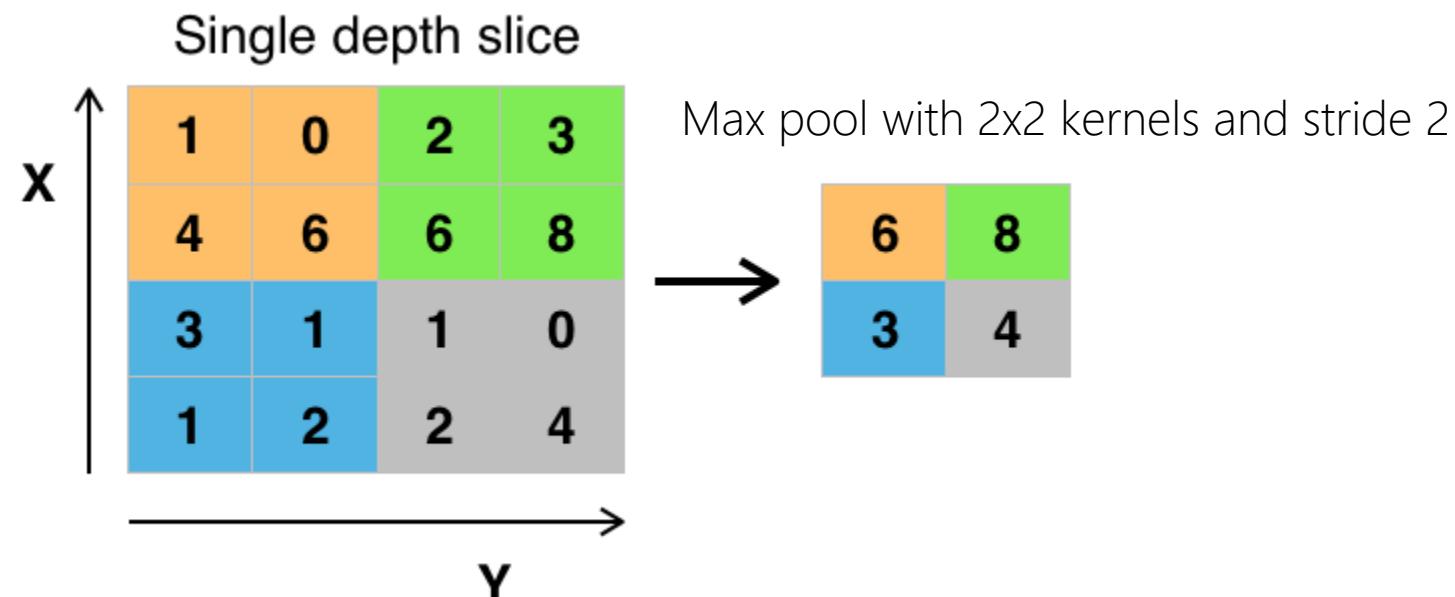
Pooling Layers

- Pooling == ConvNet way of downsampling
- “Reduce the image size” to obtain smaller but more manageable features
- Receptive fields becomes relatively larger as the image size gets smaller
- Operates over each activation map independently

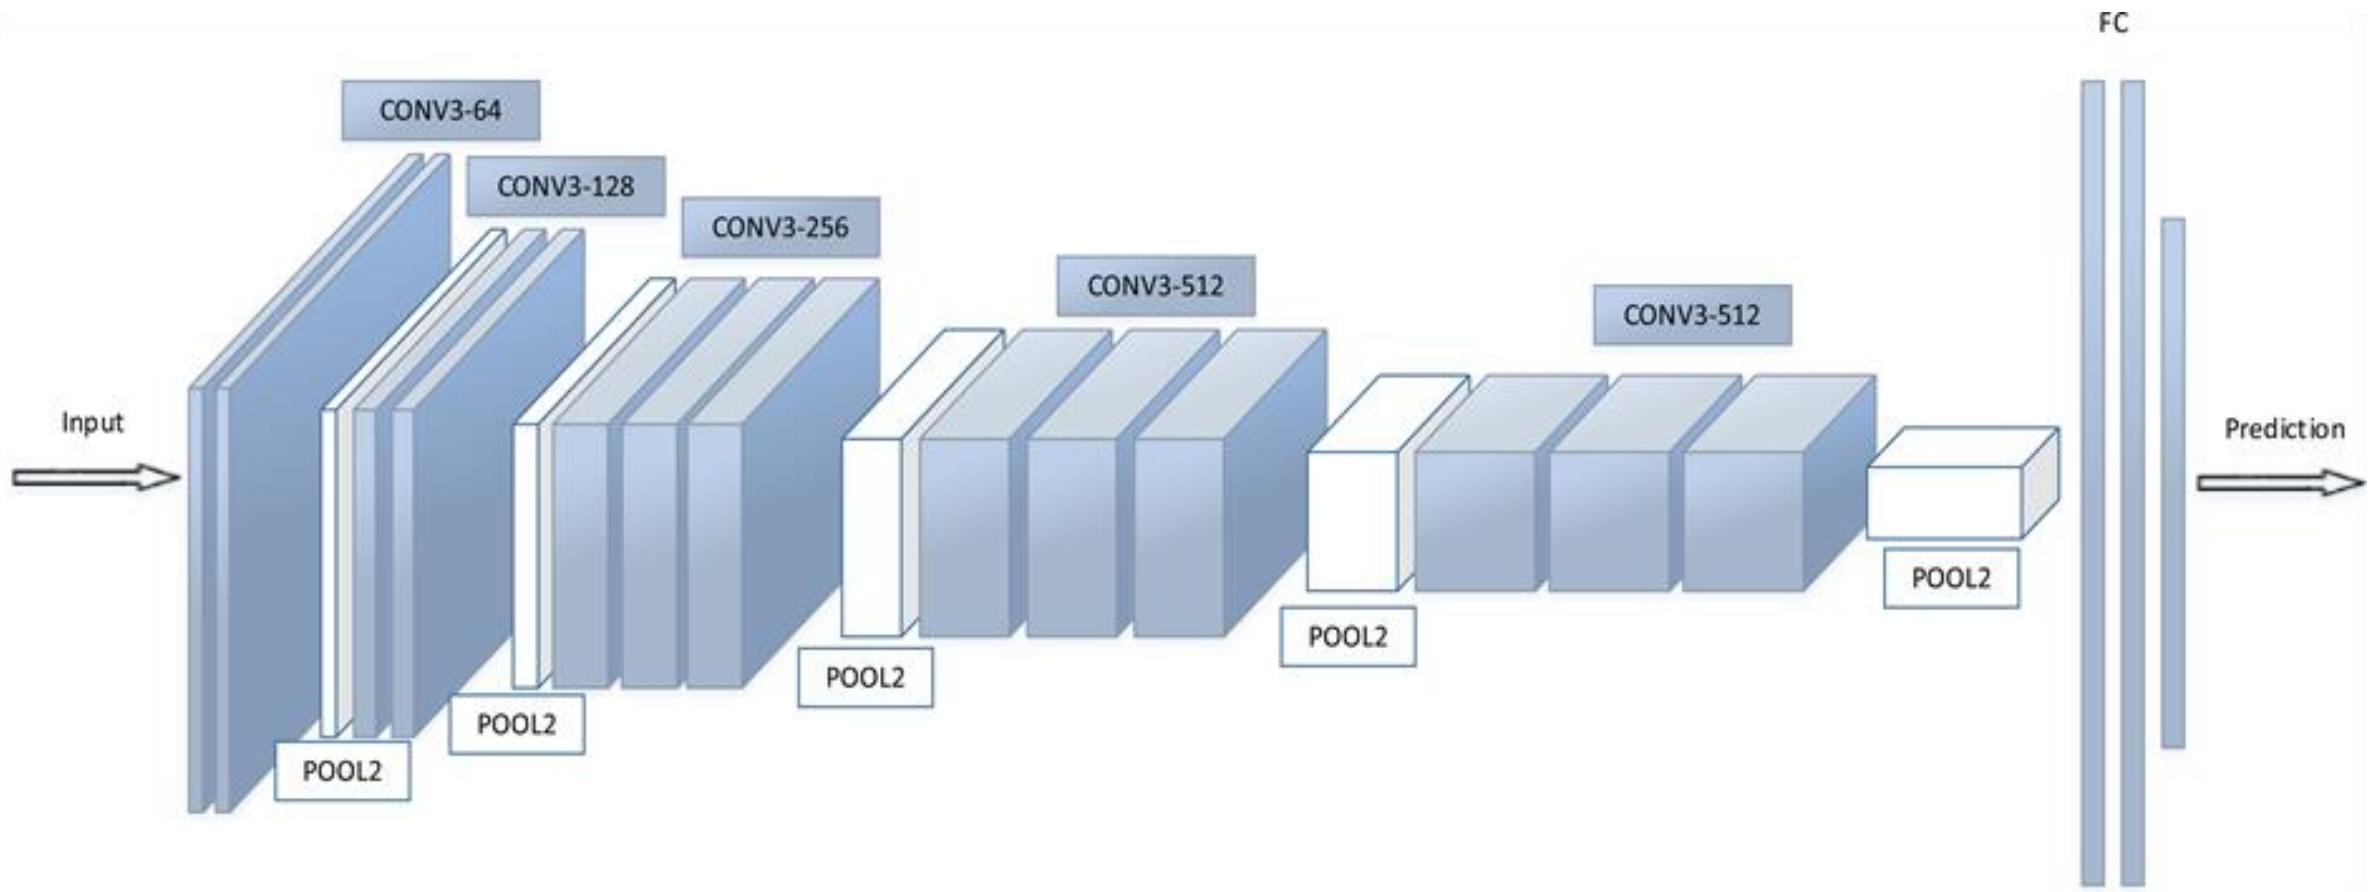


Max Pooling

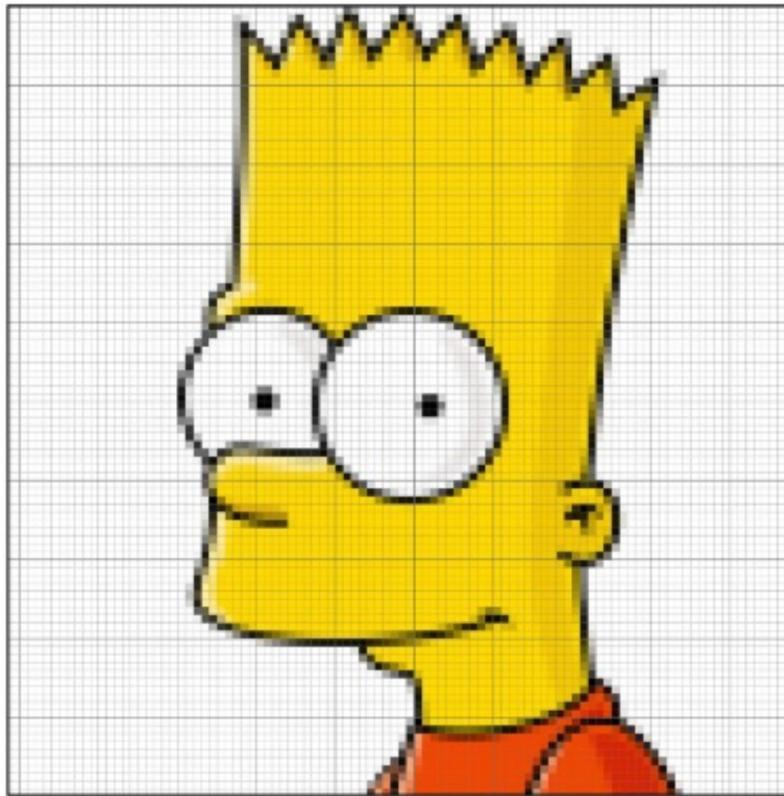
- A non-linear down-sampling method
- An image is partitioned into a set of (non-overlapping) rectangles.
- The maximum of each such sub-region is sampled.



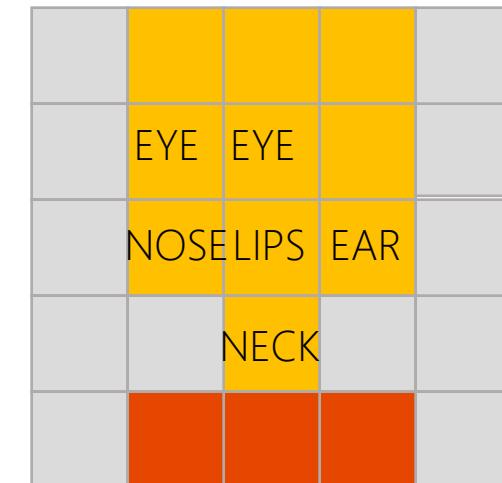
VGG Networks



Abstraction of an Image

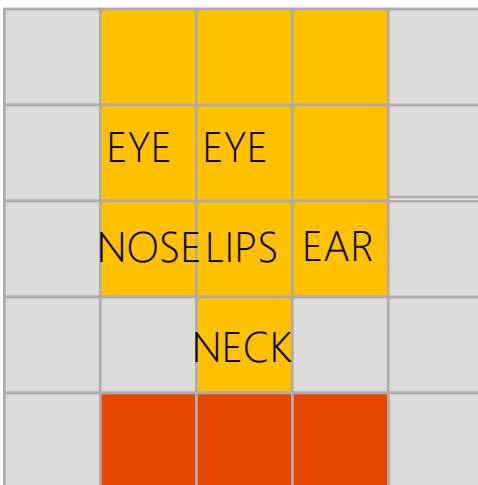


Conv2D + MaxPooling
+ Conv2D + MaxPooling
+ Conv2D + MaxPooling
+ ...



Abstraction of an Image

- Final decision is made by the fully-connected layers:



"Has two eyes, a nose, lips, an ear, and a neck
and wears t-shirts"



Cat	X
Dog	X
Person	O
Orange	X

ConvNet Architectures

Right off the shelf...



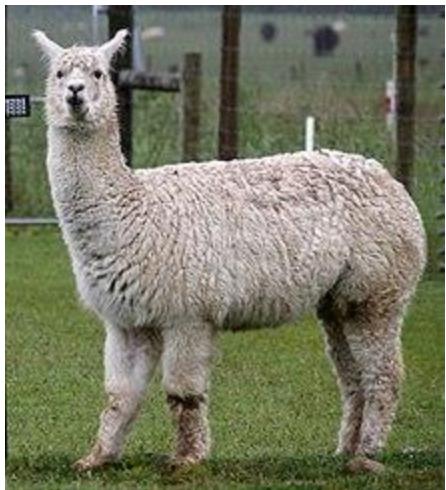
- AlexNet (2012)
- VGG (2014)
- GoogLeNet (or Inception) (2014)
- ResNet (2015)
- Inception-v4 (2016)
- DenseNet (2016)
- ResNext (2016)
- MobileNet (2017)
- NASNet (2018)
- ...

ImageNet Large Scale Visual Recognition Challenge (ILSVRC)

- ImageNet (First appearance as a poster at CVPR 2009)
 - Over 14 million images with hand-annotated labels (crowdsourced via M-turk)
 - Over 20 thousand categories
 - 1M+ images with bounding boxes.
- ImageNet Challenge
 - Since 2010
 - Uses a trimmed set of thousand non-overlapping classes.
 - Humans error is about 5%*
 - Playground of computer vision researchers.
 - Marked the start of the current AI boom.

ImageNet Large Scale Visual Recognition Challenge (ILSVRC)

- Criteria: Top-5 accuracy, Top-1 accuracy
- Say we had only four classes, namely "cat", "dog", "alpaca", and "camel". Now let's assume that the classifier returns something like this:



Your ML Model

[0.1, 0.5, 0.3, 0.1]



Top-1 class = {dog}

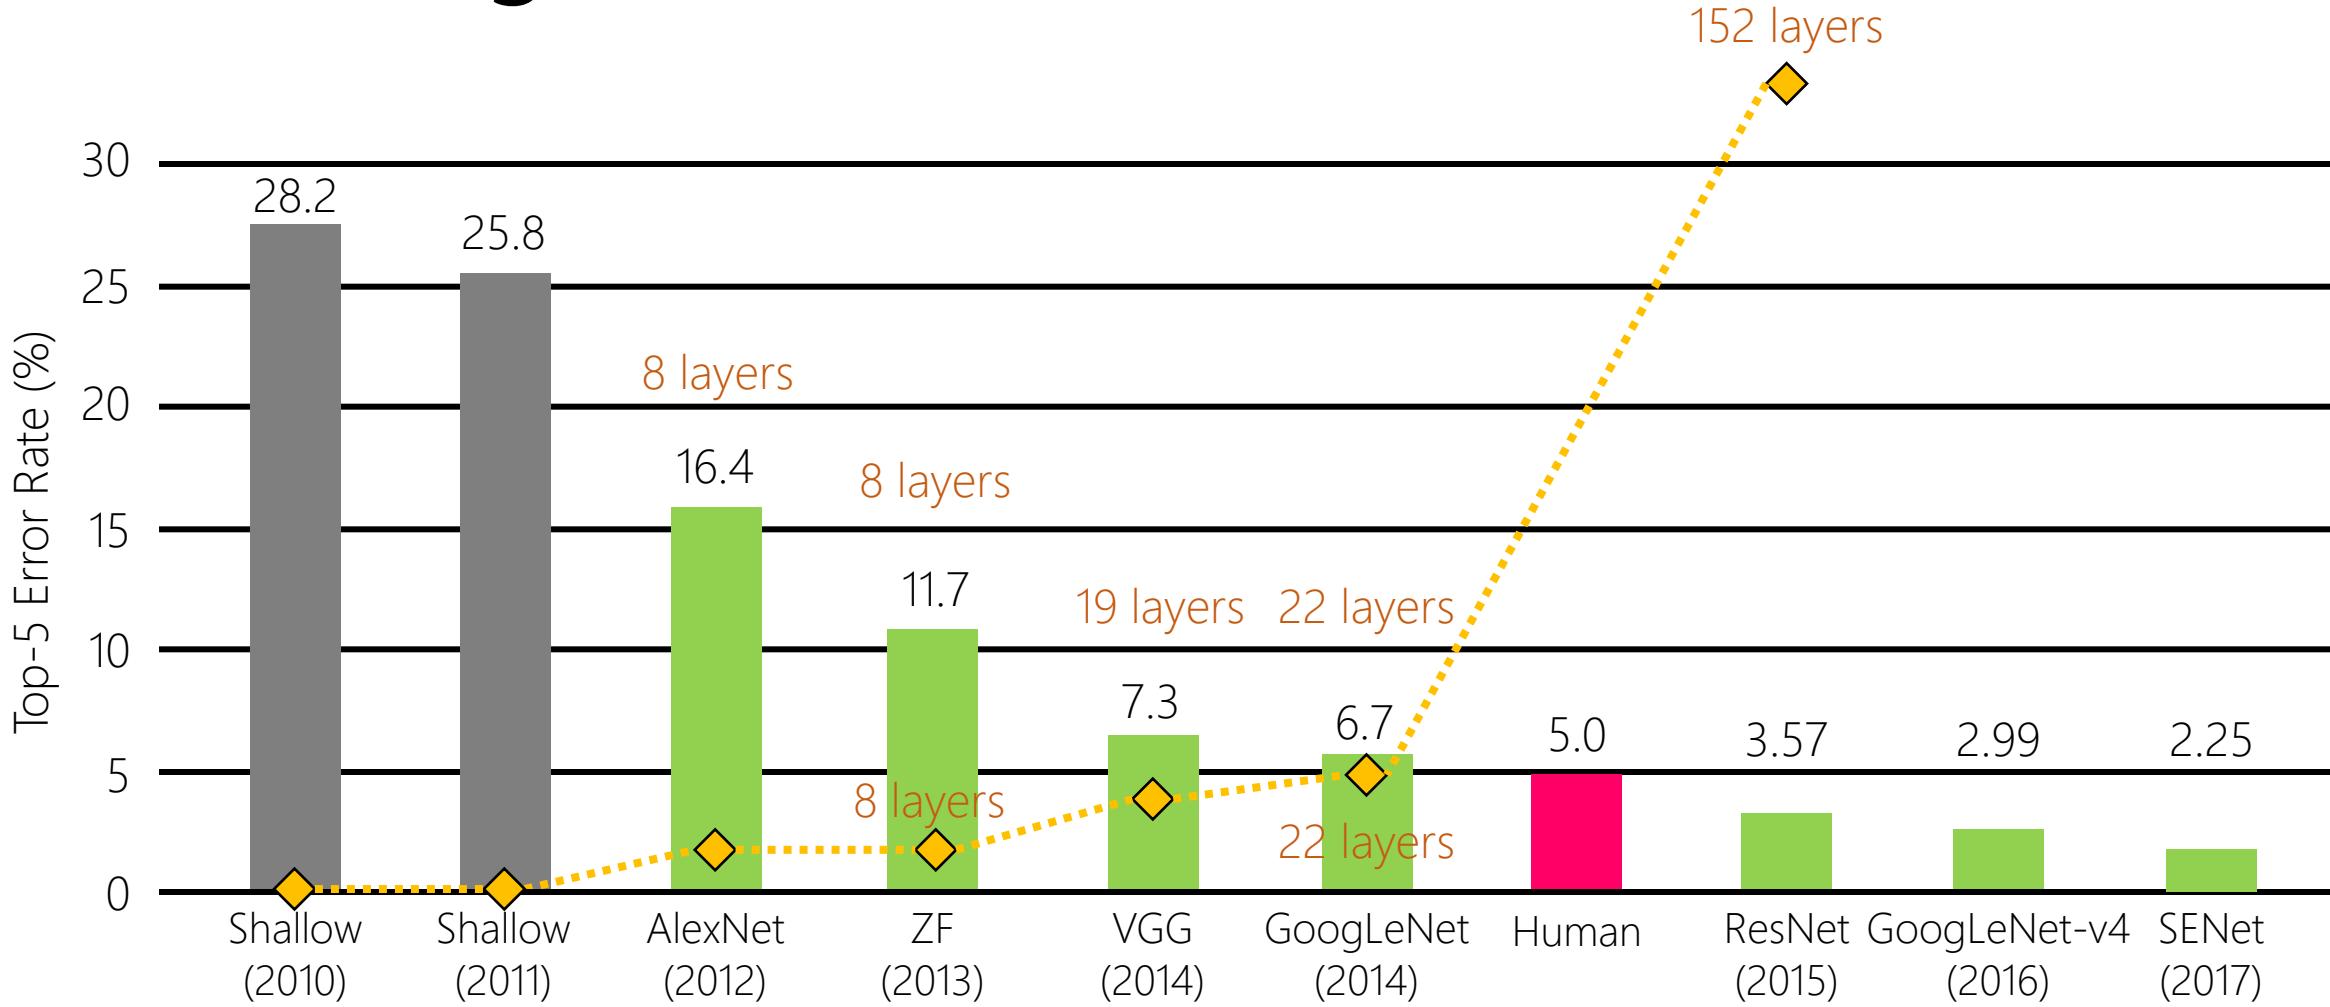
∴ Top-1 Accuracy = wrong

Top-2 class = {dog, alpaca}

∴ Top-2 Accuracy = correct

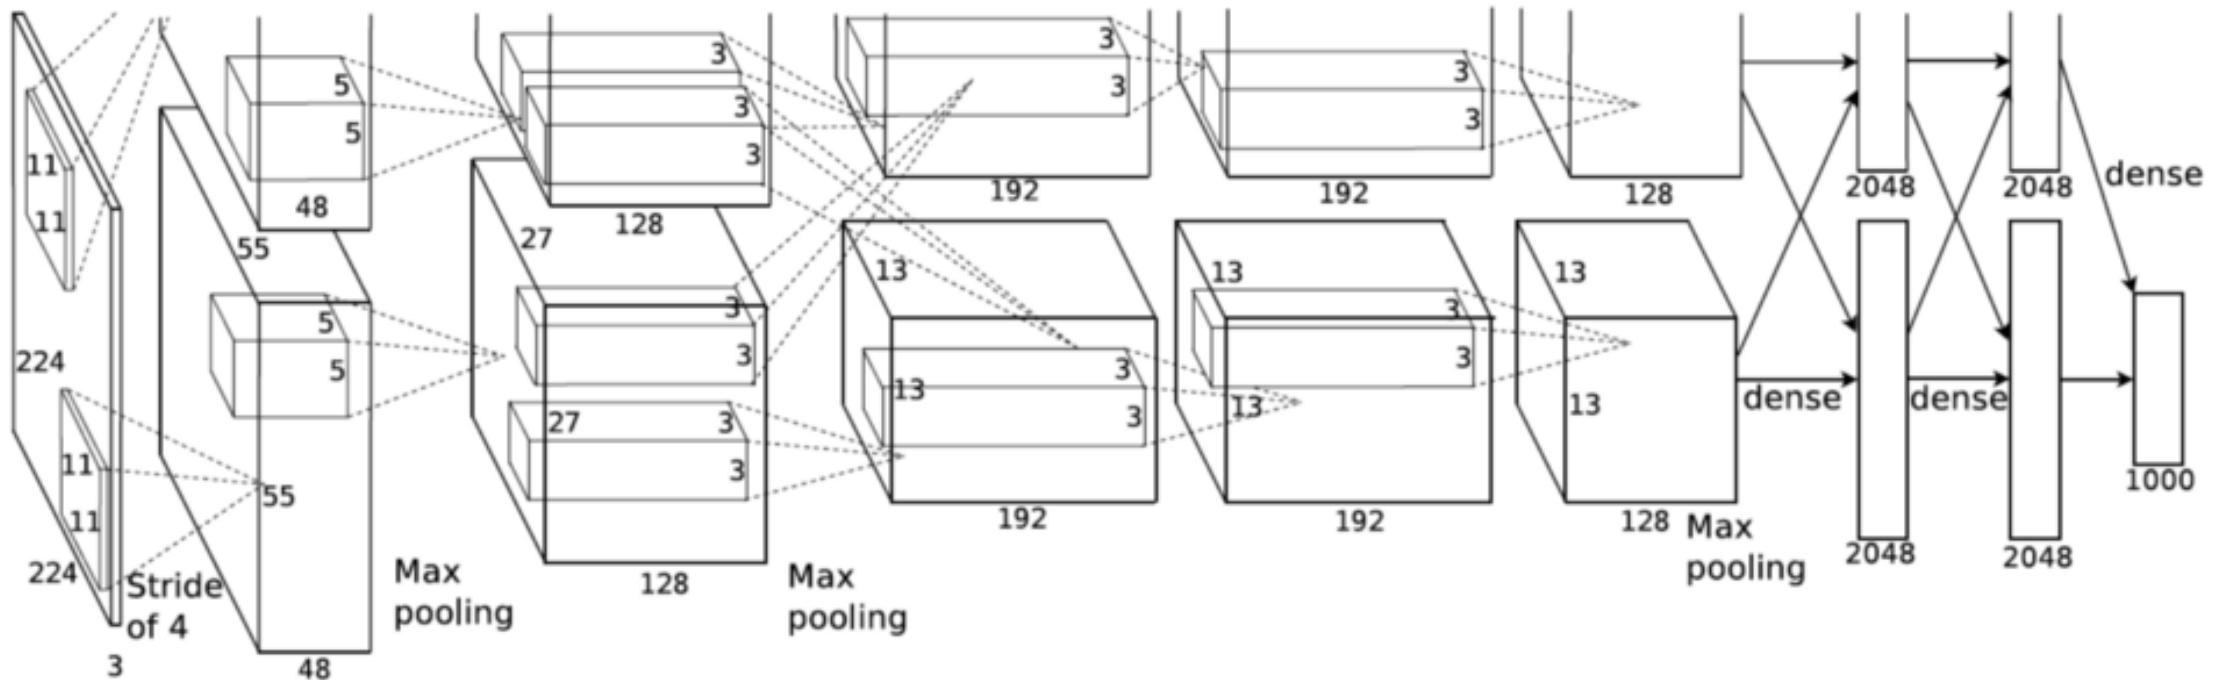
Q. Top- k accuracy for k number of classes?

ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners



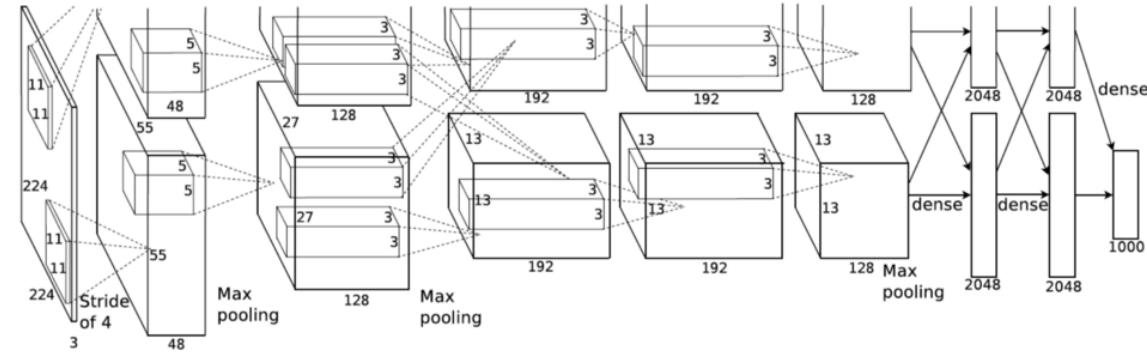
AlexNet

- Krizhevsky, Sutskever, & Hinton. (2012).



AlexNet

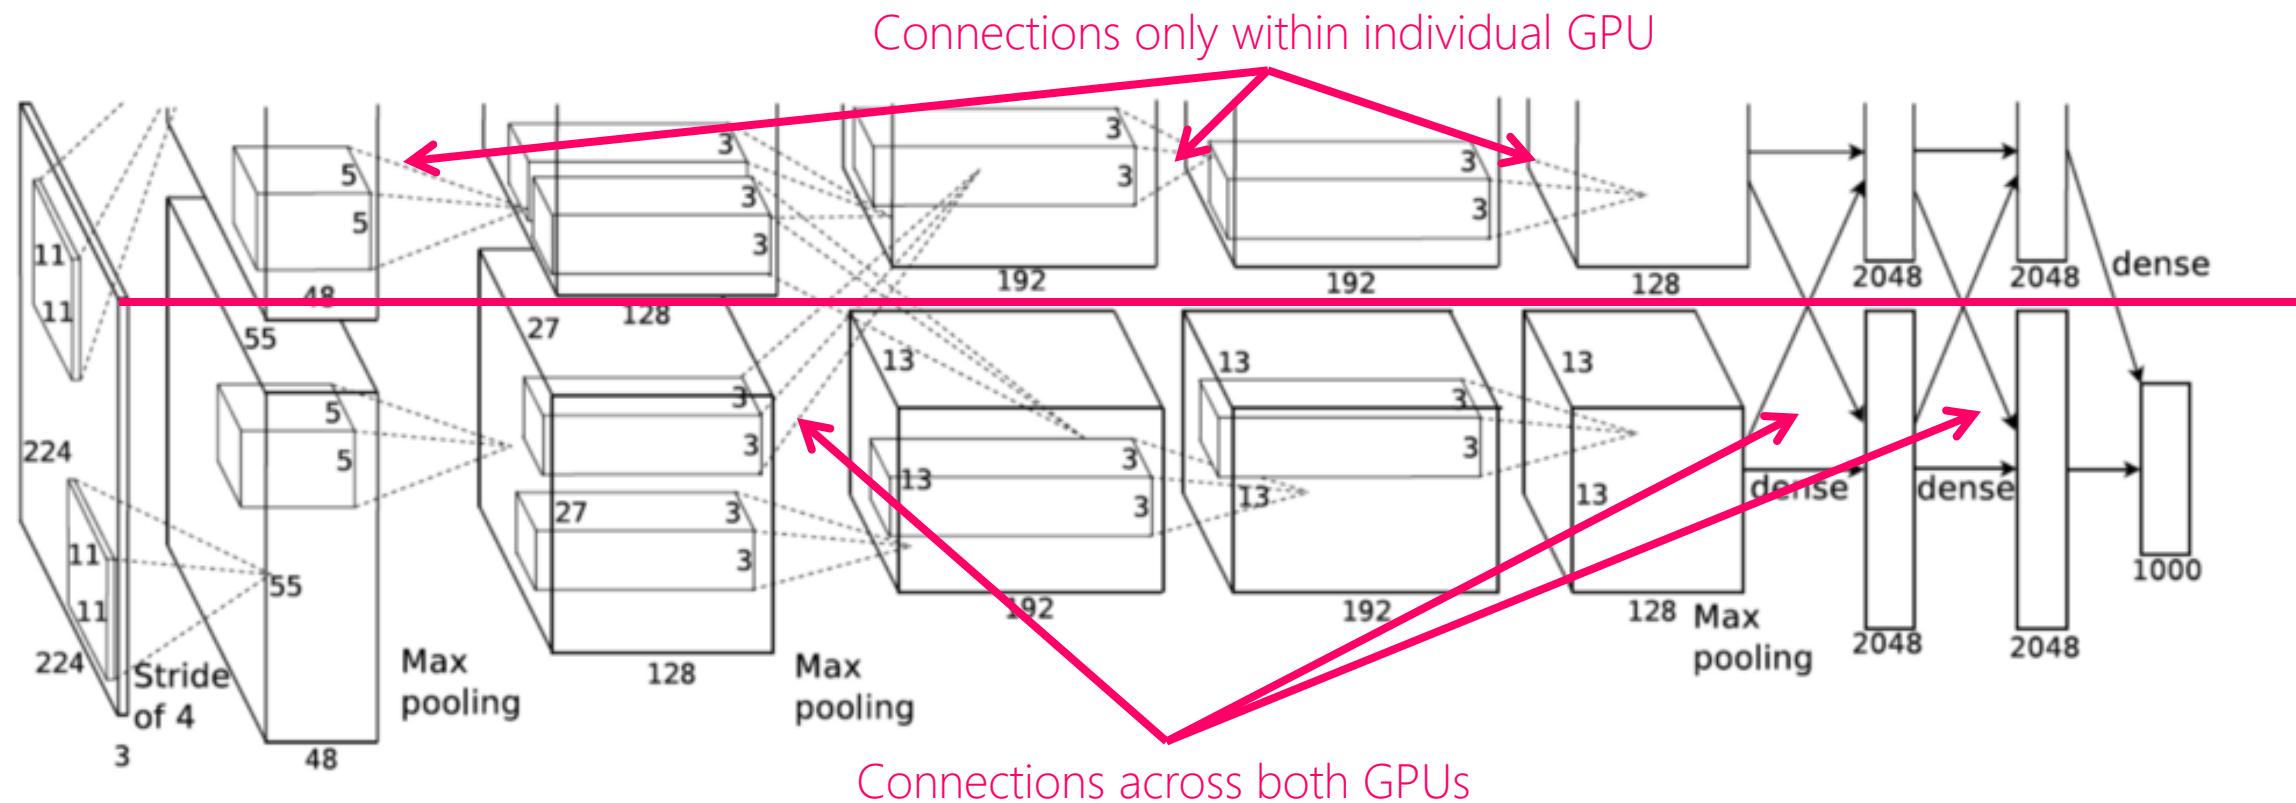
- Krizhevsky, Sutskever, & Hinton. (2012).
- 62.3M Parameters
- 1.1 billion computations in a forward pass
- Batch size 128
- SGD with Momentum 0.9
- Learning rate 0.01. The rate reduces by 10 manually when the validation accuracy plateaus
- Trained on GTX 580 GPU (3 GB memory)
→ 2x GPUs (5~6 days)
- 7 CNN ensemble (Accuracy 18.2% → 15.4%)



Conv1-96, stride 4
MaxPool3, stride 2
Norm
Conv5-256, stride 1
MaxPool3, stride 2
Norm
Conv3-384, stride 1
Conv3-384, stride 1
Conv3-256, stride 1
MaxPool3, stride 2
Dropout(0.5)
FC4096
Dropout(0.5)
FC4096
FC1000

AlexNet

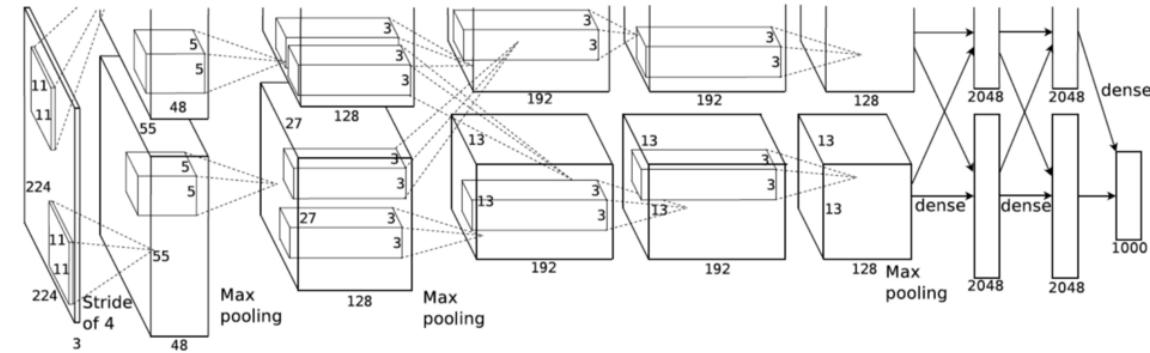
- Krizhevsky, Sutskever, & Hinton. (2012).



Dual GPU training → error rate decreases by around 1.7% (Top-1) and 1.2% (Top-5)
compared with the one trained with one GPU and half neurons.

AlexNet

- Krizhevsky, Sutskever, & Hinton. (2012).
- 96 kernels
 - top 48 kernels on GPU 1: color-agnostic
 - Bottom 48 kernels on GPU 2: color-specific



Conv1-96, stride 4
MaxPool3, stride 2
Norm
Conv5-256, stride 1
MaxPool3, stride 2
Norm
Conv3-384, stride 1
Conv3-384, stride 1
Conv3-256, stride 1
MaxPool3, stride 2
Dropout(0.5)
FC4096
Dropout(0.5)
FC4096
FC1000

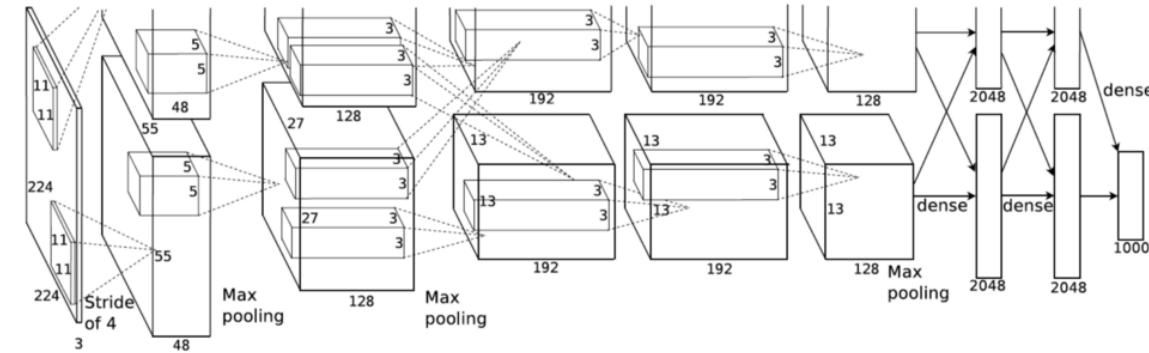
AlexNet

- Krizhevsky, Sutskever, & Hinton. (2012).
- Local Response Normalization

$$b_{x,y}^i = a_{x,y}^i / \left(k + \alpha \sum_{j=\max(0,i-n/2)}^{\min(N-1,i+n/2)} (a_{x,y}^j)^2 \right)^\beta$$

Response-normalized activity
 Activity of a neuron computed by applying kernel I at position (x,y) and then applying the ReLU nonlinearity

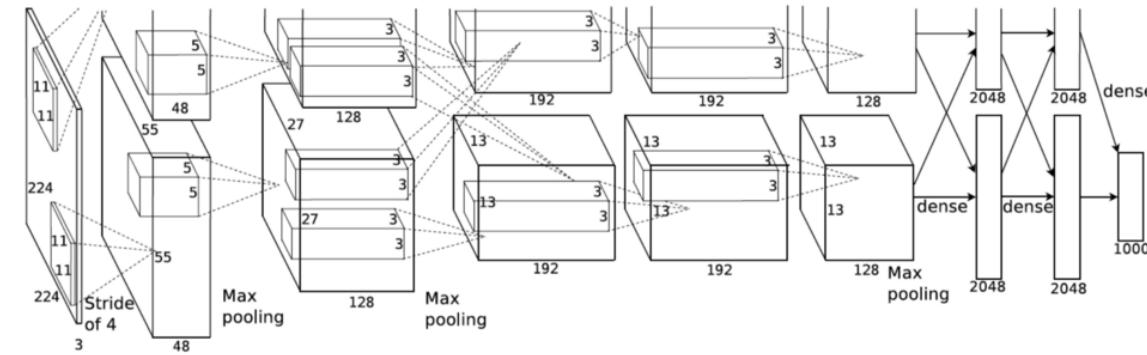
Response normalization reduces error rates by 1.4% (Top-1) & 1.2% (Top-5)



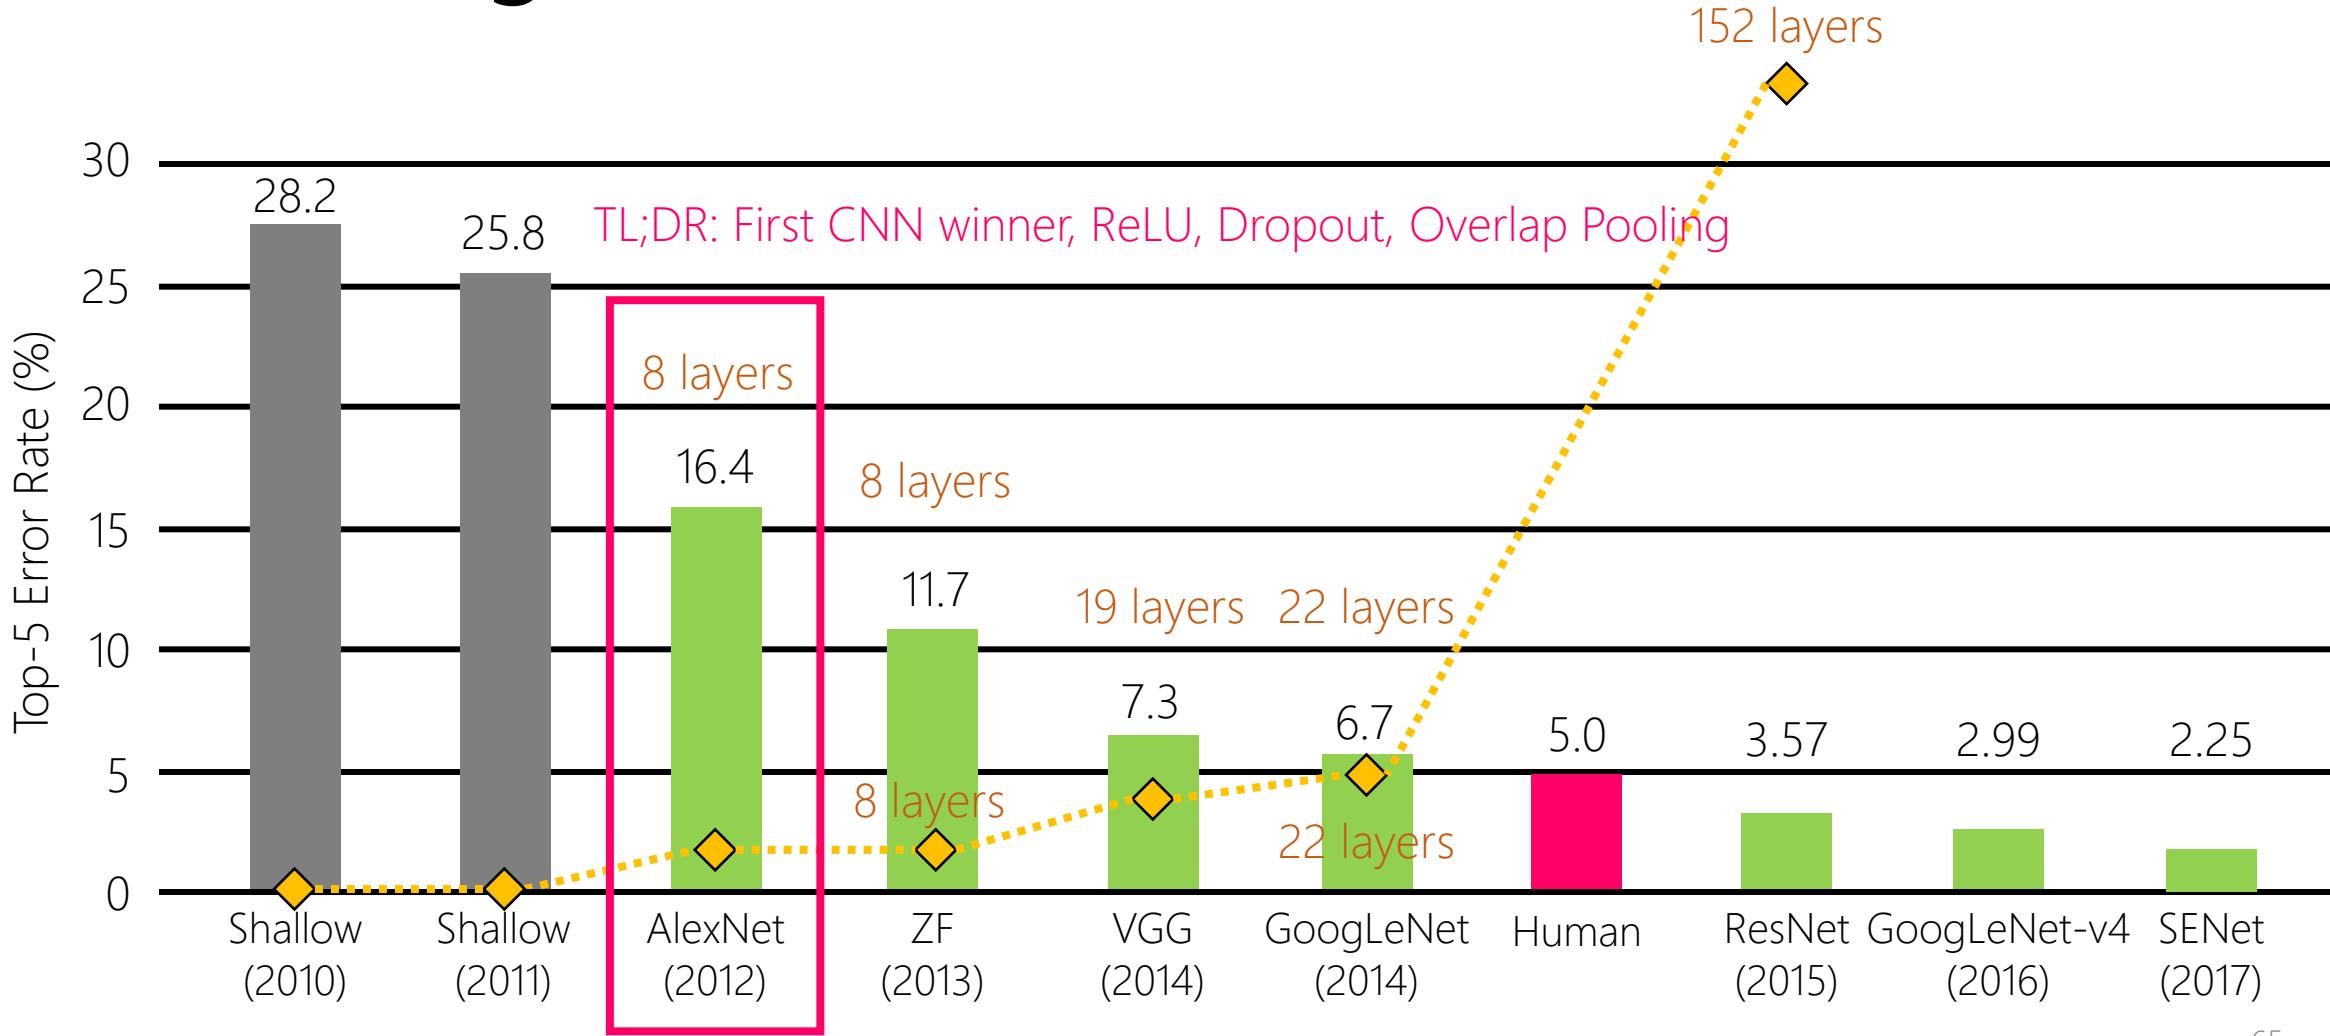
Conv1-96, stride 4
 MaxPool3, stride 2
 Norm
 Conv5-256, stride 1
 MaxPool3, stride 2
 Norm
 Conv3-384, stride 1
 Conv3-384, stride 1
 Conv3-256, stride 1
 MaxPool3, stride 2
 Dropout(0.5)
 FC4096
 Dropout(0.5)
 FC4096
 FC1000

AlexNet

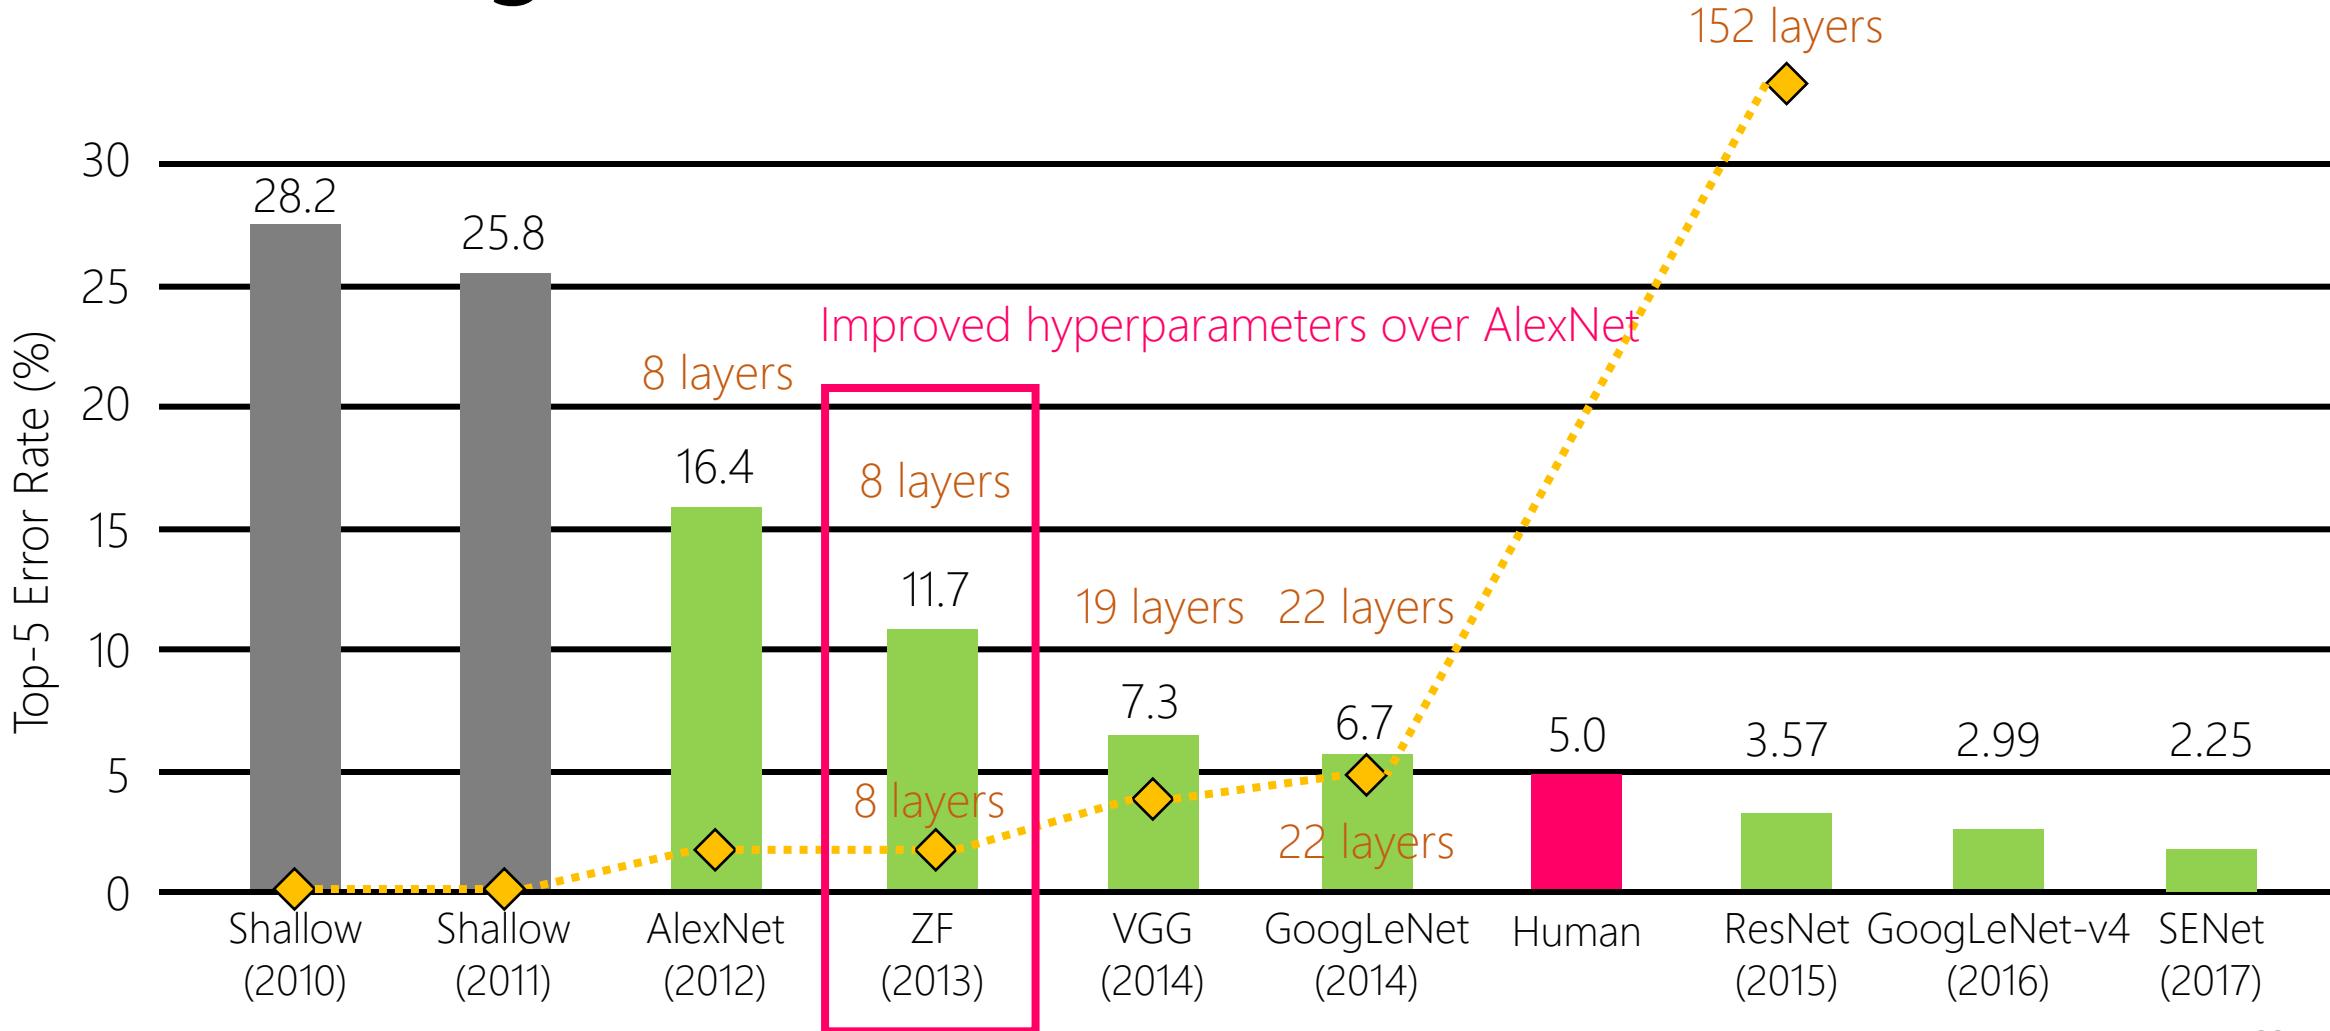
- Krizhevsky, Sutskever, & Hinton. (2012).
- Famously won the LSVRC 2012 by a large margin. (First CNN-based winner)
- First use of ReLU, instead of tanh or sigmoid.
- Dropout has been used to deal with overfitting.
- Data augmentation
- Overlap pooling → error rate decreases around 0.4% (Top-1) & 0.3% (Top-5)
- Used Norm layers (different from batch norm) → not popular anymore (turned out not very useful)



ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners

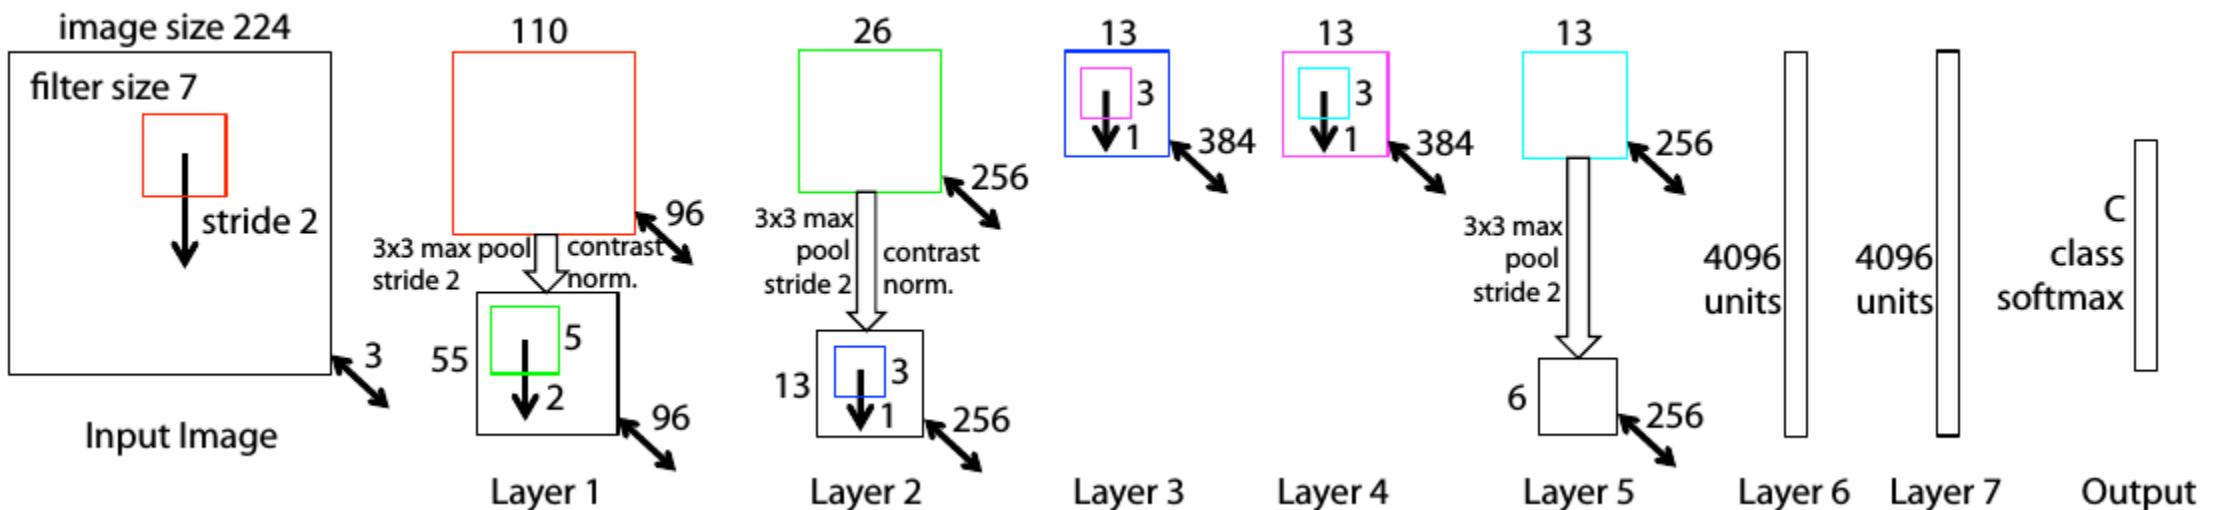


ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners

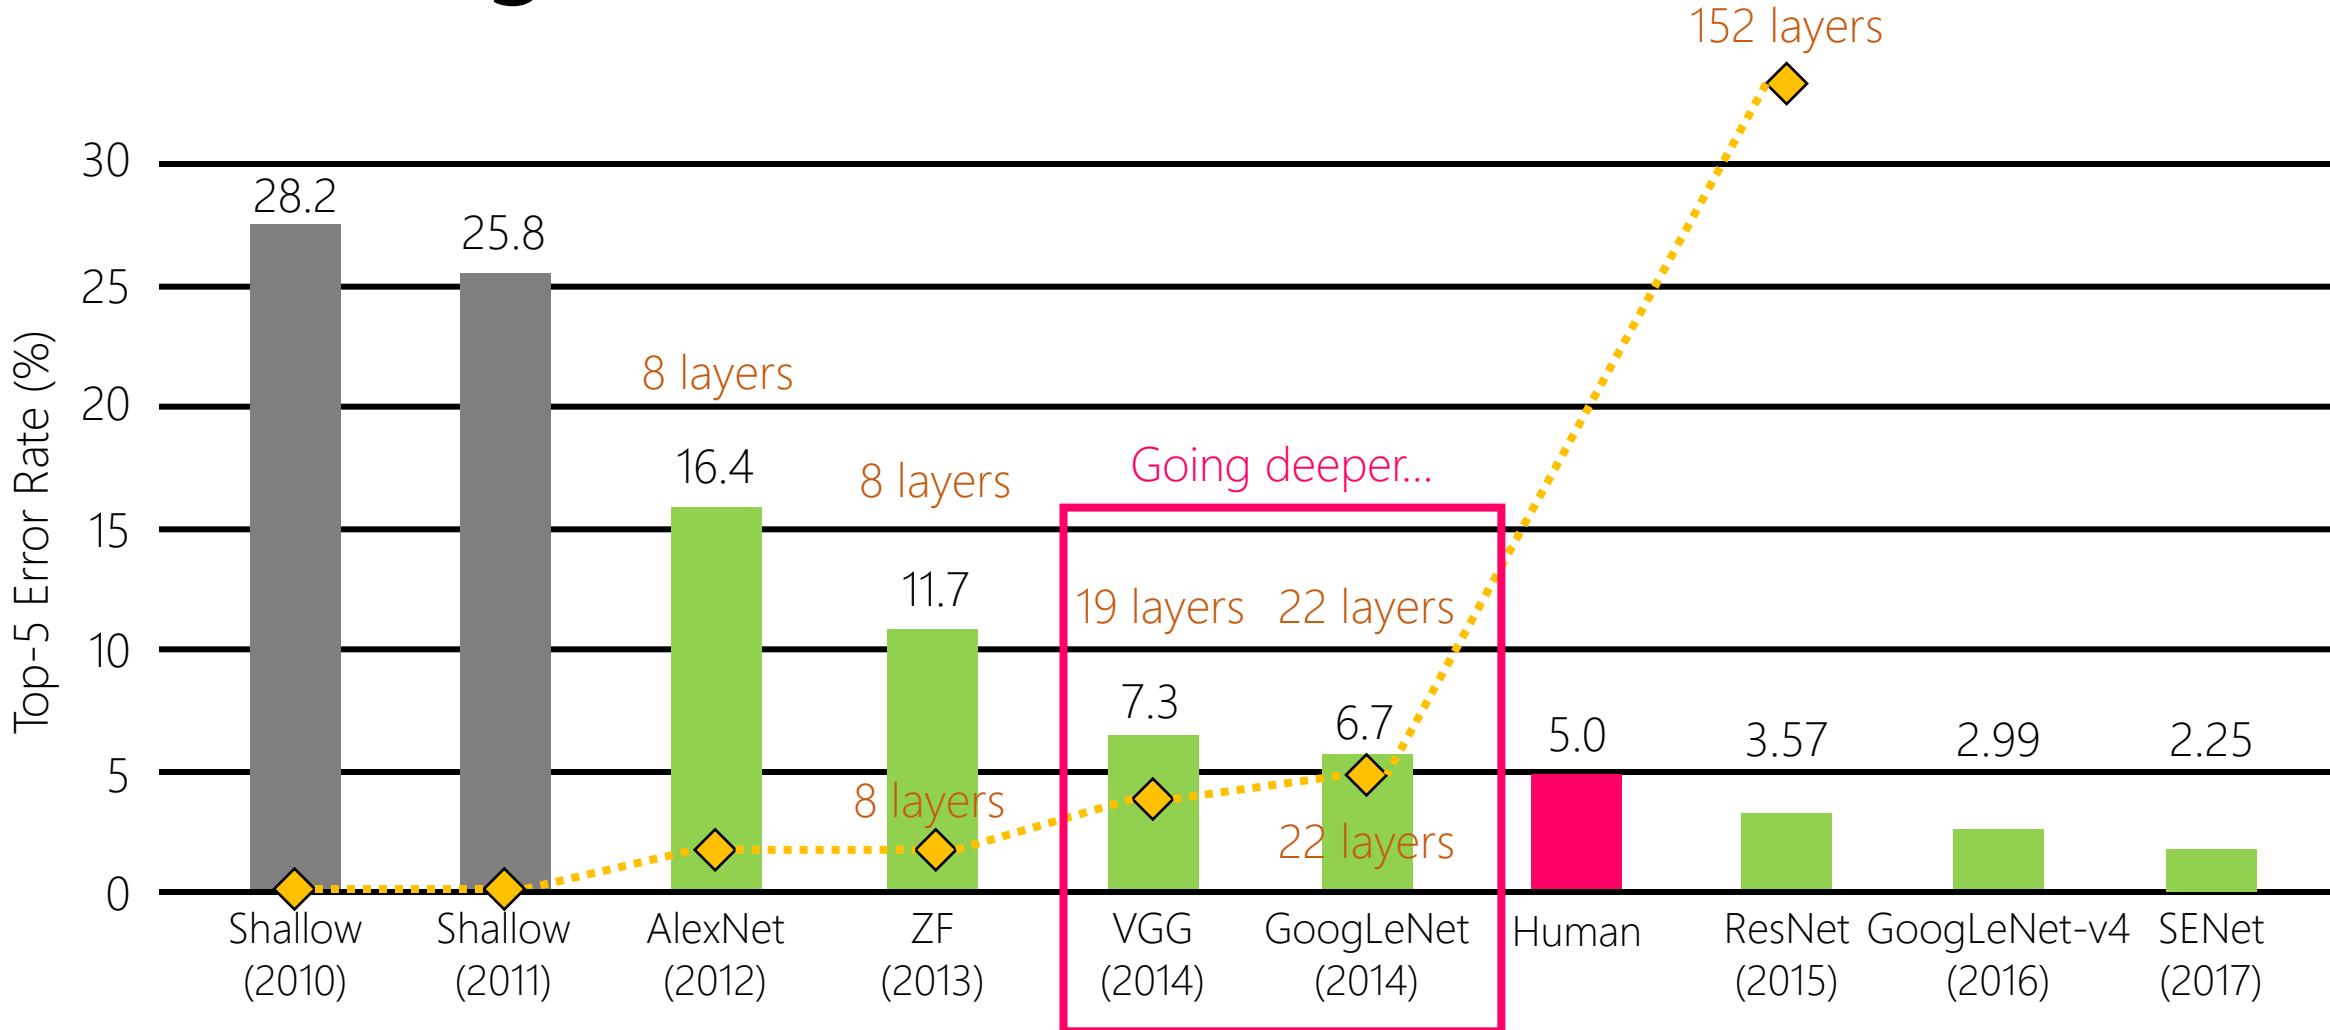


ZFNet

- Zeiler & Fergus (2013)
- Basically AlexNet, but
 - The first conv layer is changed from (11x11 stride 4) to (7x7 stride 2)
 - The third~fifth conv layer: instead of 384, 384, 256 kernels, ZFNet uses 512, 1024, 512 kernels.

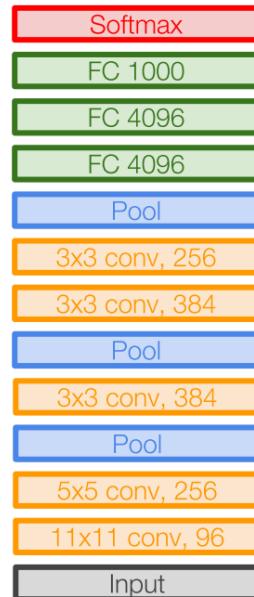


ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners

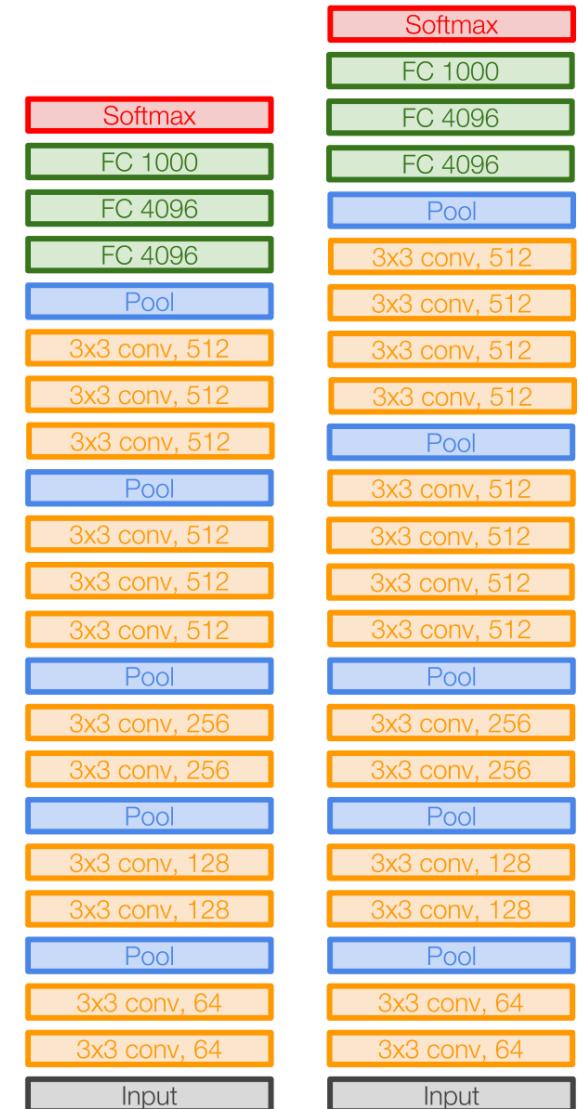


VGGNet

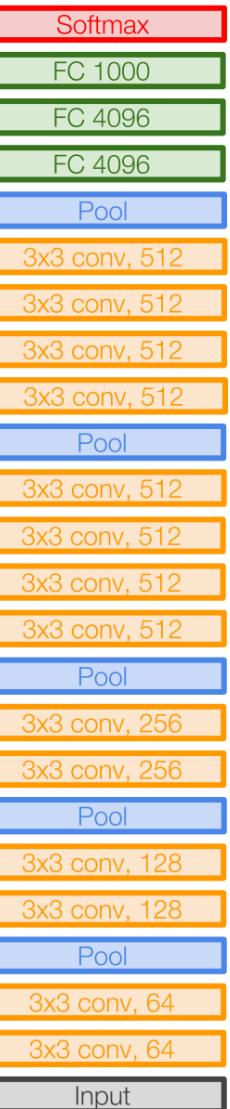
- Go deeper with smaller filters!
 - AlexNet (8 layers) → VGG16 & VGG19 (16 & 19 layers)
 - Only (Conv3-stride1-pad1) and (MaxPool2-stride2)
 - First time went below 10% error.



AlexNet



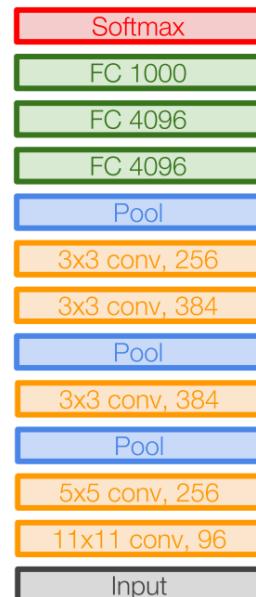
VGG16



VGG19

VGGNet

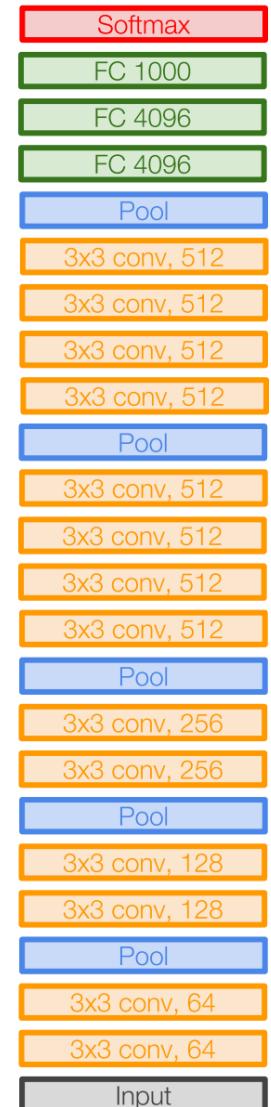
- Why use small filters?
 - Stack of three Conv3 layers covers 7x7 area in the input image (receptive field)
 - Same RF, but deeper, more nonlinearity
 - Obviously, fewer parameters! (7*7 vs 3*3*3) Softmax



AlexNet



VGG16



VGG19

VGGNet

INPUT: [224x224x3] memory: 224*224*3=150K params: 0 (not counting biases)

CONV3-64: [224x224x64] memory: 224*224*64=3.2M params: (3*3*3)*64 = 1,728

CONV3-64: [224x224x64] memory: 224*224*64=3.2M params: (3*3*64)*64 = 36,864

POOL2: [112x112x64] memory: 112*112*64=800K params: 0

CONV3-128: [112x112x128] memory: 112*112*128=1.6M params: (3*3*64)*128 = 73,728

CONV3-128: [112x112x128] memory: 112*112*128=1.6M params: (3*3*128)*128 = 147,456

POOL2: [56x56x128] memory: 56*56*128=400K params: 0

CONV3-256: [56x56x256] memory: 56*56*256=800K params: (3*3*128)*256 = 294,912

CONV3-256: [56x56x256] memory: 56*56*256=800K params: (3*3*256)*256 = 589,824

CONV3-256: [56x56x256] memory: 56*56*256=800K params: (3*3*256)*256 = 589,824

POOL2: [28x28x256] memory: 28*28*256=200K params: 0

CONV3-512: [28x28x512] memory: 28*28*512=400K params: (3*3*256)*512 = 1,179,648

CONV3-512: [28x28x512] memory: 28*28*512=400K params: (3*3*512)*512 = 2,359,296

CONV3-512: [28x28x512] memory: 28*28*512=400K params: (3*3*512)*512 = 2,359,296

POOL2: [14x14x512] memory: 14*14*512=100K params: 0

CONV3-512: [14x14x512] memory: 14*14*512=100K params: (3*3*512)*512 = 2,359,296

CONV3-512: [14x14x512] memory: 14*14*512=100K params: (3*3*512)*512 = 2,359,296

CONV3-512: [14x14x512] memory: 14*14*512=100K params: (3*3*512)*512 = 2,359,296

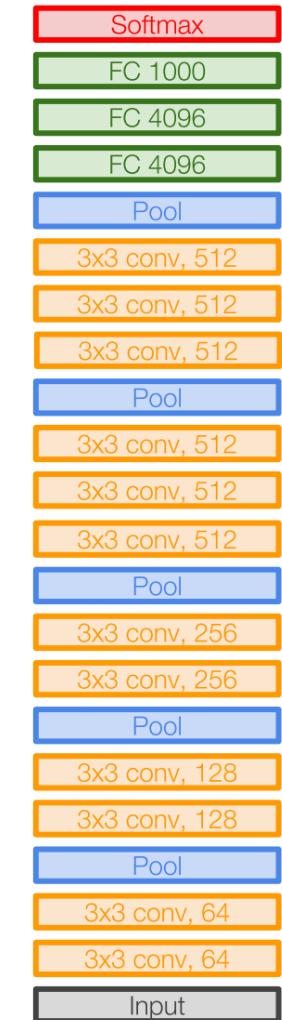
POOL2: [7x7x512] memory: 7*7*512=25K params: 0

FC: [1x1x4096] memory: 4096 params: 7*7*512*4096 = 102,760,448

FC: [1x1x4096] memory: 4096 params: 4096*4096 = 16,777,216

FC: [1x1x1000] memory: 1000 params: 4096*1000 = 4,096,000

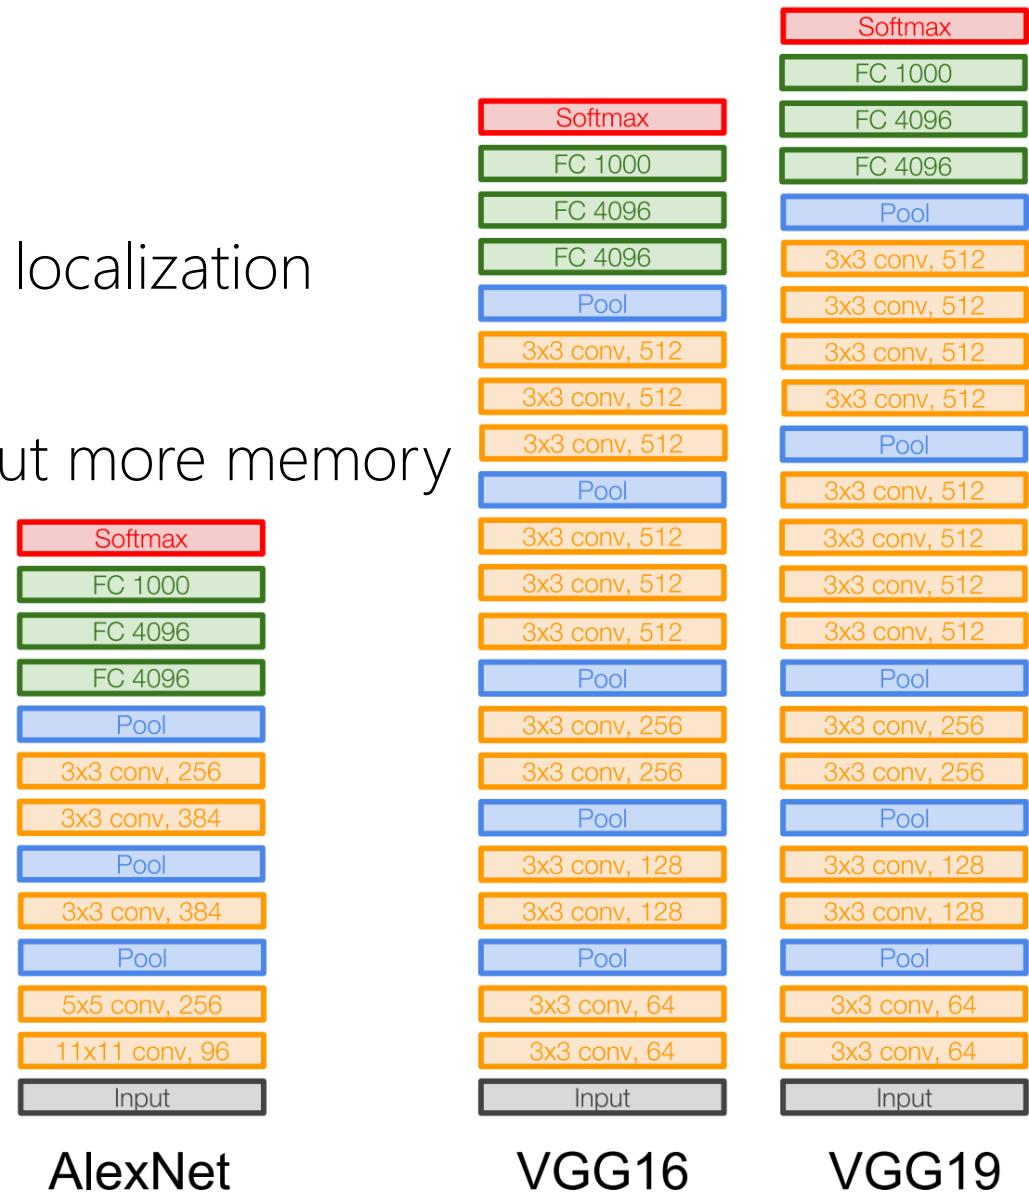
TOTAL memory: 24M * 4 bytes ~= 96MB / image (only forward! ~*2 for bwd)
 TOTAL params: 138M parameters



VGG16

VGGNet

- ILSVRC'14 runner-up in classification, 1st in localization
- No local response normalization
- VGG19 is only slightly better than VGG16 but more memory
- Simple and straightforward, the second FC layer generalizes quite well when transferred → “garlic salt”

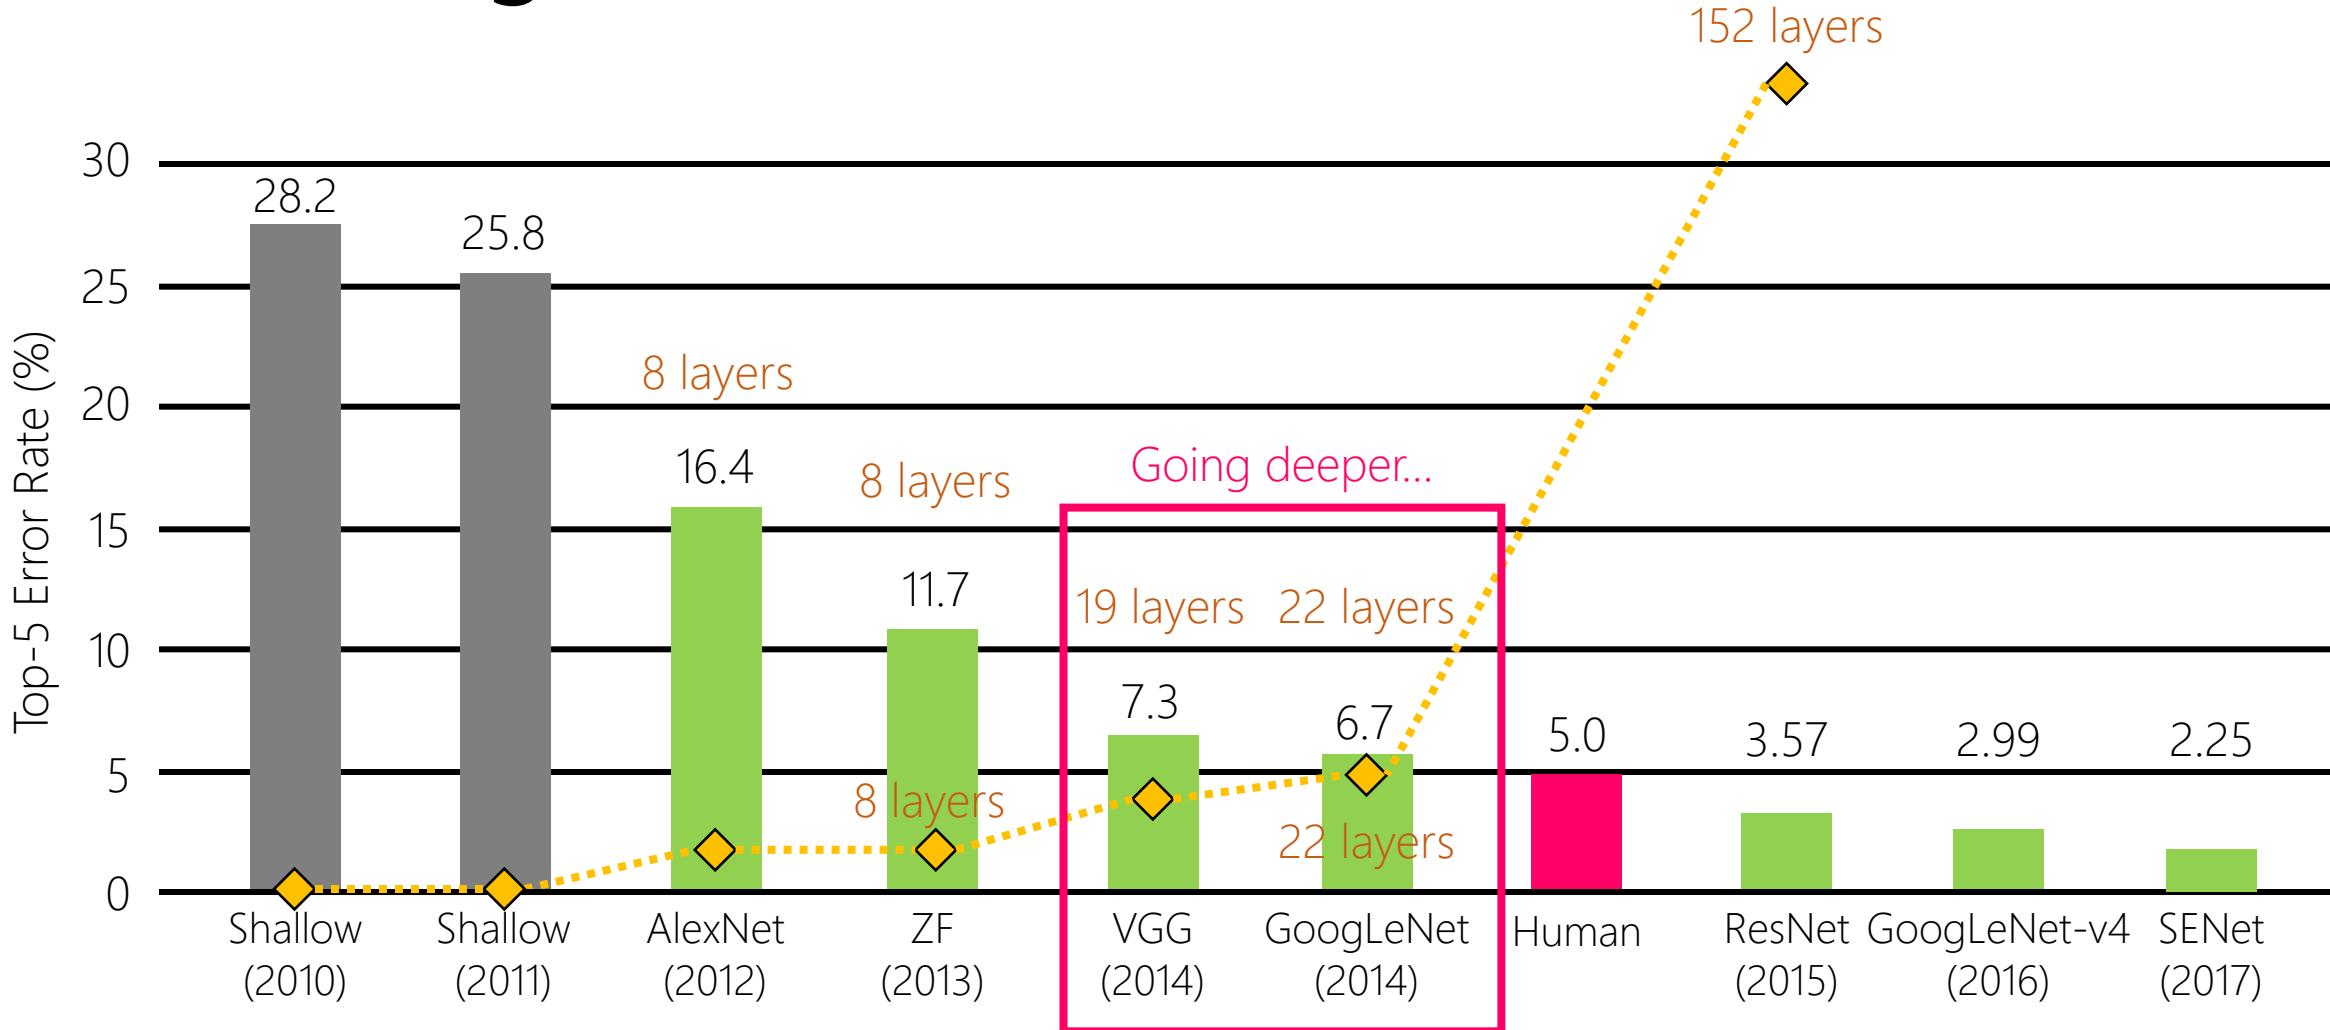


AlexNet

VGG16

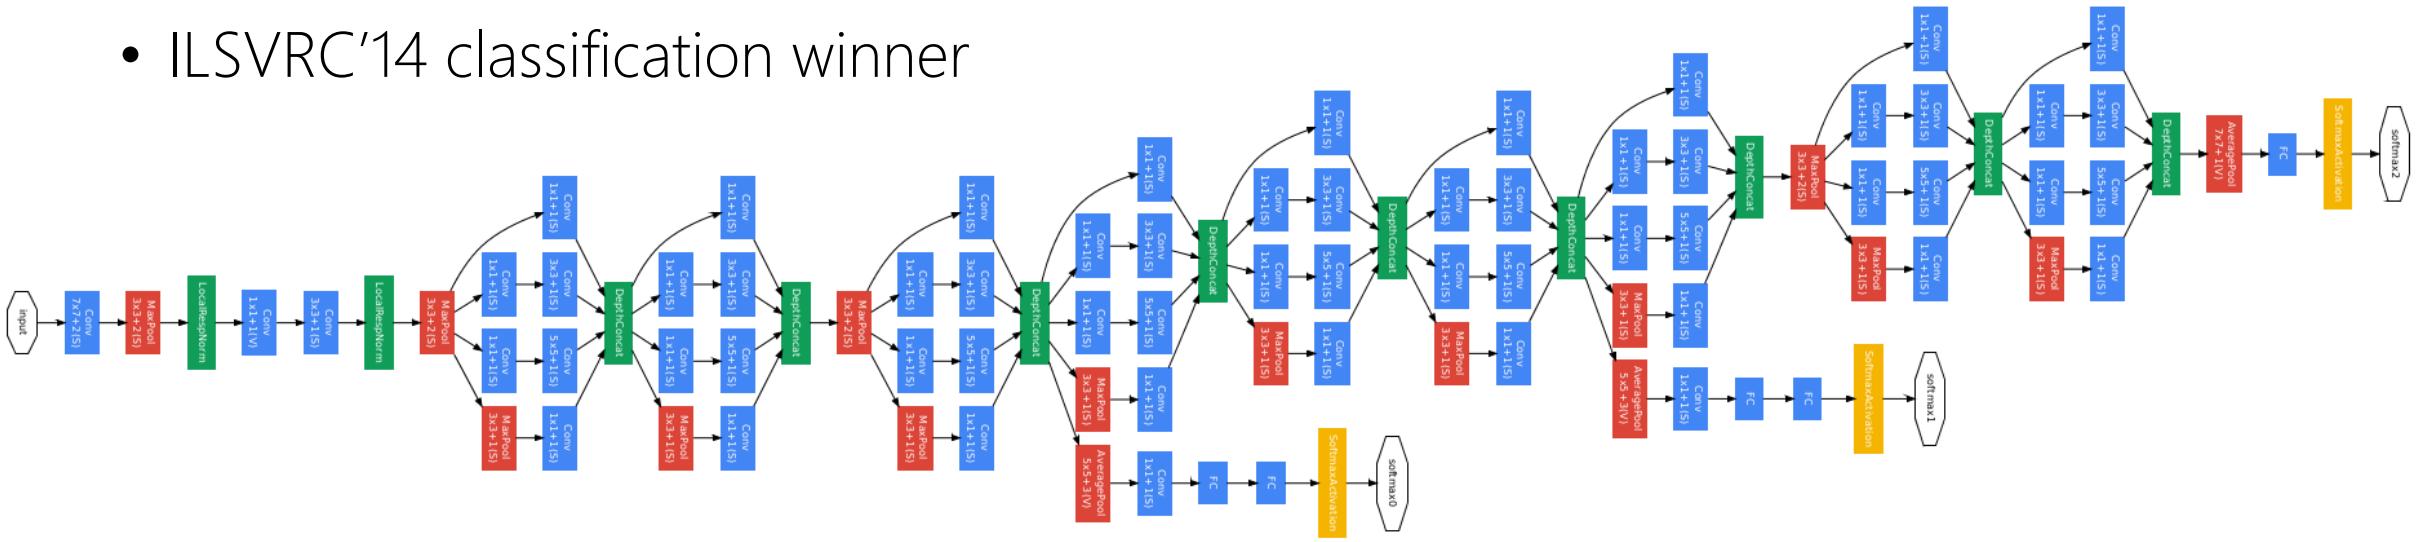
VGG19

ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners



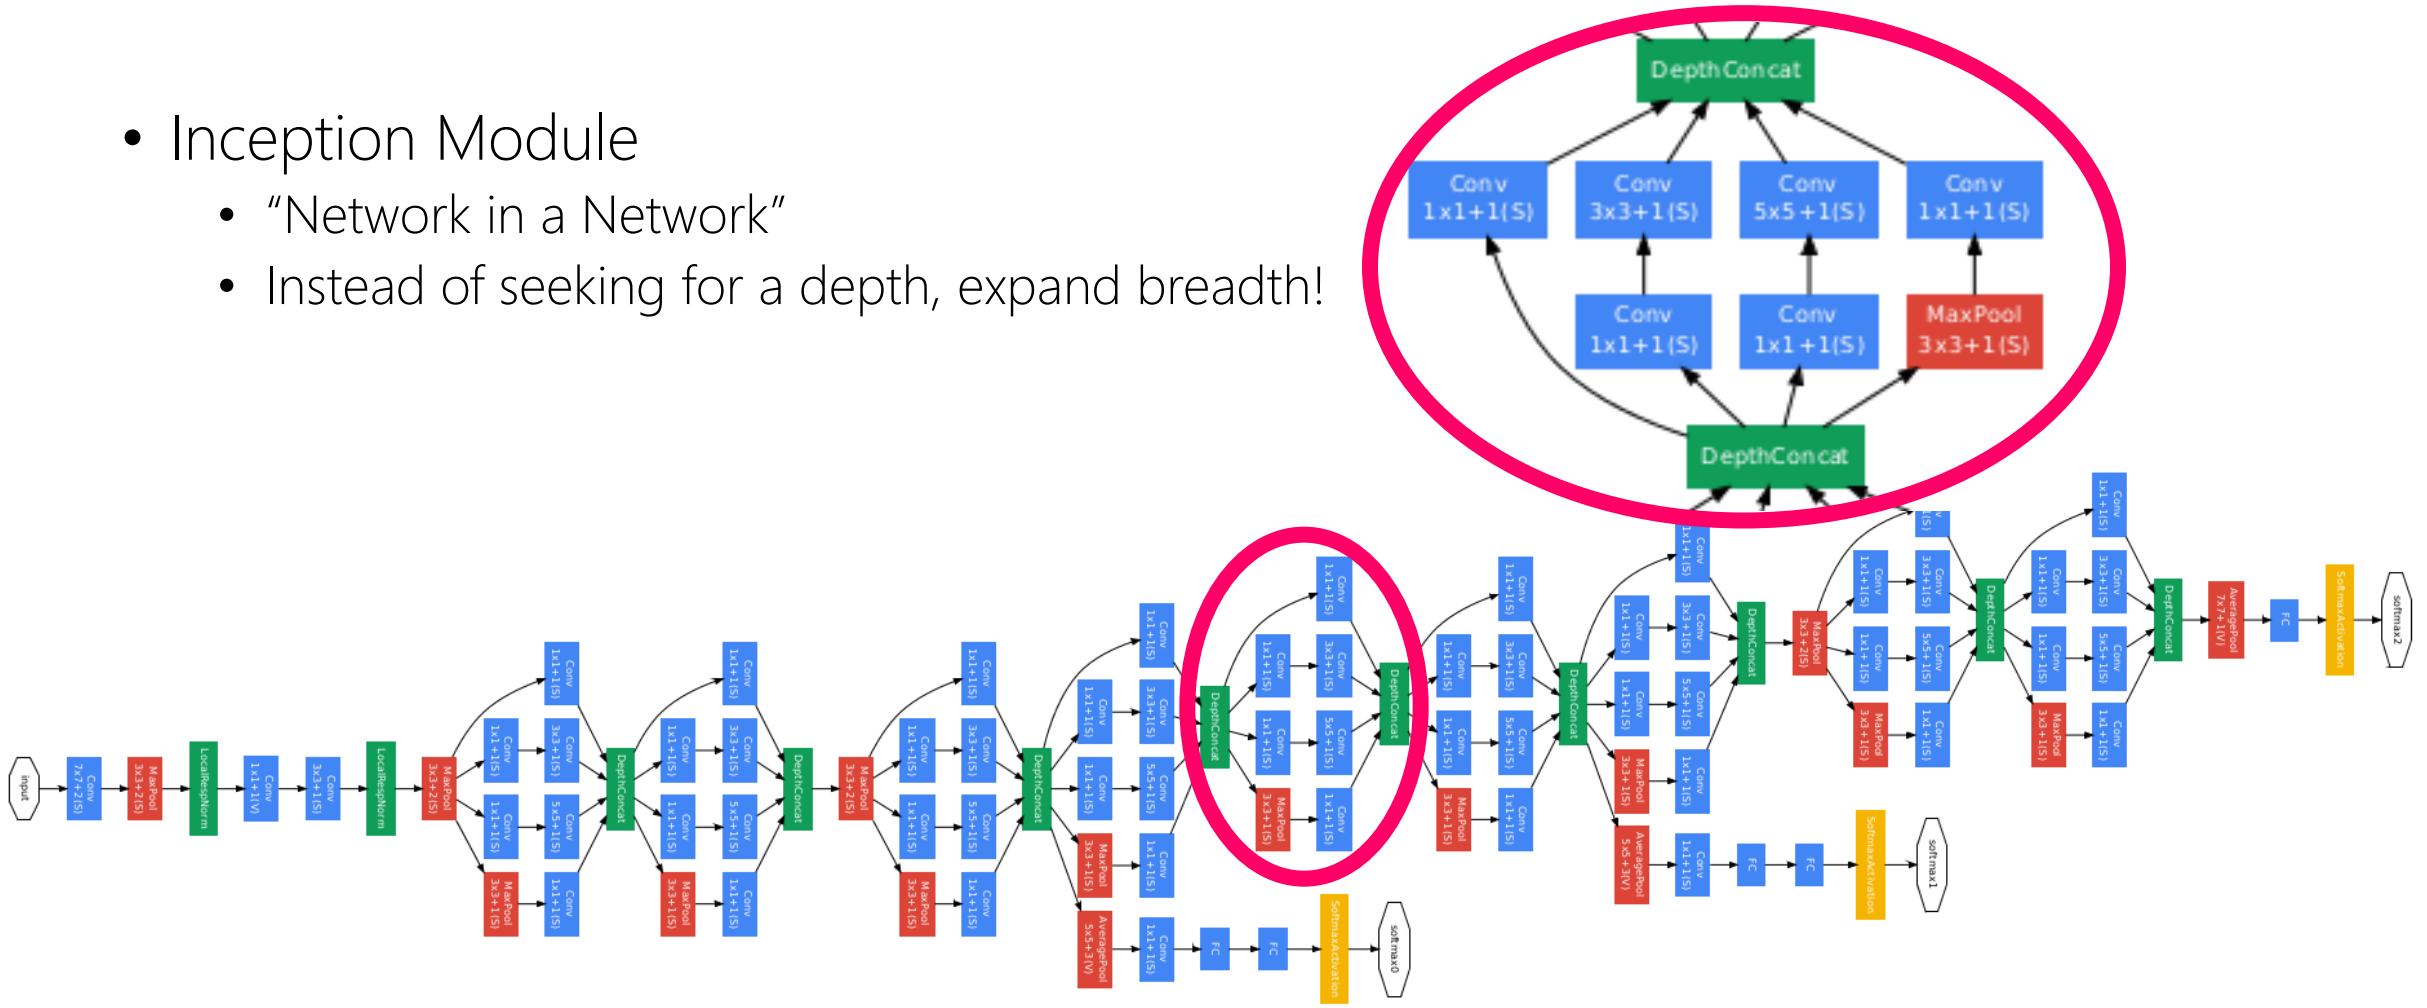
GoogLeNet a.k.a. Inception Network

- Deeper but efficient computation!
 - Efficiency via inception module
 - Reduced FC layers
 - Only 5M, as opposed to 62M (AlexNet) & 138M (VGG16)
 - ILSVRC'14 classification winner



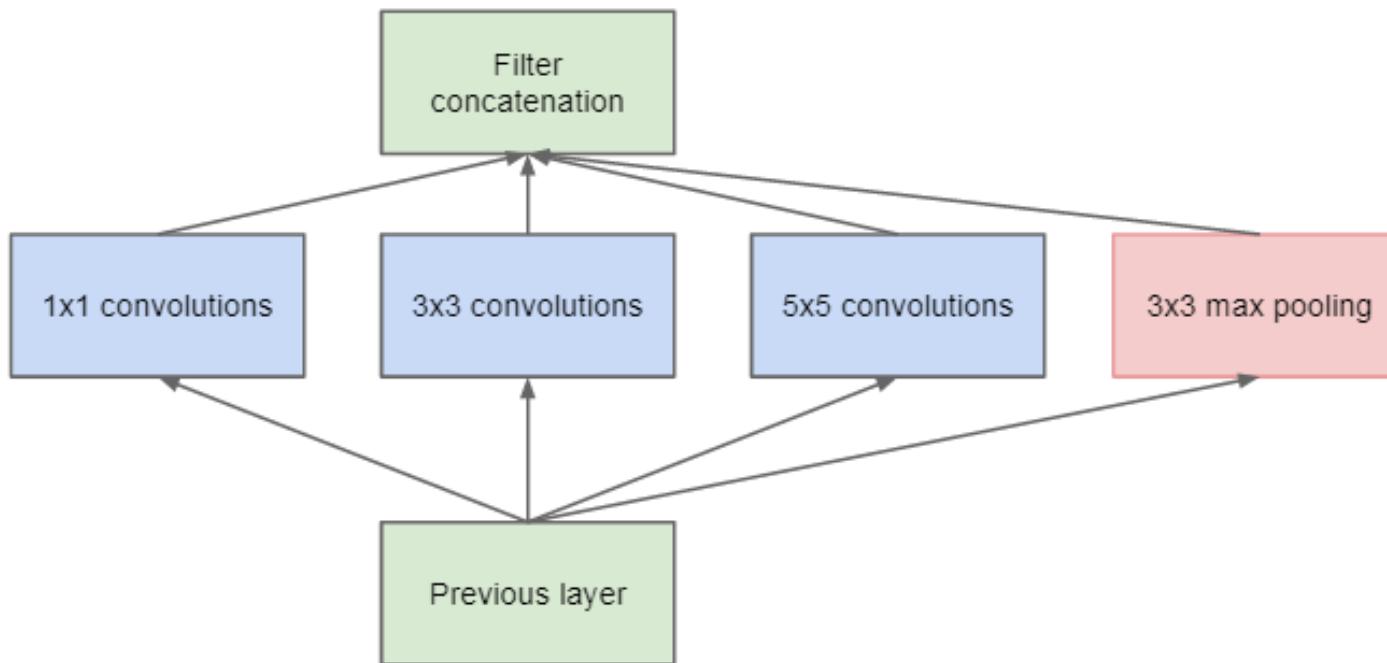
GoogLeNet a.k.a. Inception Network

- Inception Module
 - “Network in a Network”
 - Instead of seeking for a depth, expand breadth!



GoogLeNet a.k.a. Inception Network

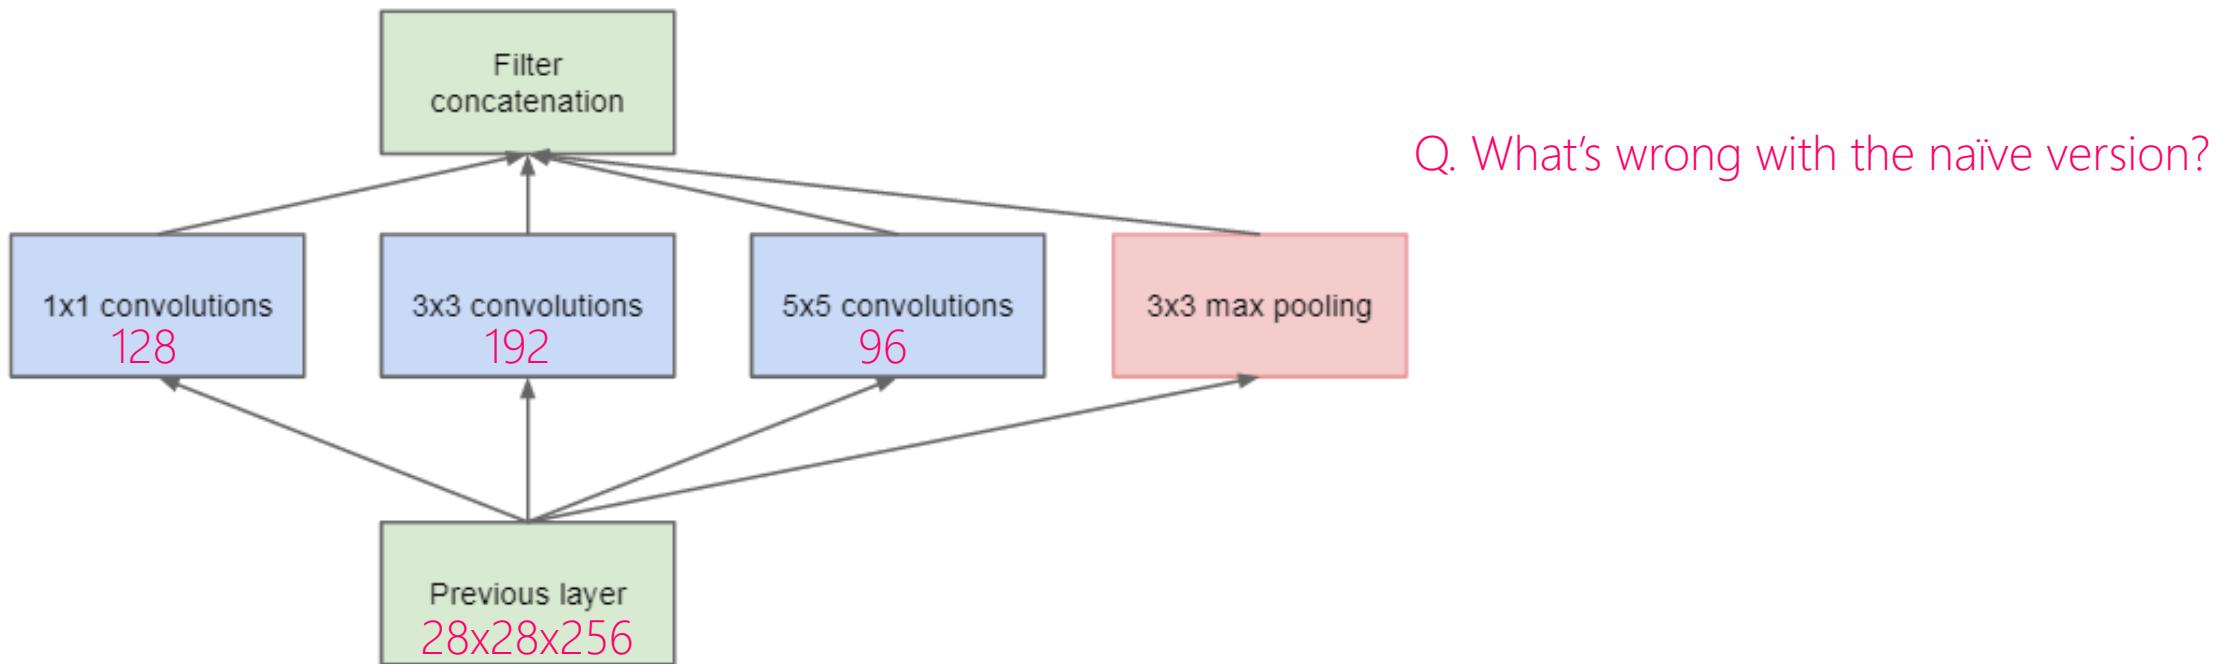
- Naïve Inception Module
 - “Why bother to decide which size of convolution to use when you can have them all?”



(a) Inception module, naïve version

GoogLeNet a.k.a. Inception Network

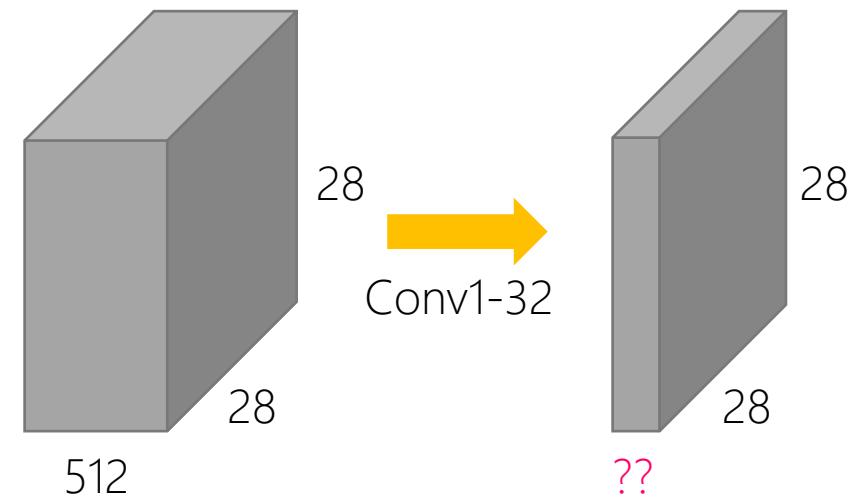
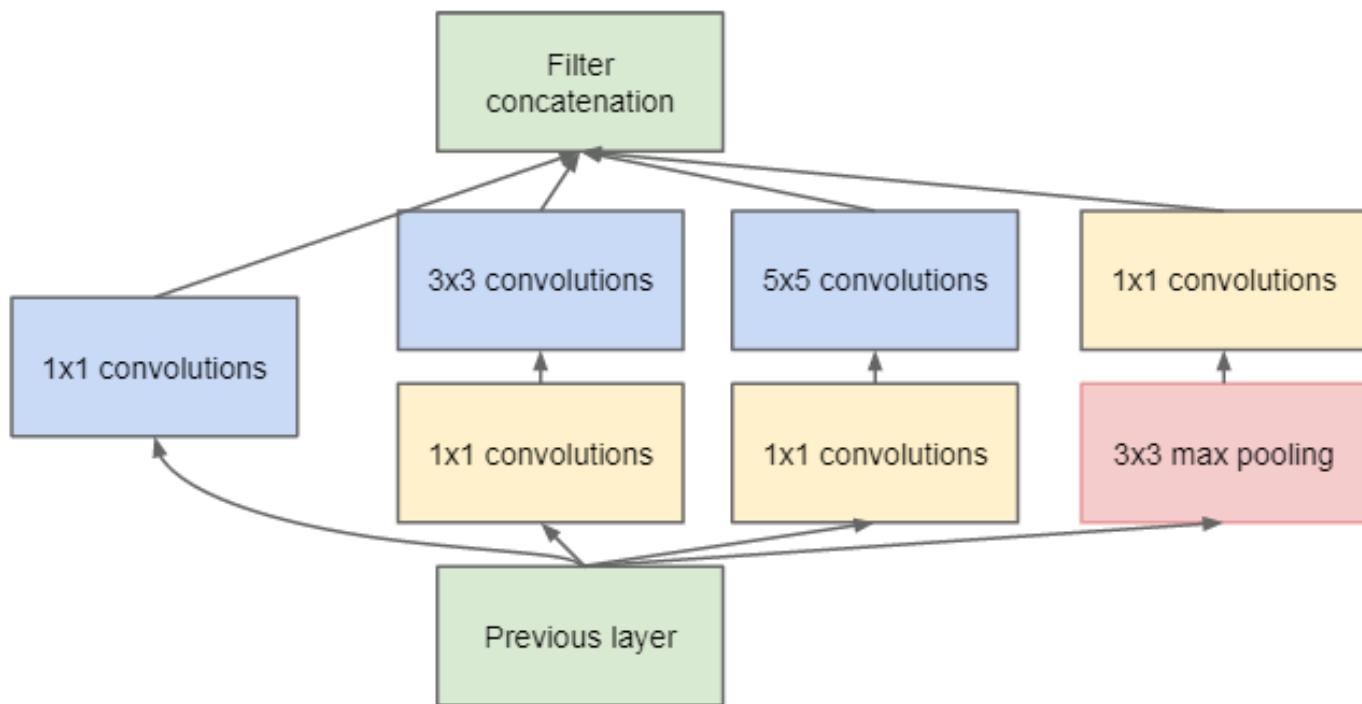
- Naïve Inception Module
 - “Why bother to decide which size of convolution to use when you can have them all?”



(a) Inception module, naïve version

GoogLeNet a.k.a. Inception Network

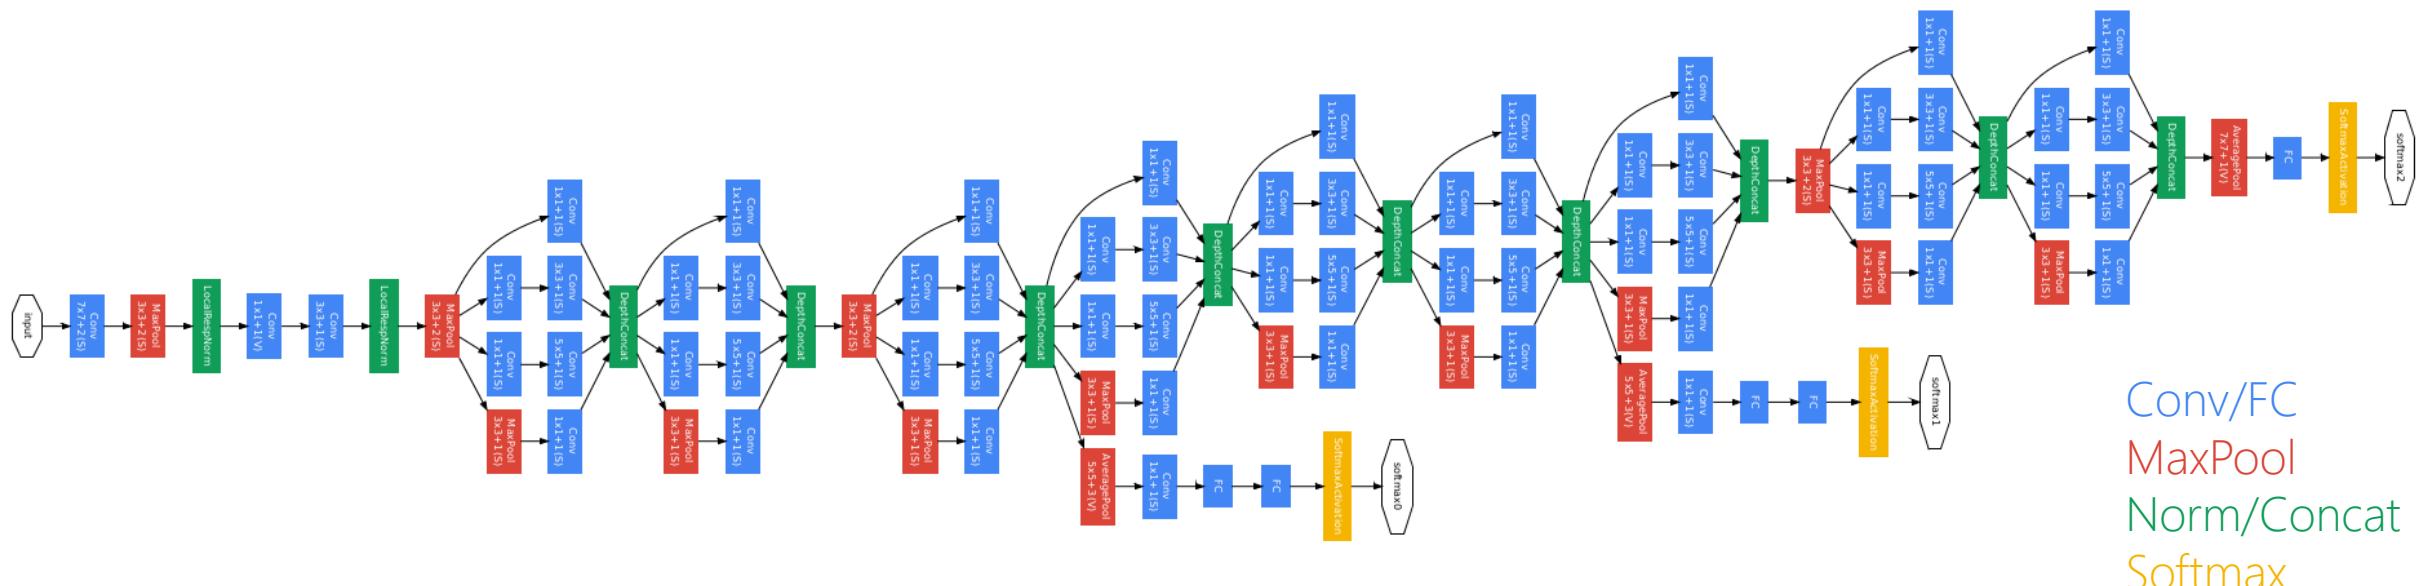
- Inception module with bottleneck layers



(b) Inception module with dimension reductions

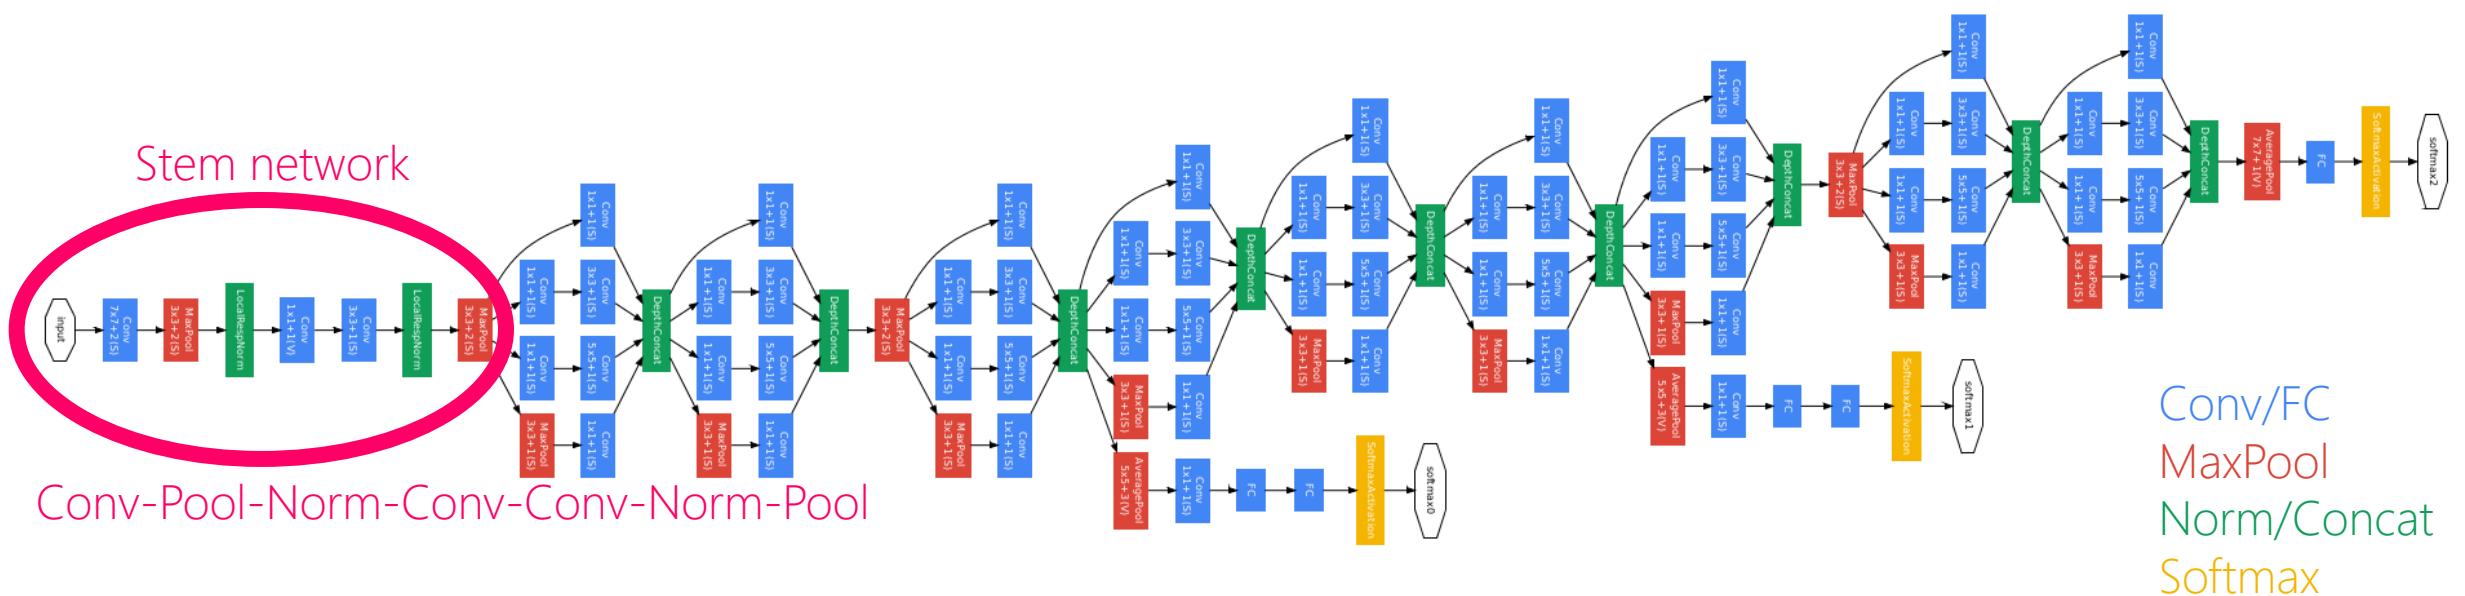
GoogLeNet a.k.a. Inception Network

- Putting them all together...



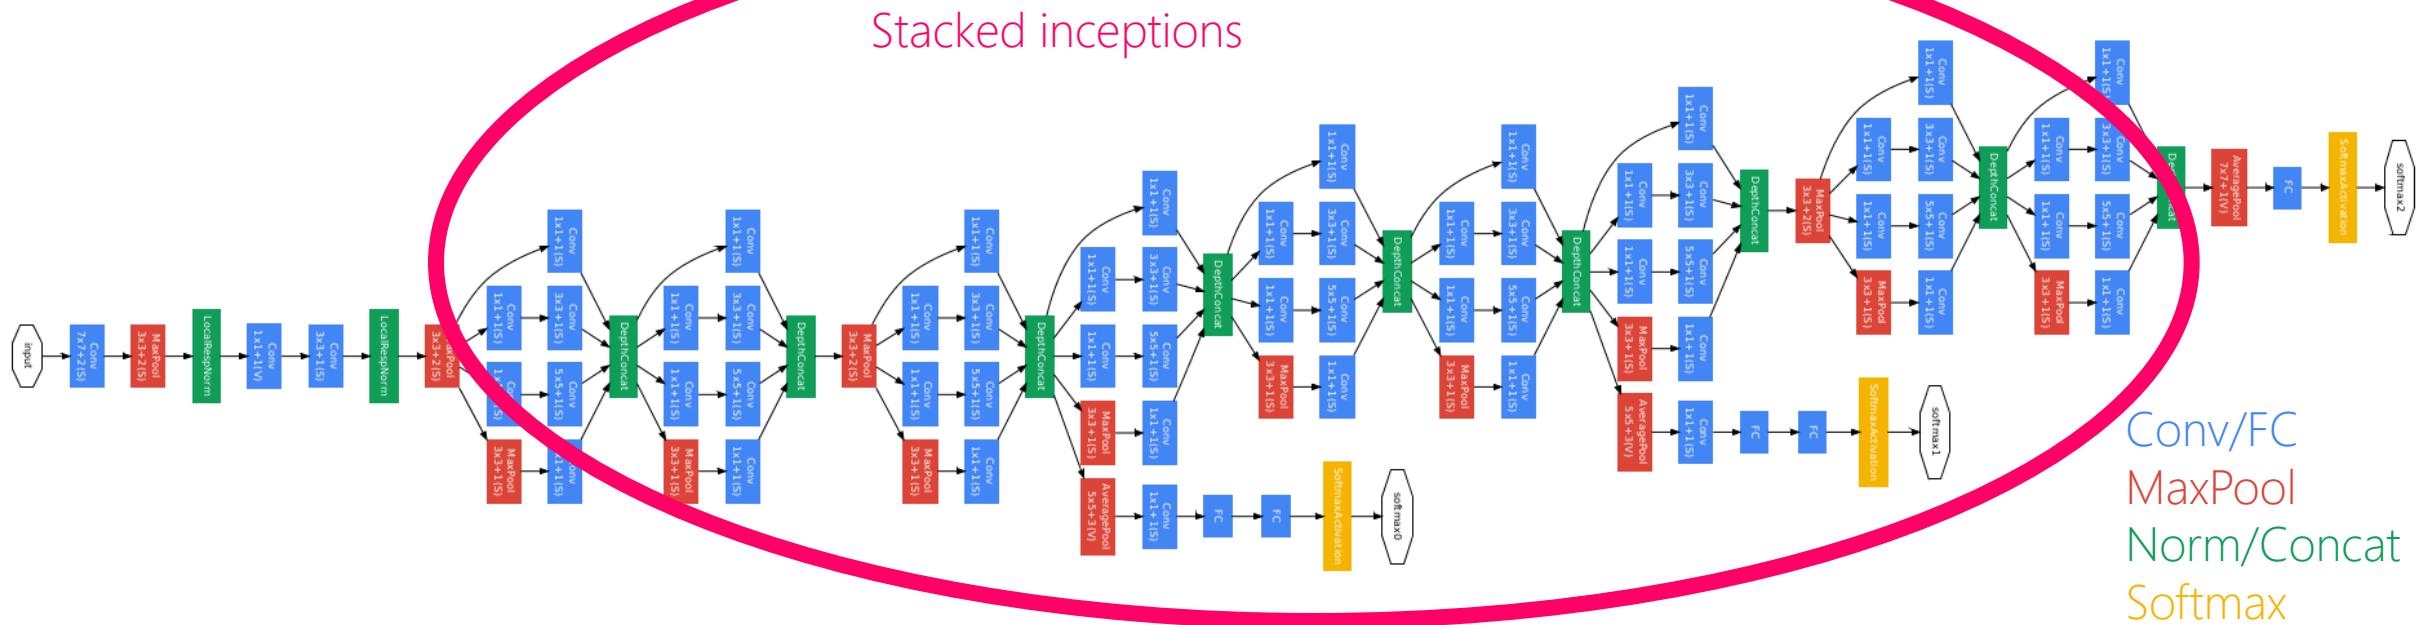
GoogLeNet a.k.a. Inception Network

- Putting them all together...



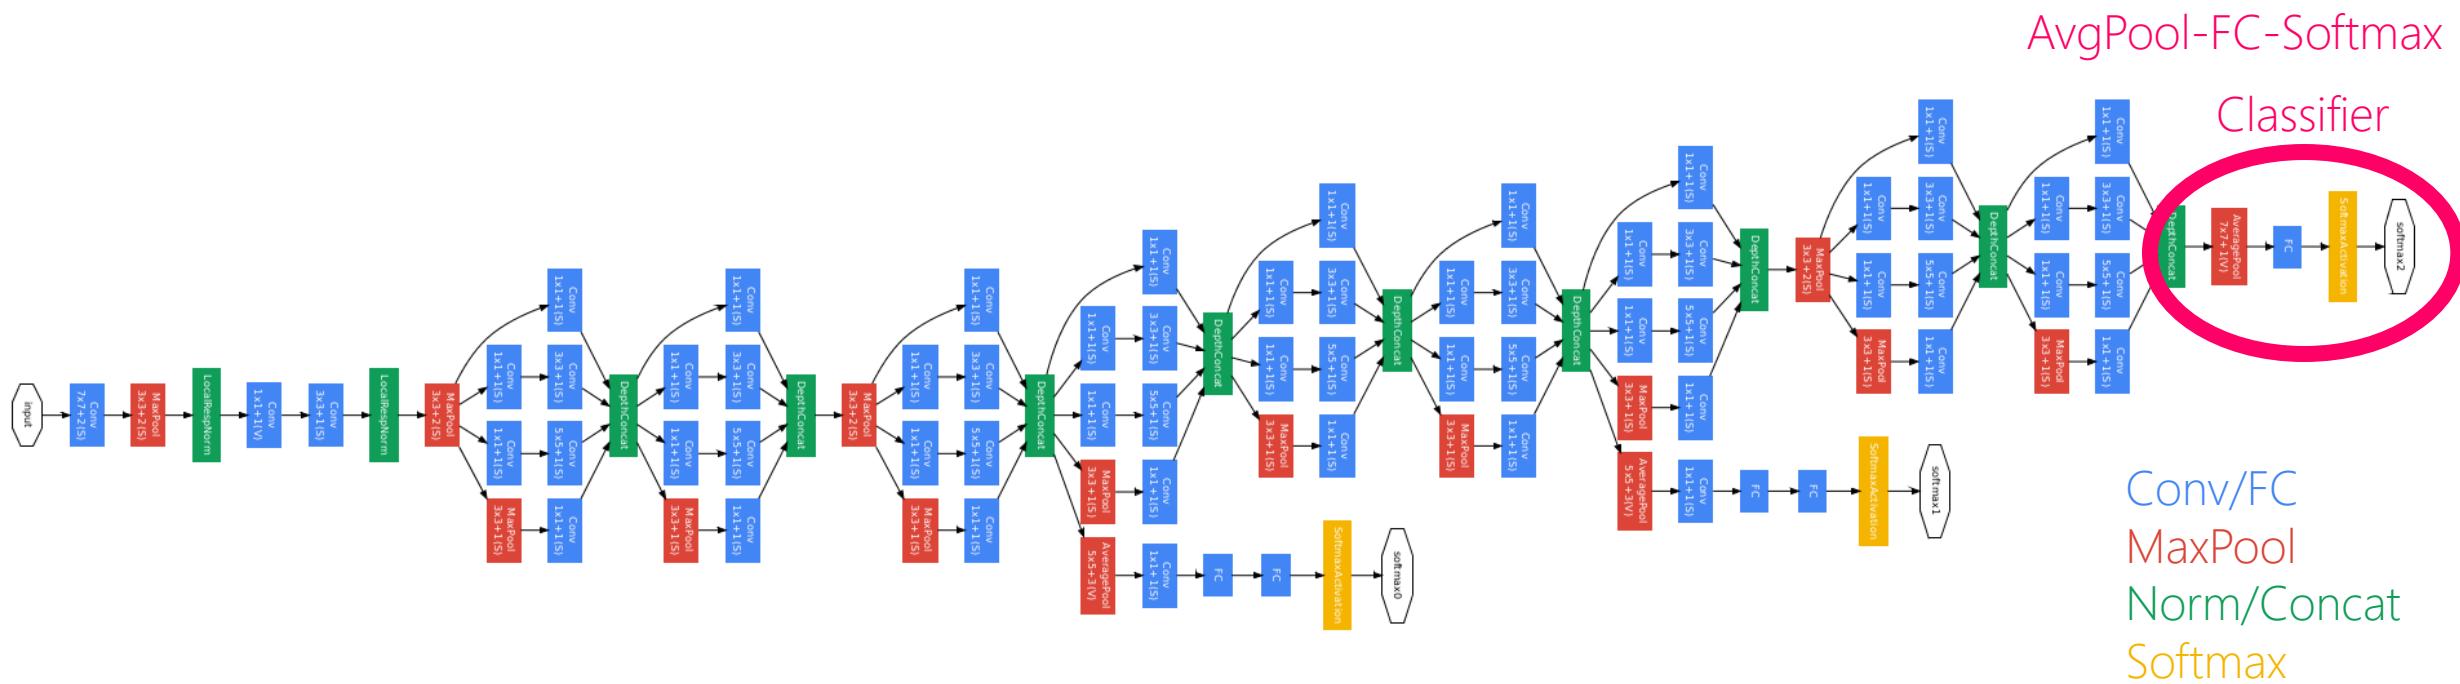
GoogLeNet a.k.a. Inception Network

- Putting them all together...



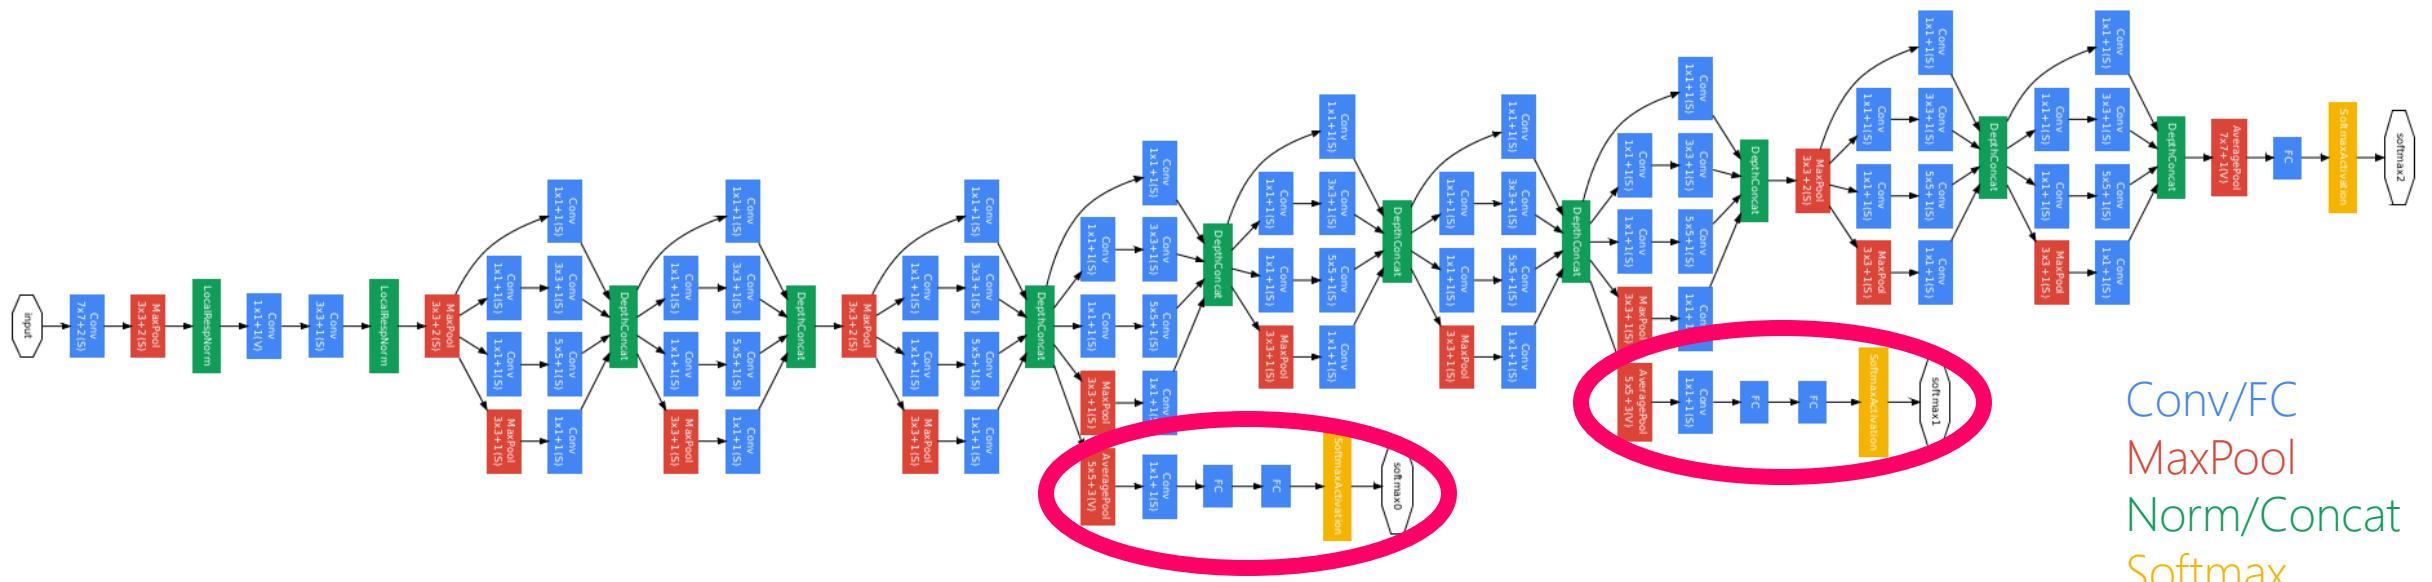
GoogLeNet a.k.a. Inception Network

- Putting them all together...



GoogLeNet a.k.a. Inception Network

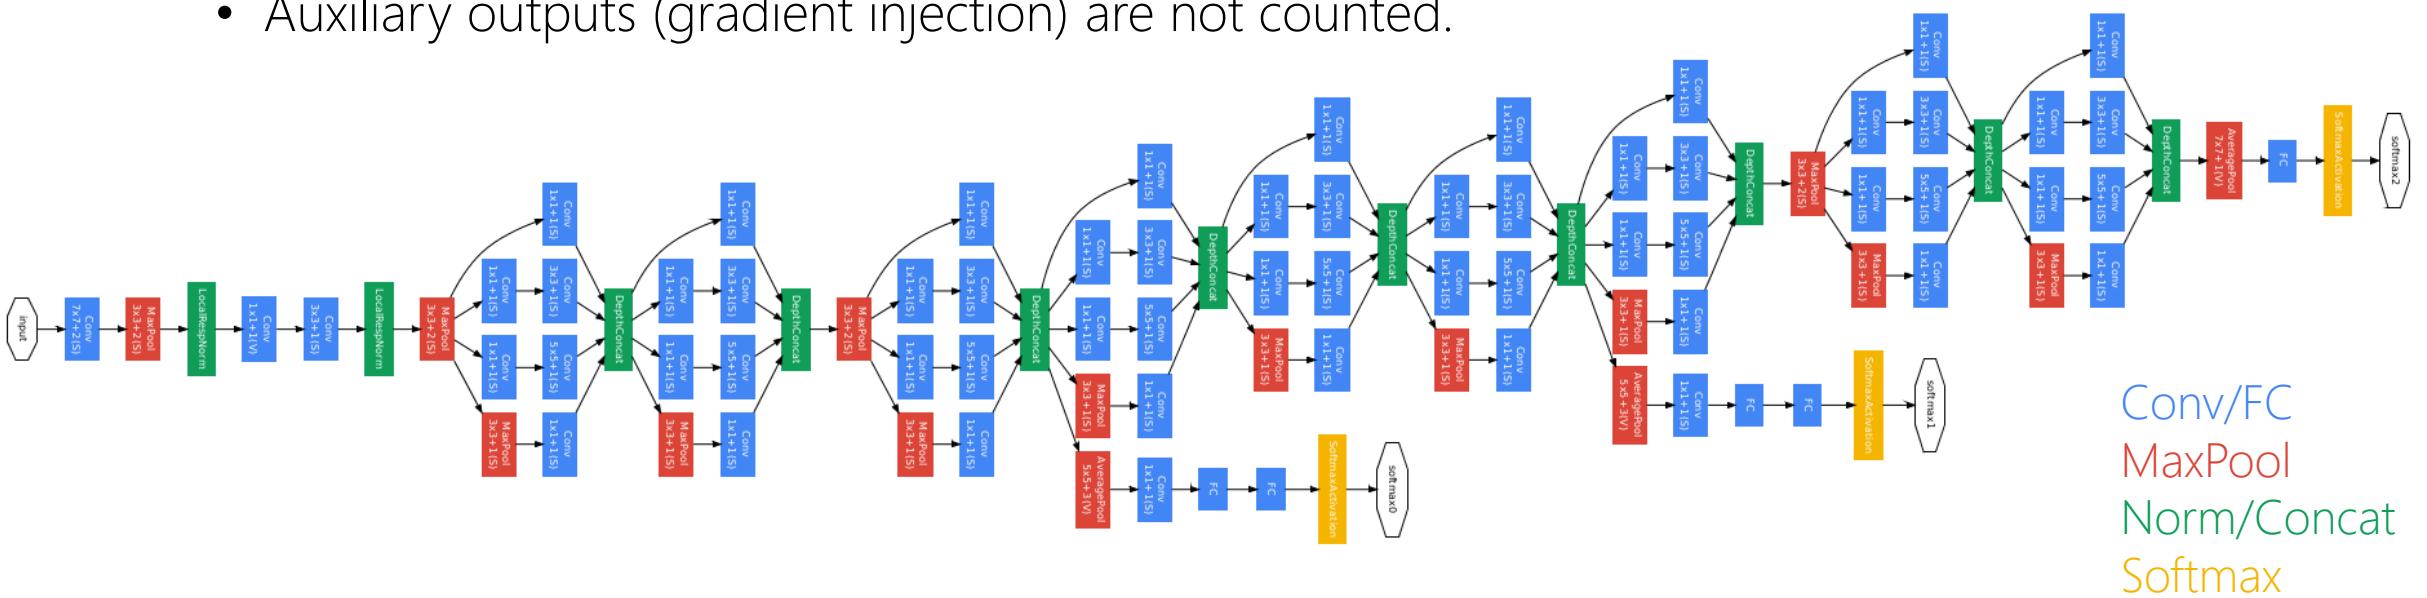
- Putting them all together...



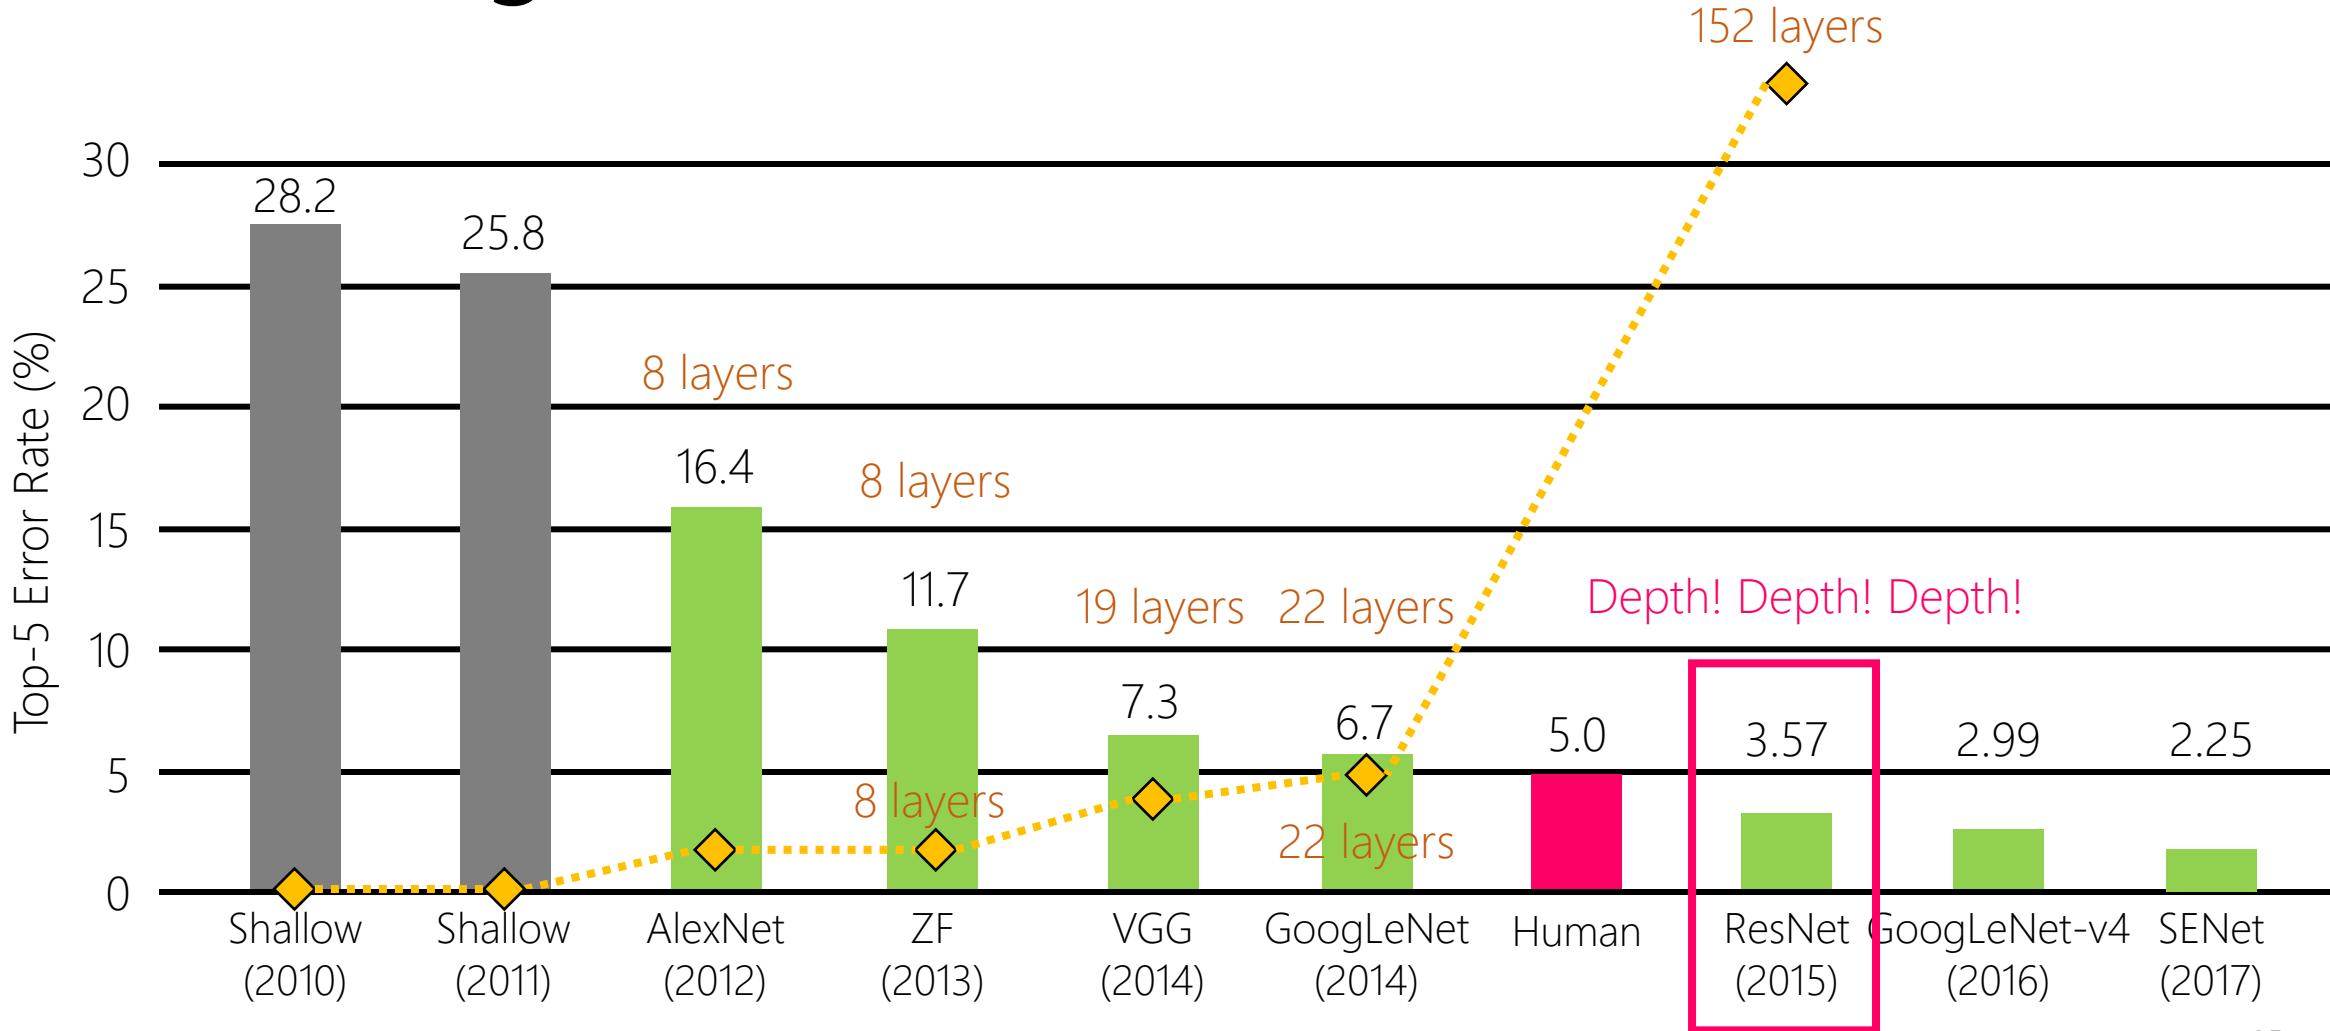
Additional classification losses to inject additional gradients in order to prevent vanishing gradient.

GoogLeNet a.k.a. Inception Network

- How to count 22 layers?
 - Parallel layers are considered as one. (i.e. 2 layers per inception)
 - Auxiliary outputs (gradient injection) are not counted.

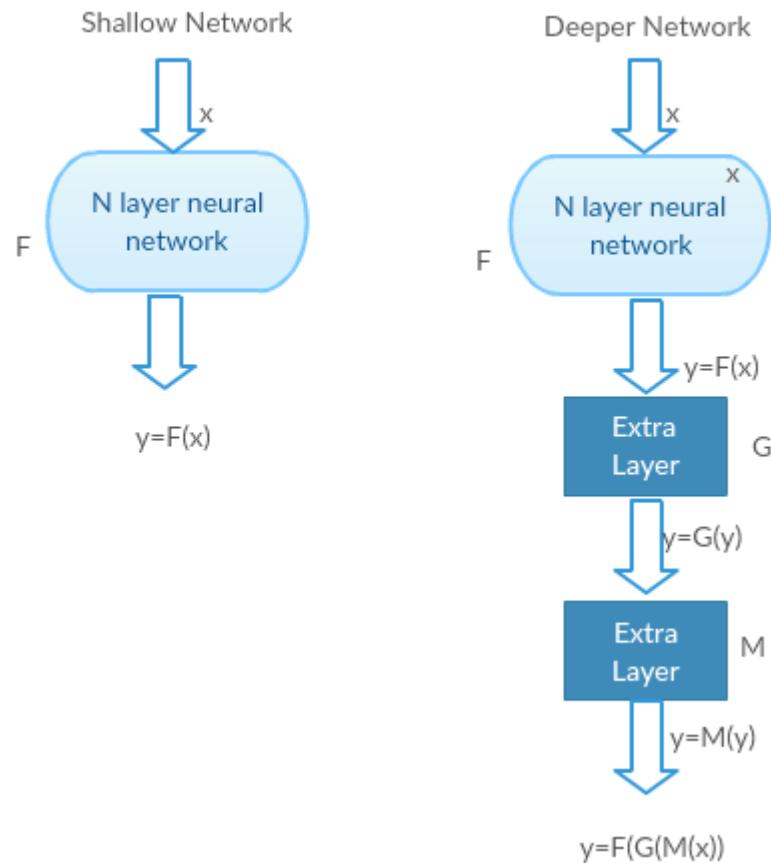


ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners



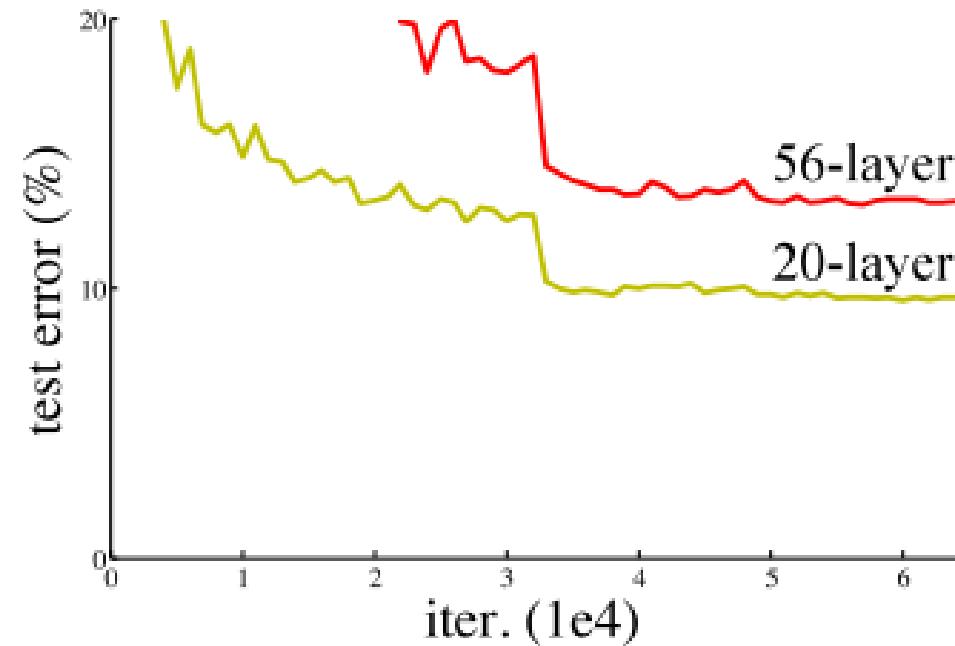
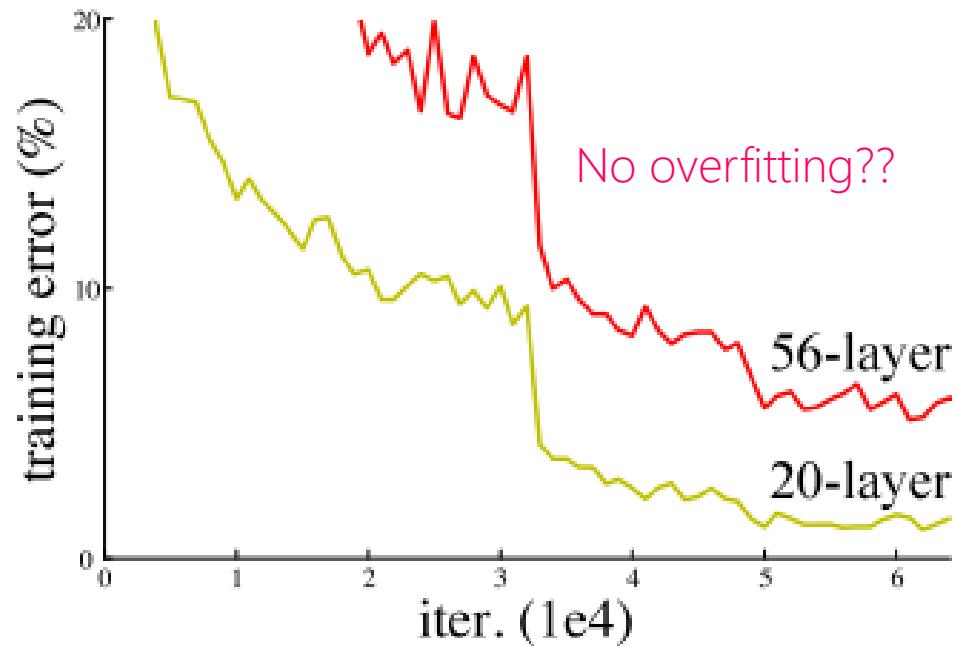
Adding more layers seems to work well, but...

- Deeper models should be able to perform at least as good as shallow models...
- Why?



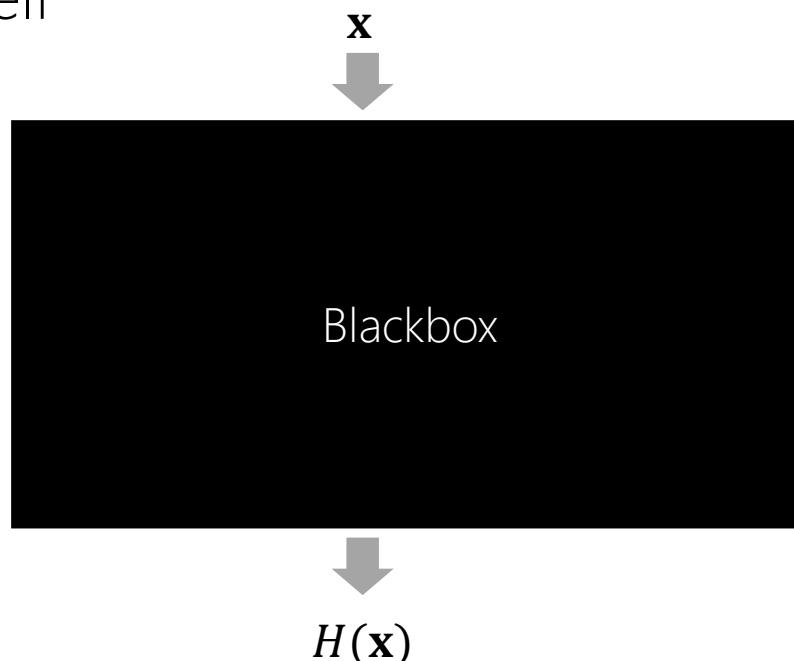
Adding more layers seems to work well, but...

- Degradation of accuracy with more layers...



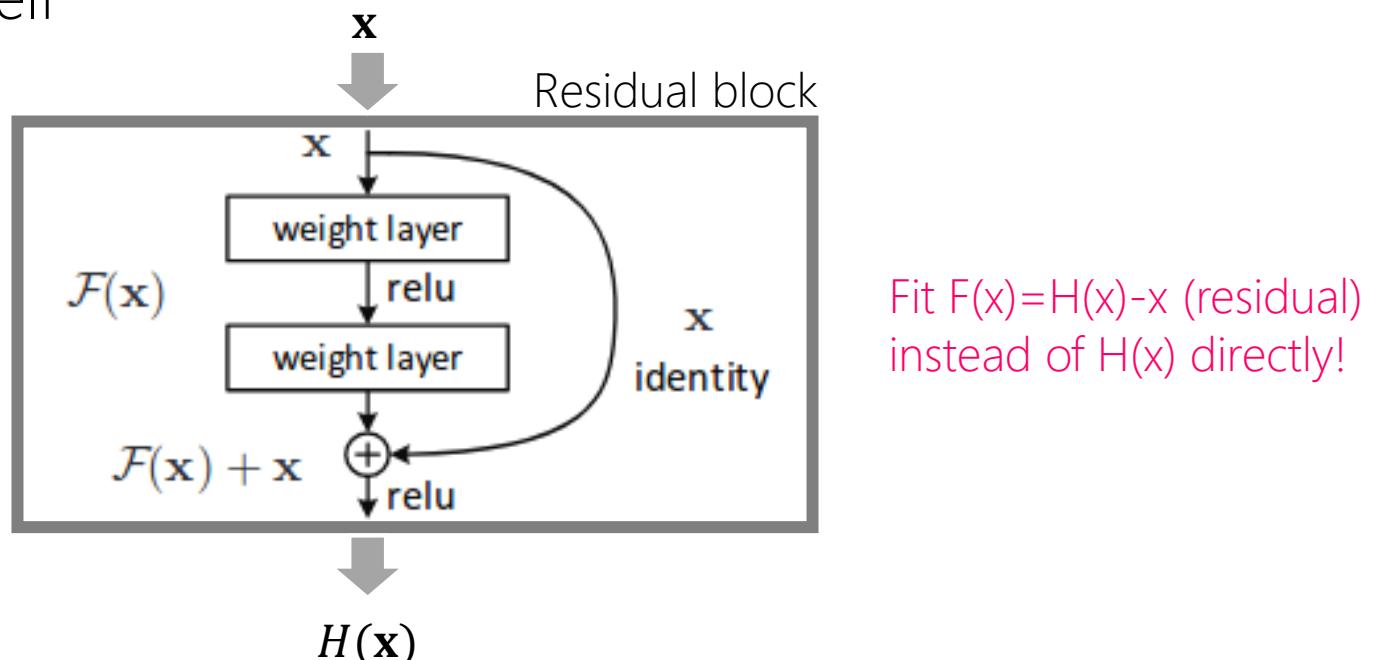
ResNet

- Hypothesis: the degradation problem is due to an optimization problem (deeper models are harder to optimize)
- Proposed solution: Let's fit the residual (i.e. the extra accuracy gain of a conv block compared to the identity mapping) instead of directly fitting the underlying mapping itself



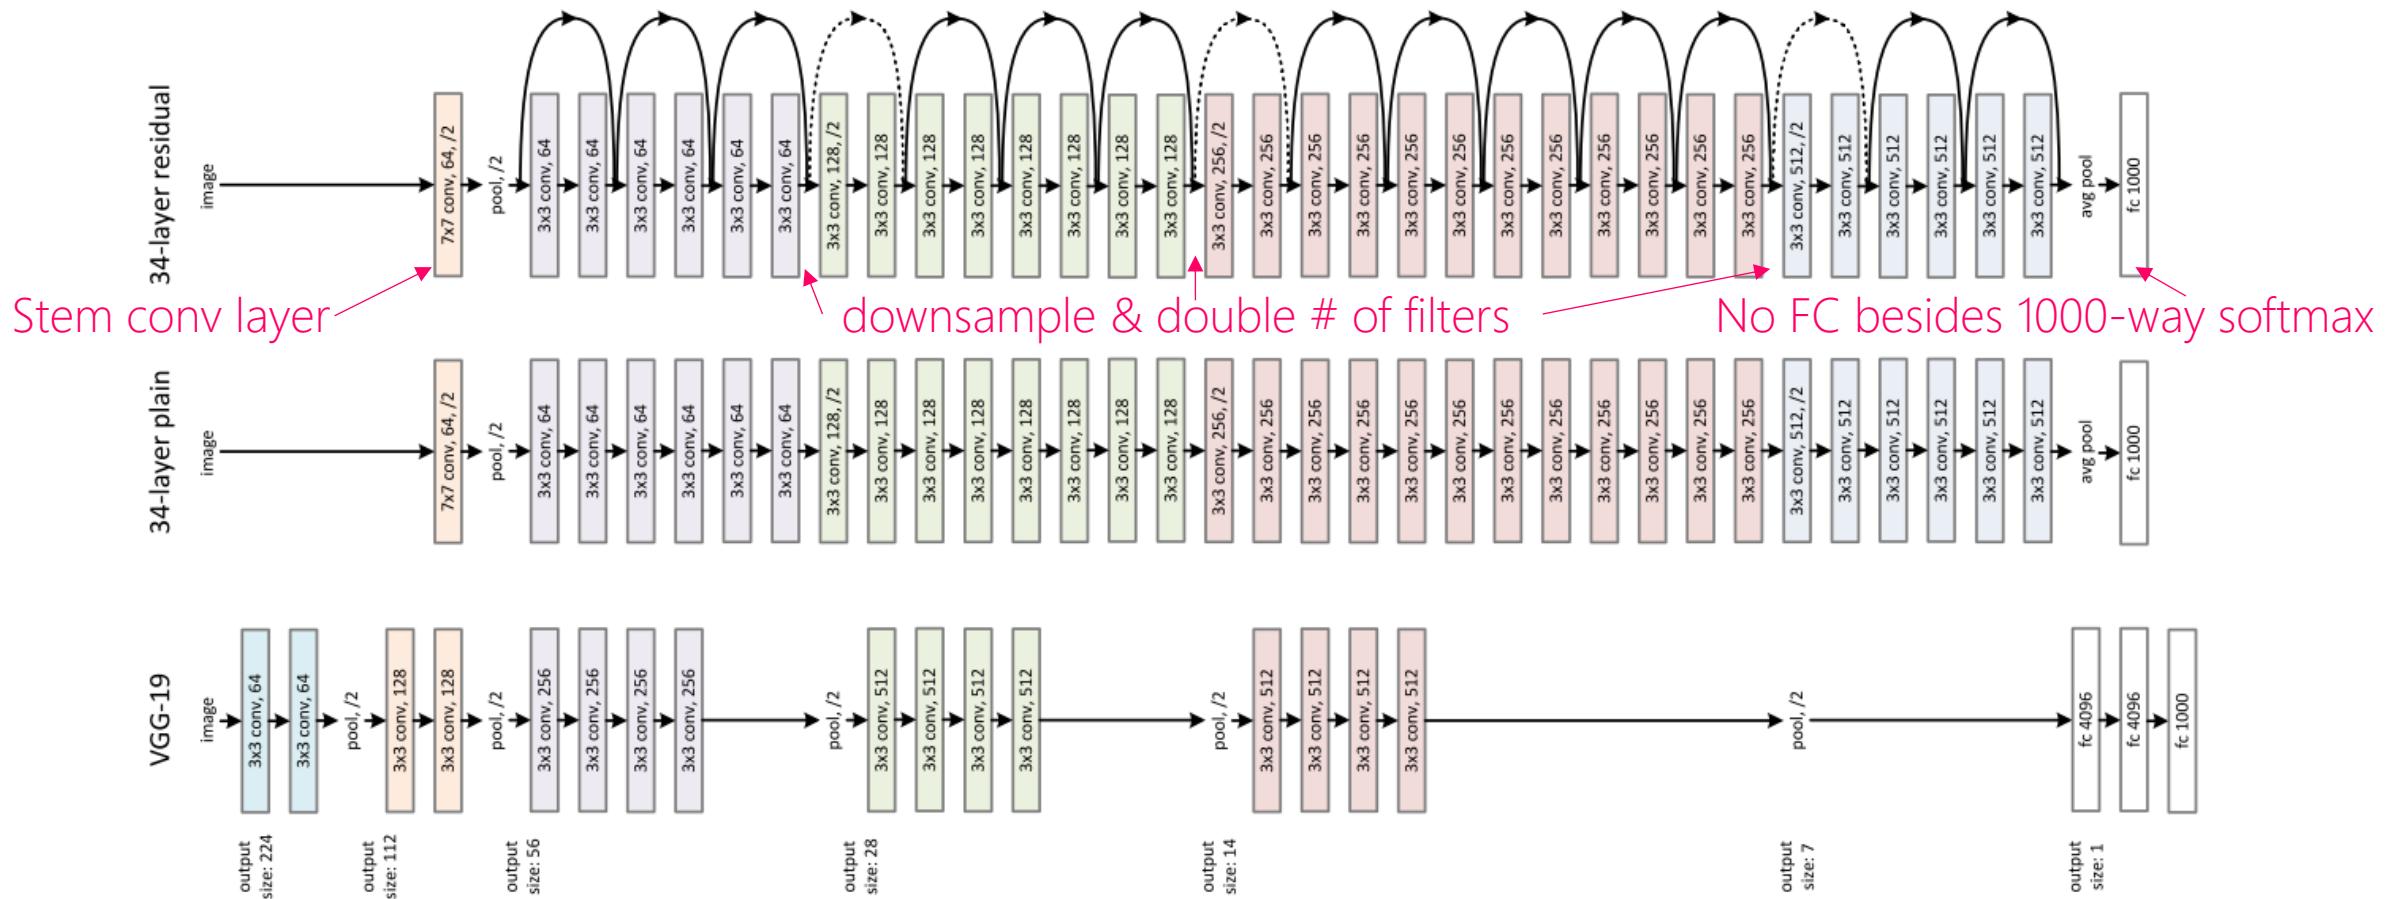
ResNet

- Hypothesis: the degradation problem is due to an optimization problem (deeper models are harder to optimize)
- Proposed solution: Let's fit the residual (i.e. the extra accuracy gain of a conv block compared to the identity mapping) instead of directly fitting the underlying mapping itself



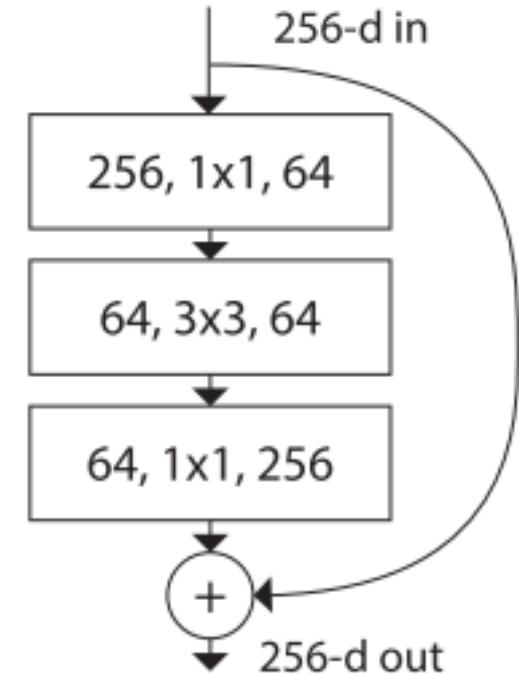
ResNet

- Full ResNet architecture: stack residual blocks each of which has two Conv3.



ResNet

- ResNet-34, ResNet-50, ResNet-101, ResNet-152 available...
 - ResNet-50+ uses bottleneck layers to improve efficiency
 - Batch normalization after every conv
 - SGD with Momentum 0.9
 - Learning rate: 0.1, divided by 10 when val error plateaus
 - Mini-batch size 256
 - No dropout
- Can be trained without degradation



ResNet

- Swept 1st place in all ILSVRC and COCO 2015 Competitions

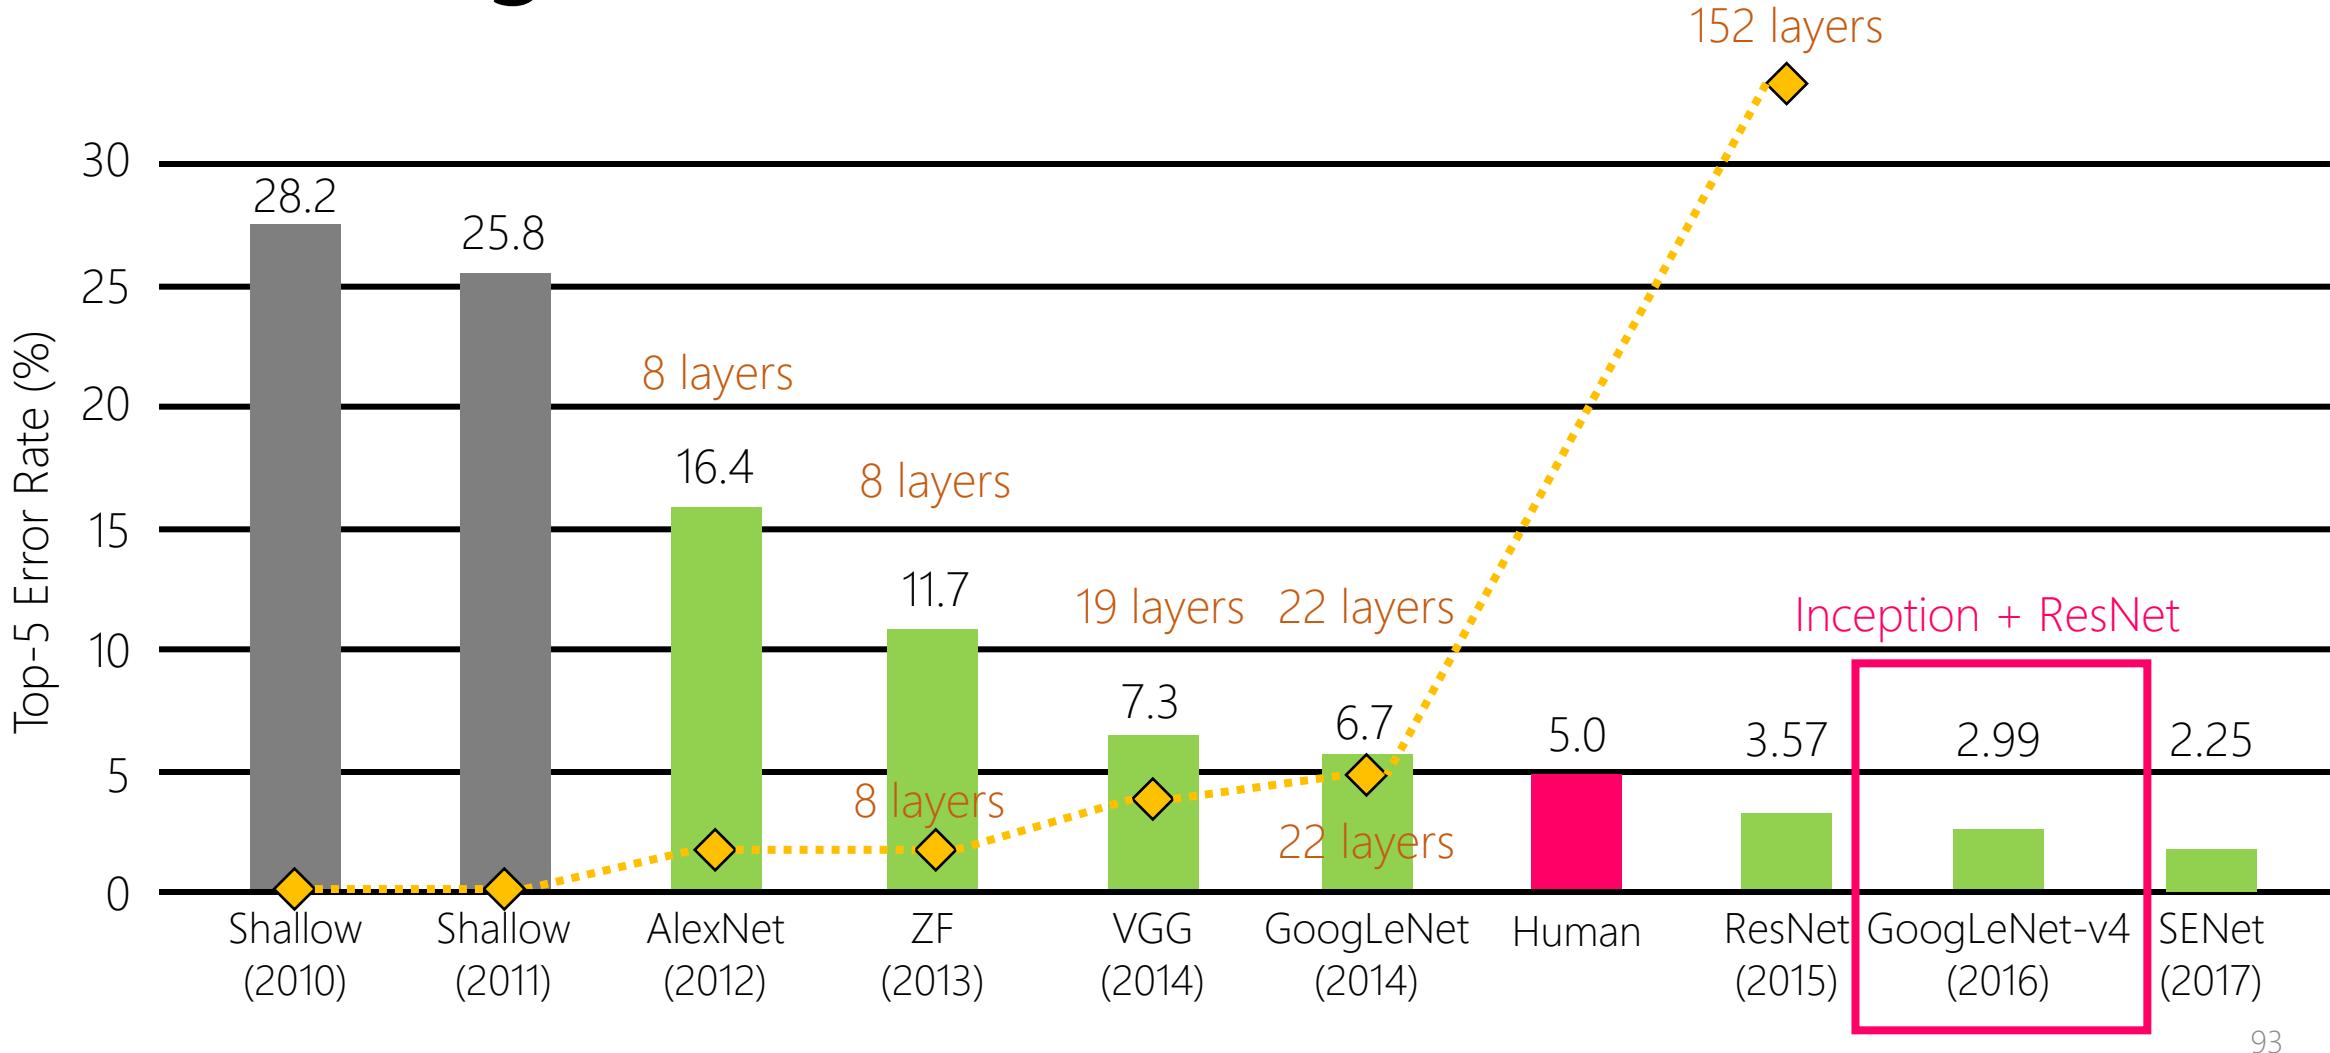
MSRA @ ILSVRC & COCO 2015 Competitions

- **1st places in all five main tracks**

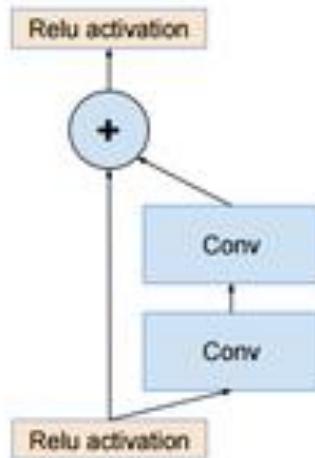
- ImageNet Classification: "*Ultra-deep*" (quote Yann) **152-layer nets**
- ImageNet Detection: **16%** better than 2nd
- ImageNet Localization: **27%** better than 2nd
- COCO Detection: **11%** better than 2nd
- COCO Segmentation: **12%** better than 2nd

- Better than human performance now!

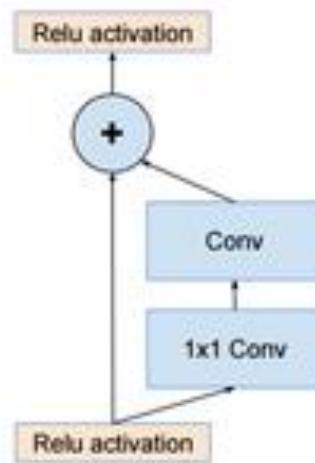
ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners



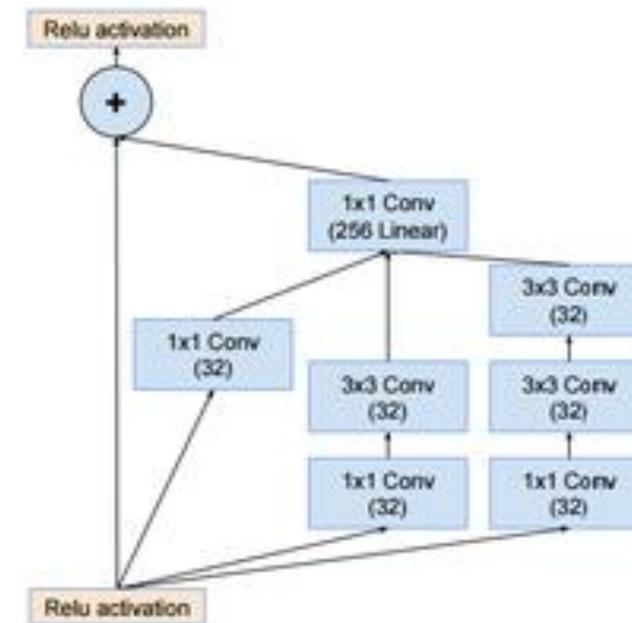
Inception v4 (a.k.a. Inception-ResNet)



(a) Residual module introduced by He et al. (2015)

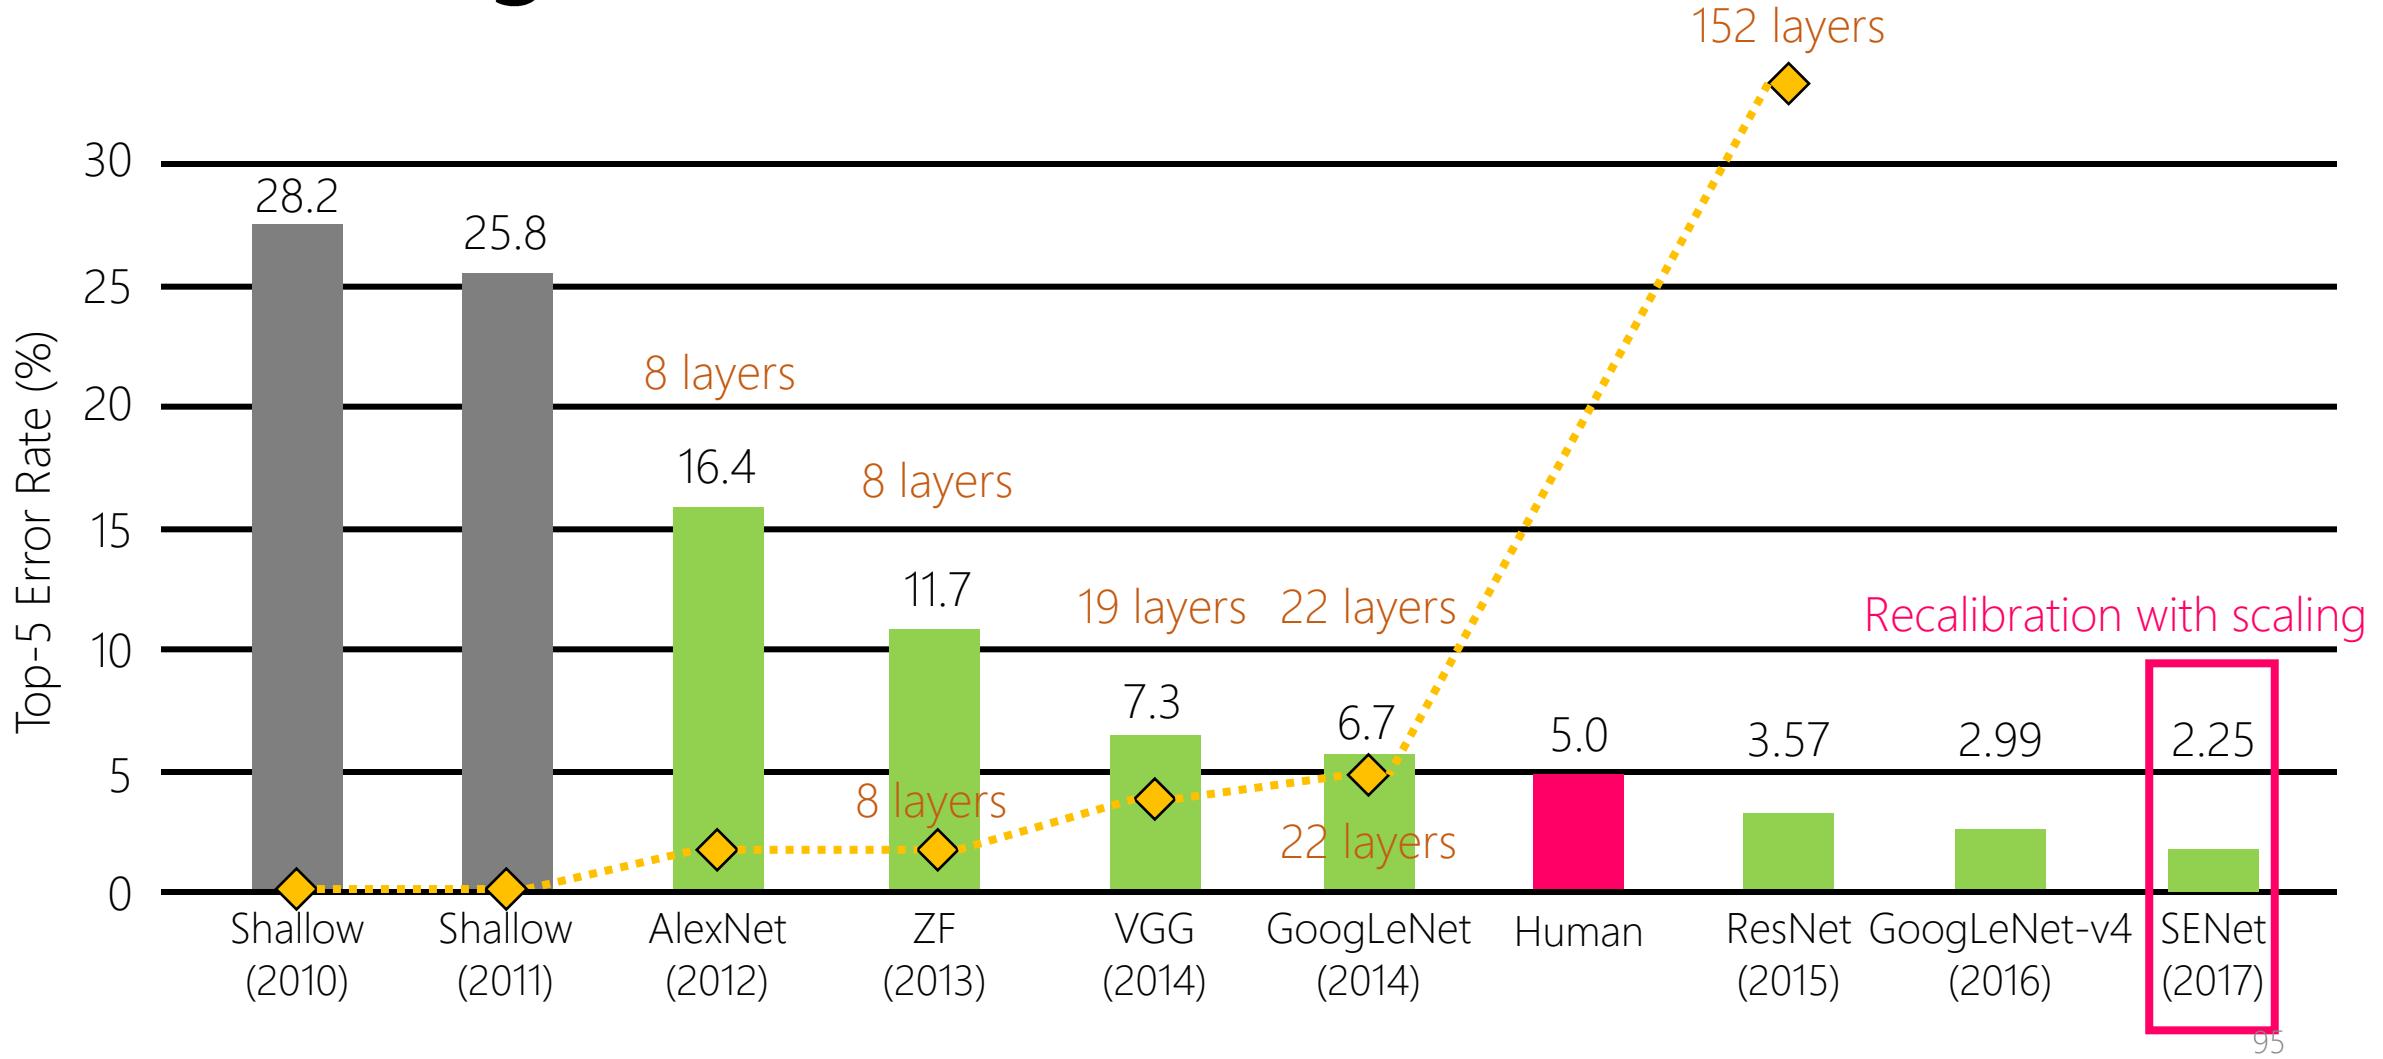


(b) Optimized version of residual modules to shield computation, like in Lin et al. (2013).



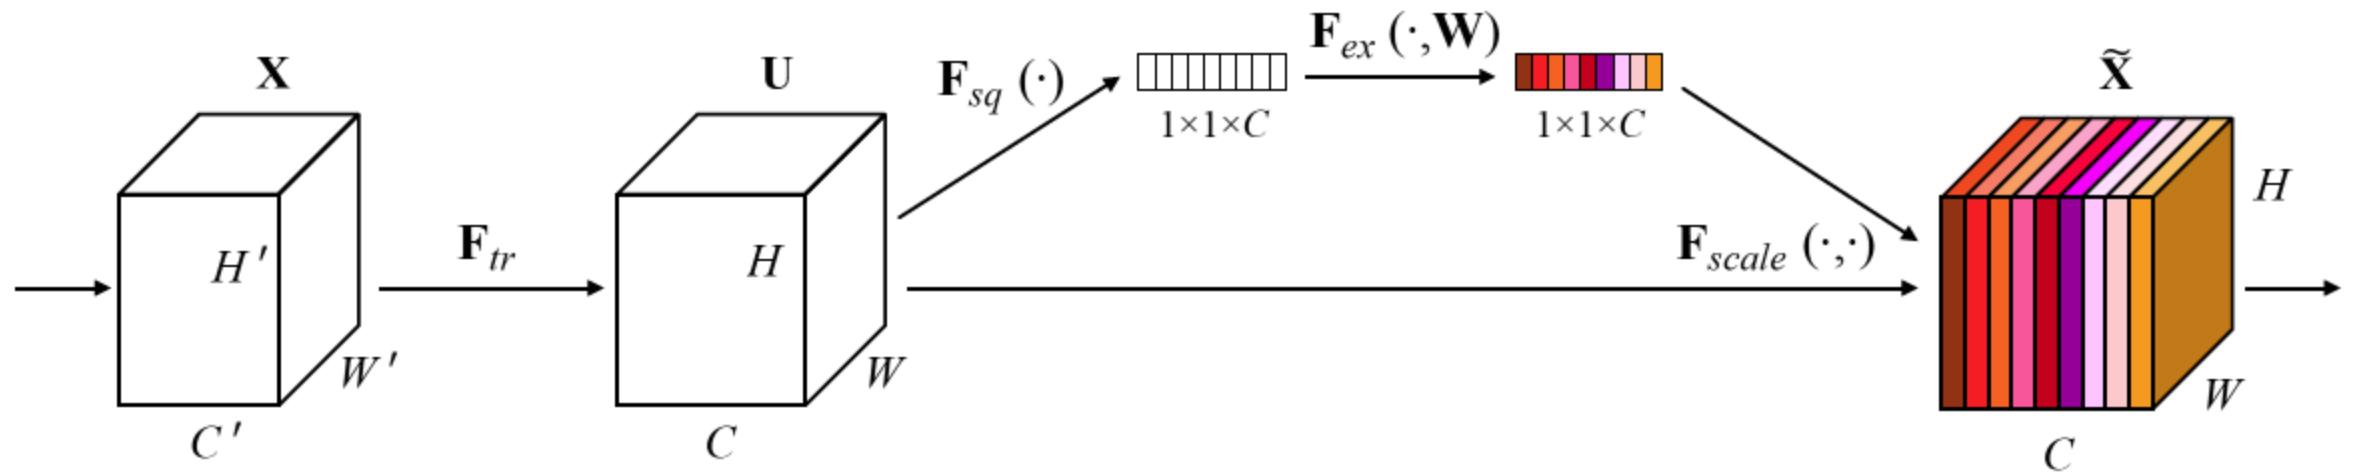
(c) Inception-ResNet example module for the 35×35 grid.

ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners



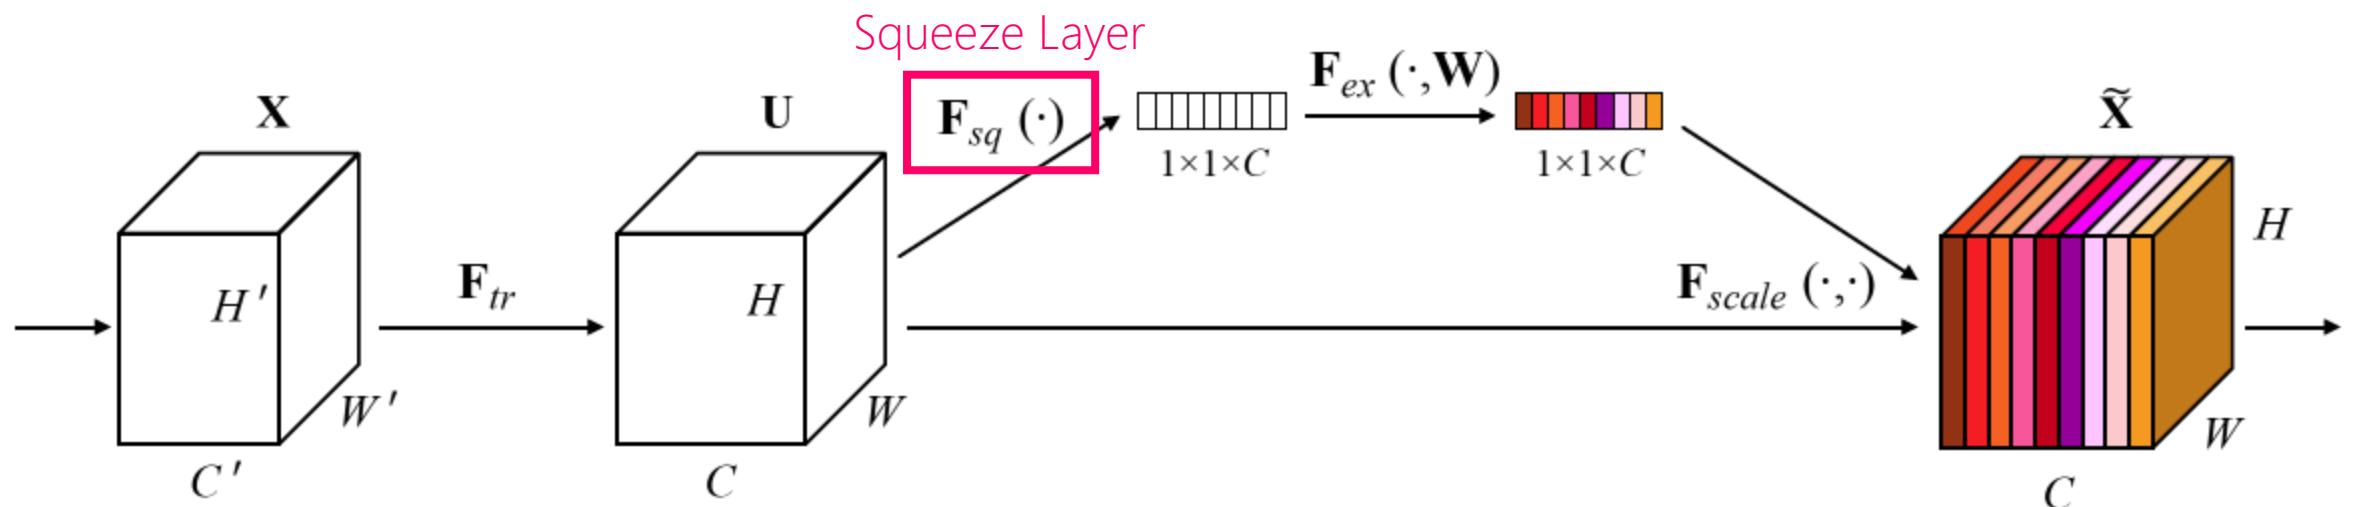
Squeeze and Excitation (SE) Network

- Recalibration of the activation maps with the global average



Squeeze and Excitation (SE) Network

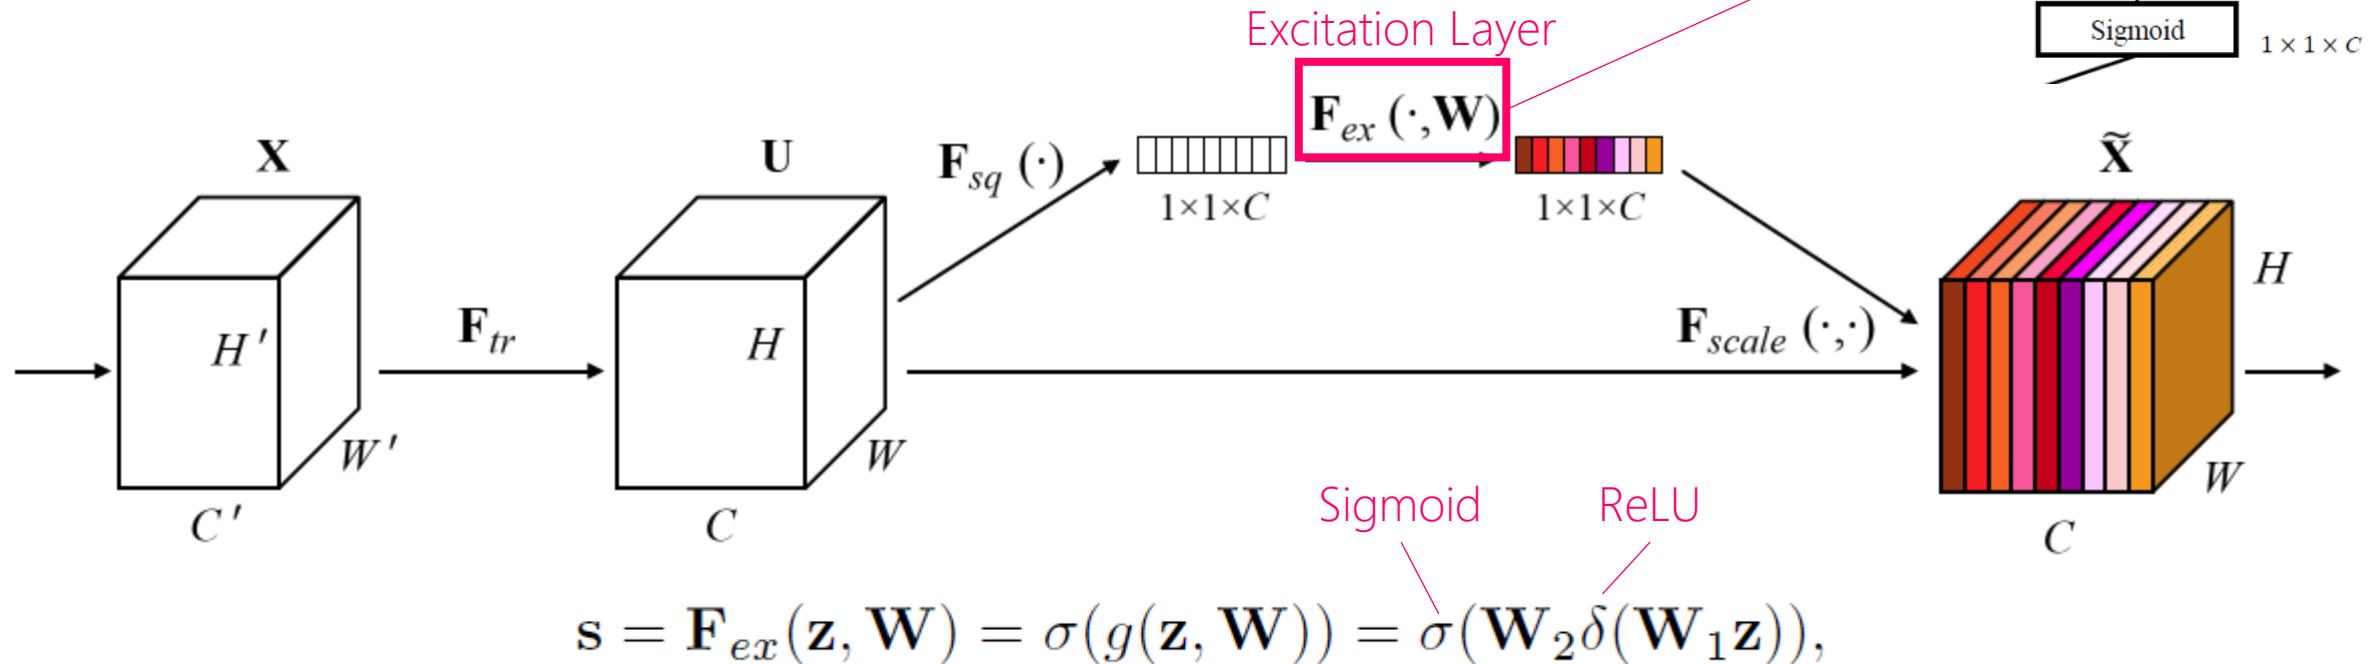
- Recalibration of the activation maps with the global average



$$z_c = \mathbf{F}_{sq}(\mathbf{u}_c) = \frac{1}{H \times W} \sum_{i=1}^H \sum_{j=1}^W u_c(i, j).$$

Squeeze and Excitation (SE) Network

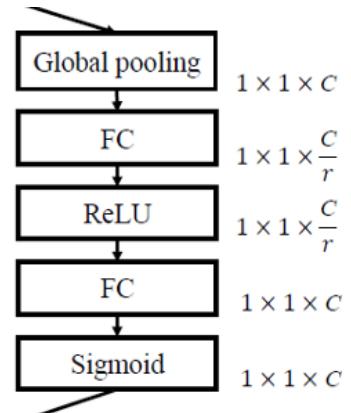
- Recalibration of the activation maps with the global average



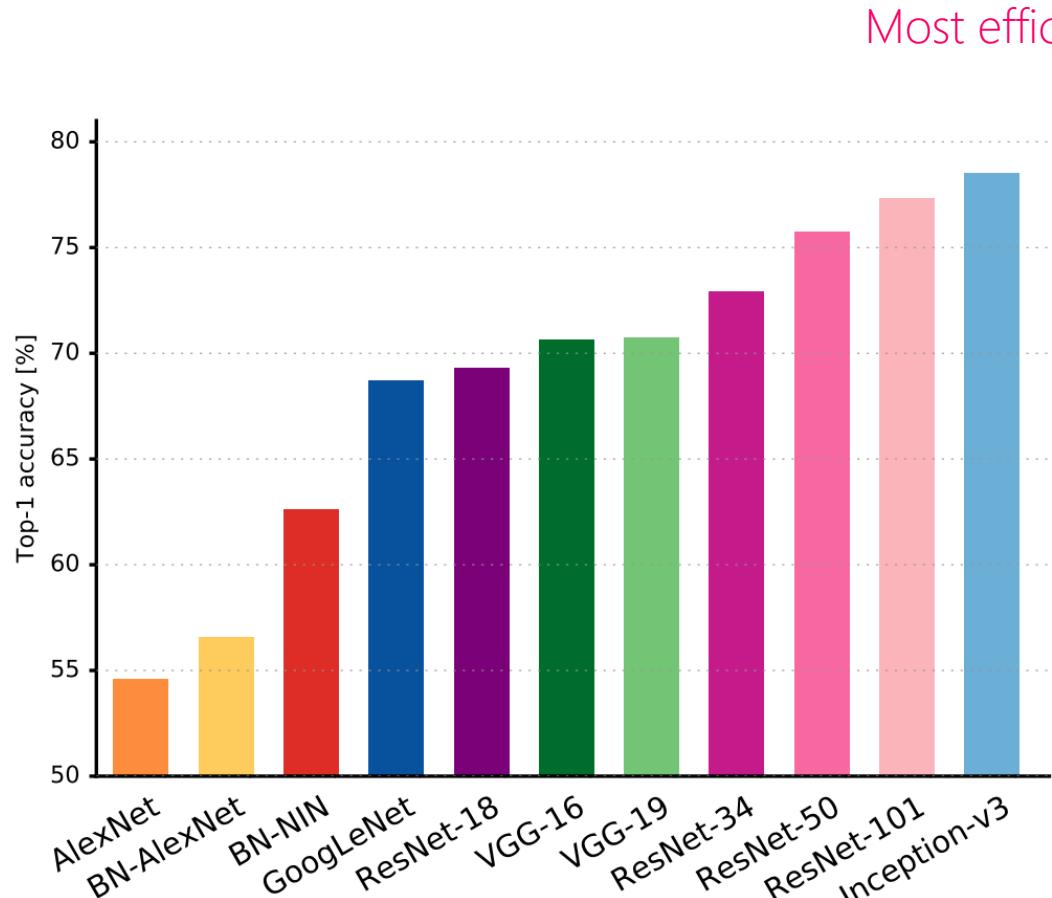
r: Dimensionality reduction ratio (4, 8, 16, 32)

Captures channel-wise dependencies

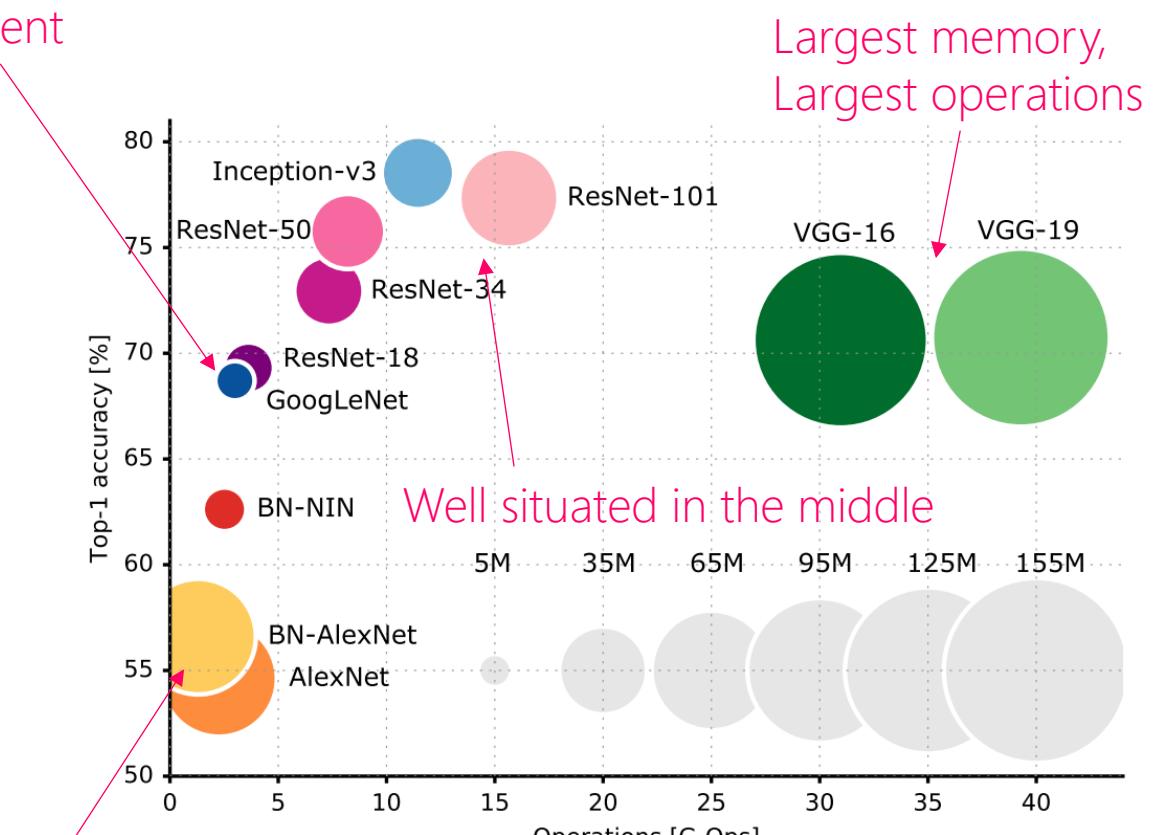
Learns a nonlinear and non-mutually-exclusive relationship between channels.



Alex, VGG, ResNet, Inception at a glance

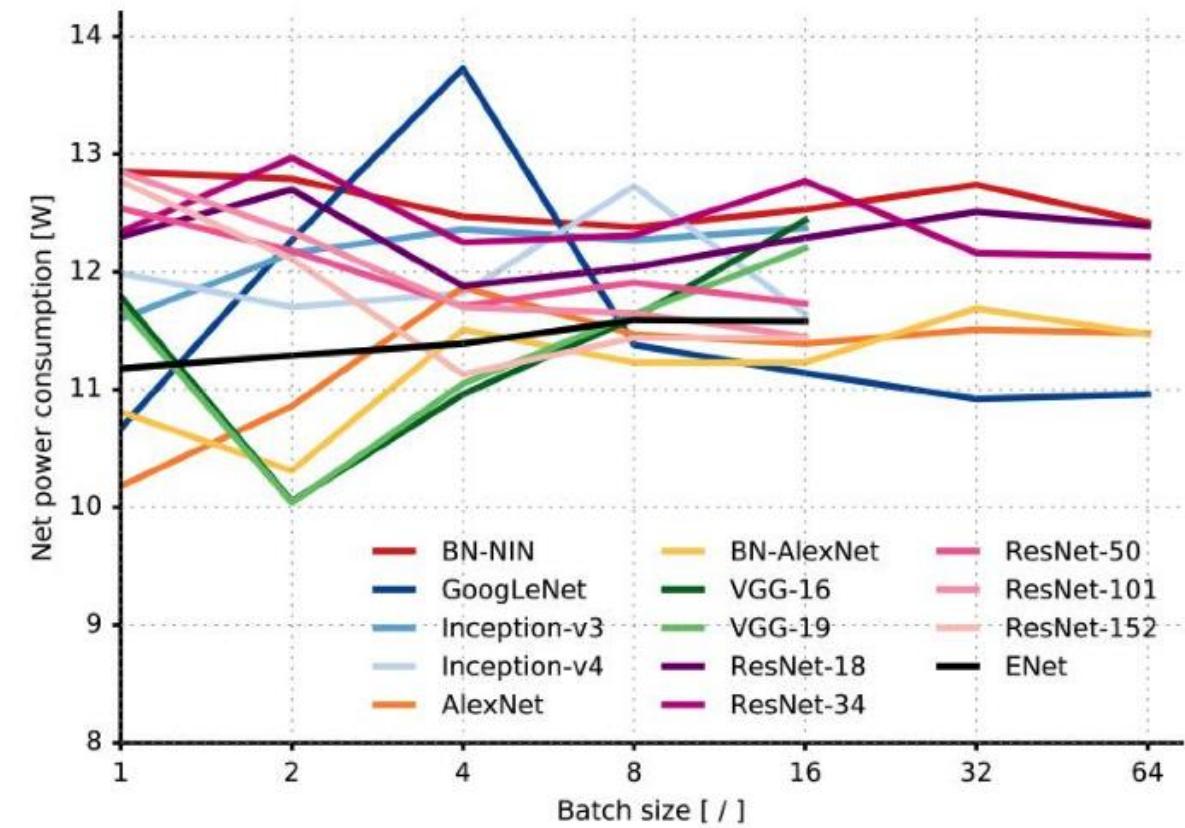
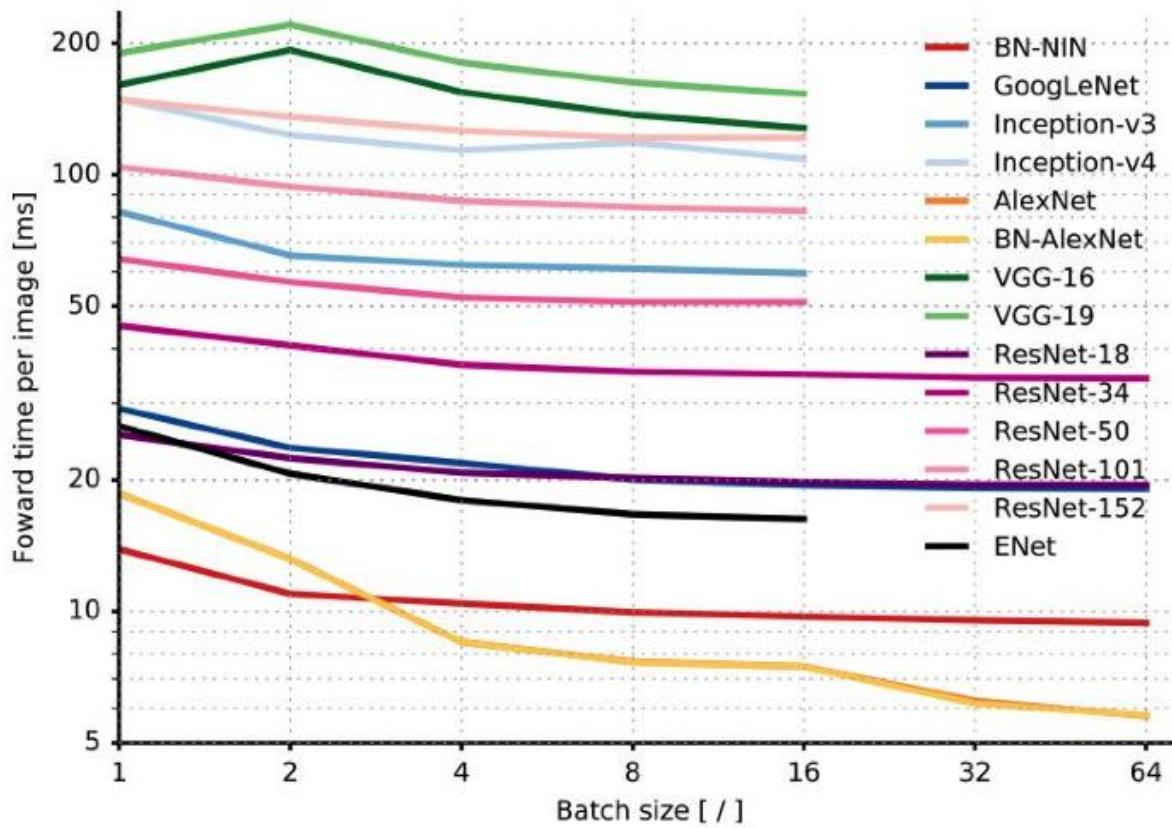


Most efficient



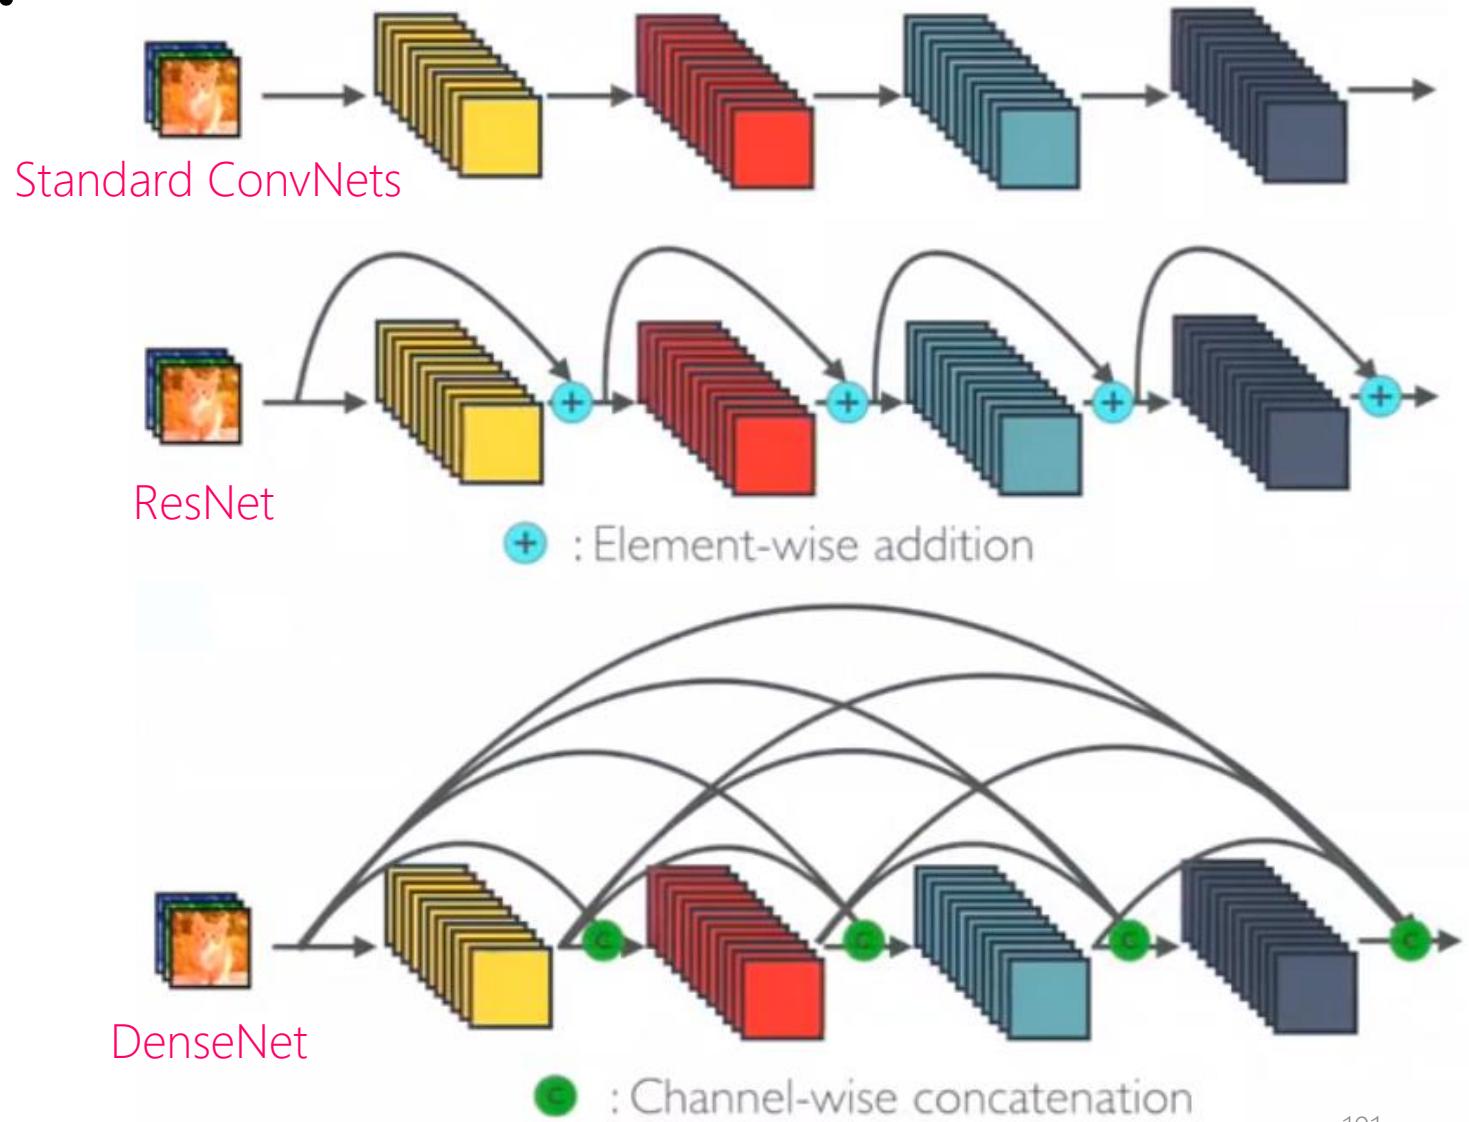
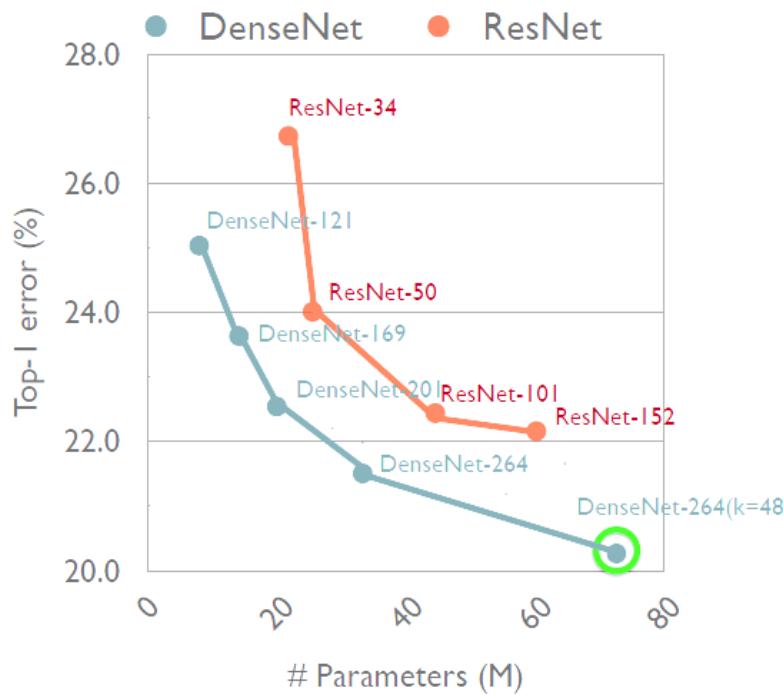
Smaller computation, but memory heavy
Lower accuracy

Alex, VGG, ResNet, Inception at a glance



Other Variants...

- DenseNet (CVPR 2017)



Other Variants...

- ResNeXt (CVPR 2017)

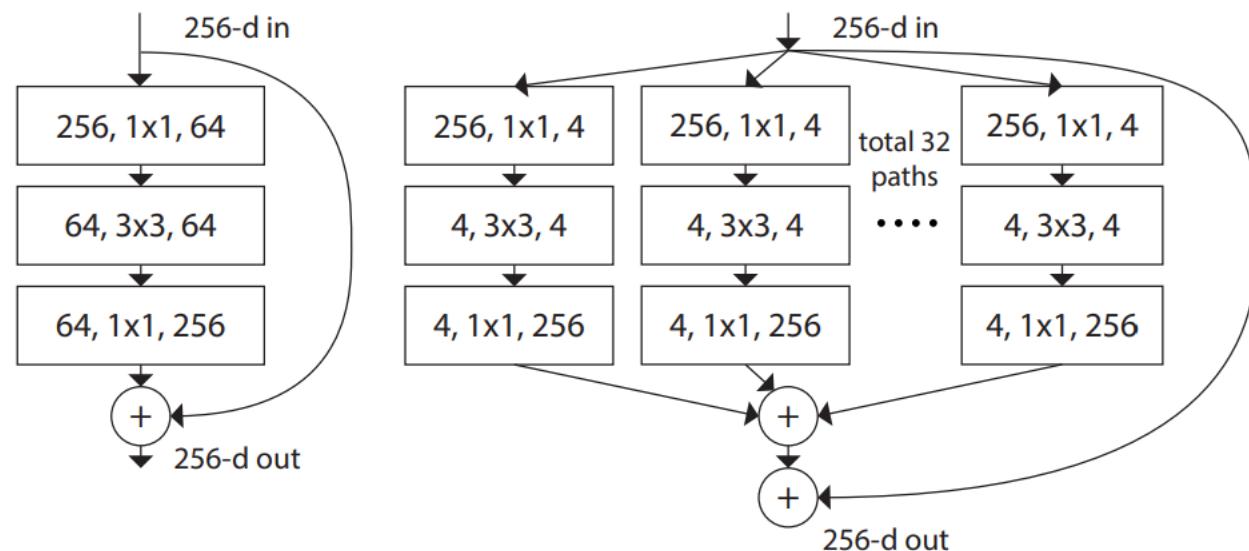


Figure 1. **Left:** A block of ResNet [14]. **Right:** A block of ResNeXt with cardinality = 32, with roughly the same complexity. A layer is shown as (# in channels, filter size, # out channels).

Other Variants...

- MobileNet-V1 (2017)

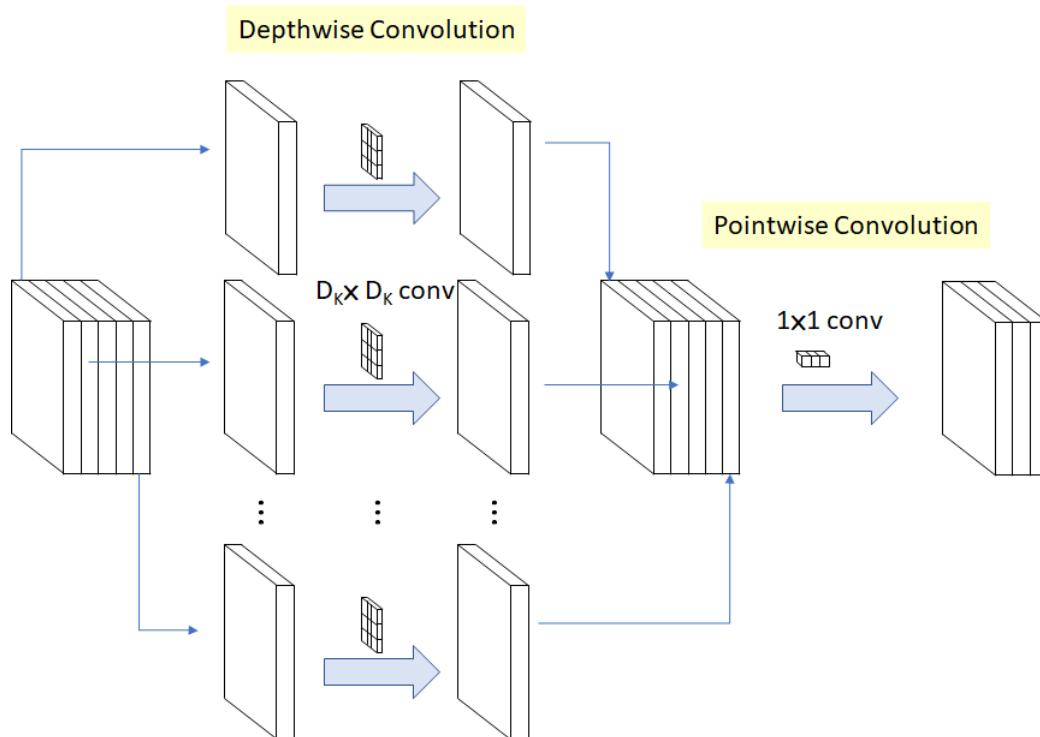
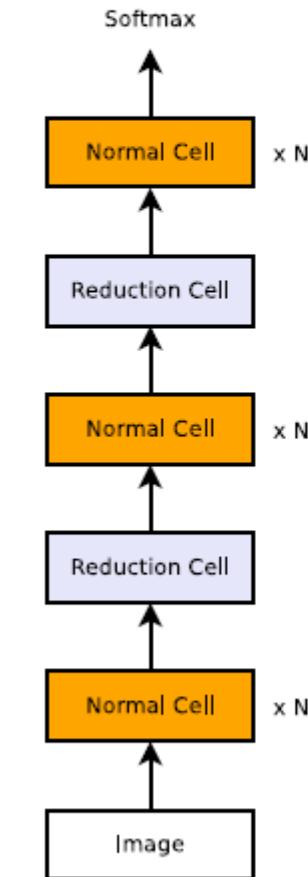


Table 1. MobileNet Body Architecture

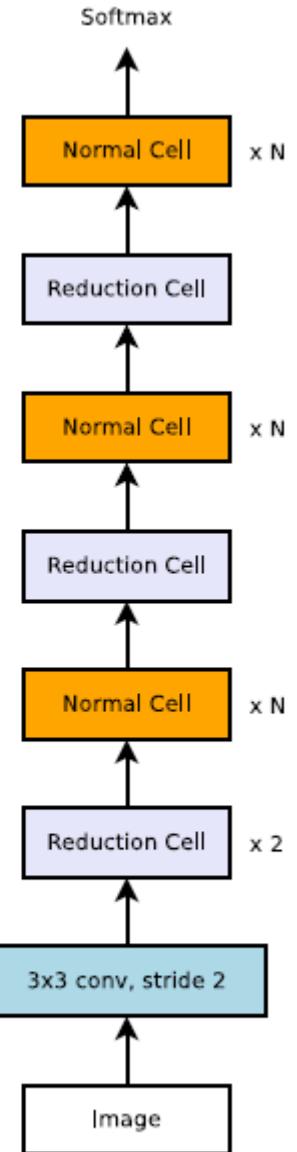
Type / Stride	Filter Shape	Input Size
Conv / s2	$3 \times 3 \times 3 \times 32$	$224 \times 224 \times 3$
Conv dw / s1	$3 \times 3 \times 32$ dw	$112 \times 112 \times 32$
Conv / s1	$1 \times 1 \times 32 \times 64$	$112 \times 112 \times 32$
Conv dw / s2	$3 \times 3 \times 64$ dw	$112 \times 112 \times 64$
Conv / s1	$1 \times 1 \times 64 \times 128$	$56 \times 56 \times 64$
Conv dw / s1	$3 \times 3 \times 128$ dw	$56 \times 56 \times 128$
Conv / s1	$1 \times 1 \times 128 \times 128$	$56 \times 56 \times 128$
Conv dw / s2	$3 \times 3 \times 128$ dw	$56 \times 56 \times 128$
Conv / s1	$1 \times 1 \times 128 \times 256$	$28 \times 28 \times 128$
Conv dw / s1	$3 \times 3 \times 256$ dw	$28 \times 28 \times 256$
Conv / s1	$1 \times 1 \times 256 \times 256$	$28 \times 28 \times 256$
Conv dw / s2	$3 \times 3 \times 256$ dw	$28 \times 28 \times 256$
Conv / s1	$1 \times 1 \times 256 \times 512$	$14 \times 14 \times 256$
5× Conv dw / s1	$3 \times 3 \times 512$ dw	$14 \times 14 \times 512$
Conv / s1	$1 \times 1 \times 512 \times 512$	$14 \times 14 \times 512$
Conv dw / s2	$3 \times 3 \times 512$ dw	$14 \times 14 \times 512$
Conv / s1	$1 \times 1 \times 512 \times 1024$	$7 \times 7 \times 512$
Conv dw / s2	$3 \times 3 \times 1024$ dw	$7 \times 7 \times 1024$
Conv / s1	$1 \times 1 \times 1024 \times 1024$	$7 \times 7 \times 1024$
Avg Pool / s1	Pool 7×7	$7 \times 7 \times 1024$
FC / s1	1024×1000	$1 \times 1 \times 1024$
Softmax / s1	Classifier	$1 \times 1 \times 1000$

Other Variants...

- NASNet (CVPR 2018)
 - Neural Architecture Search
 - Overall architecture is predefined.
Specific blocks and cells are not.
 - Normal Cell: Conv layers that return a feature map of the same size.
 - Reduction Cell: Conv layers that return a feature map of the size reduced by a factor of two
 - Structures of the Normal and Reduction Cells are searched by a recurrent neural network (Lecture 9)



CIFAR10
Architecture

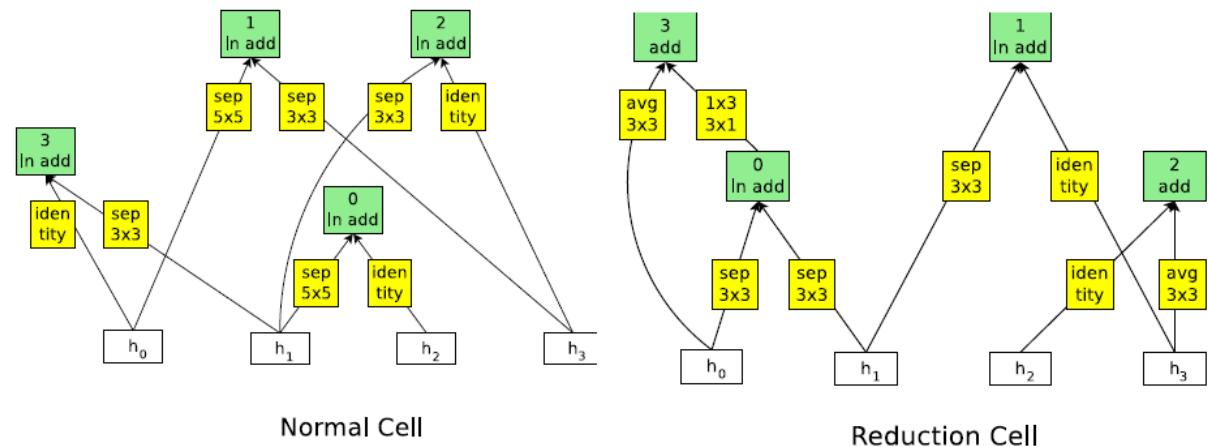


ImageNet
Architecture
104

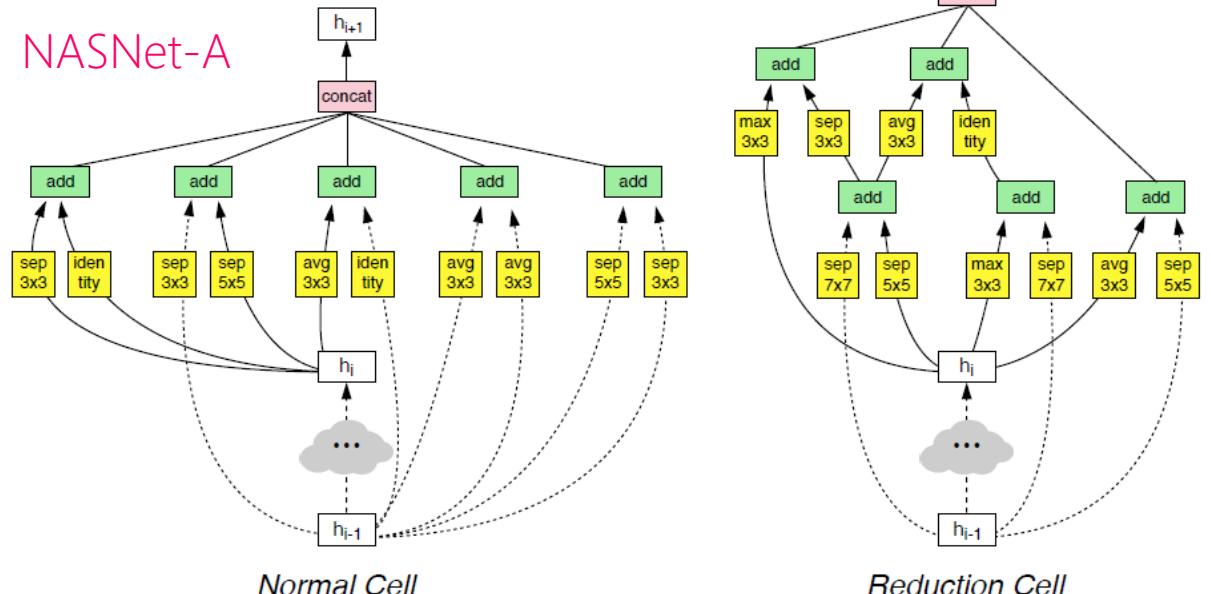
Other Variants...

- NASNet (CVPR 2018)

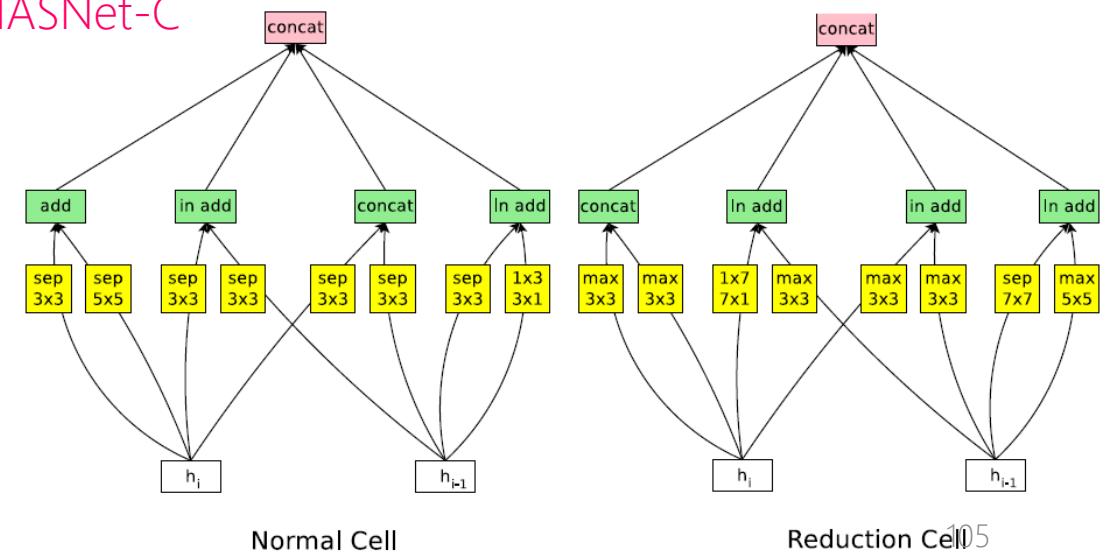
NASNet-B



NASNet-A

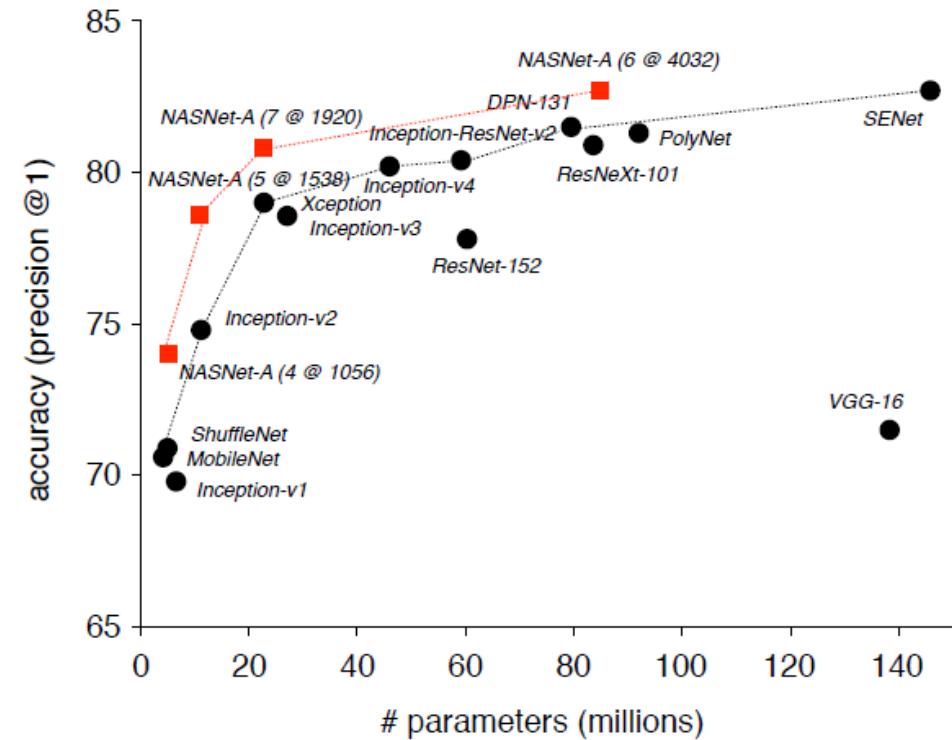
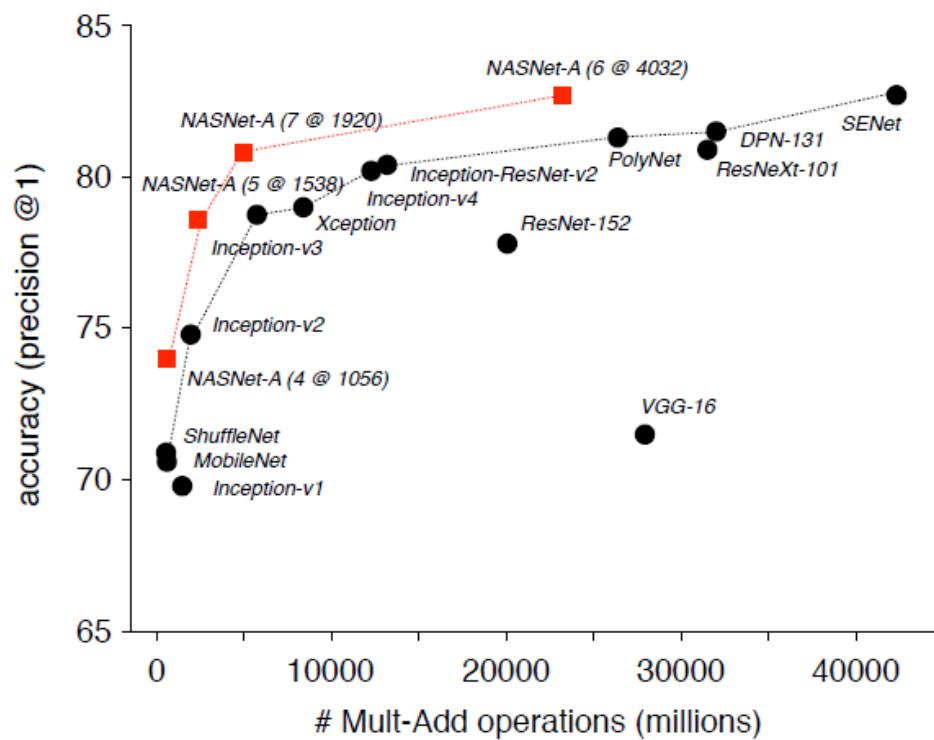


NASNet-C



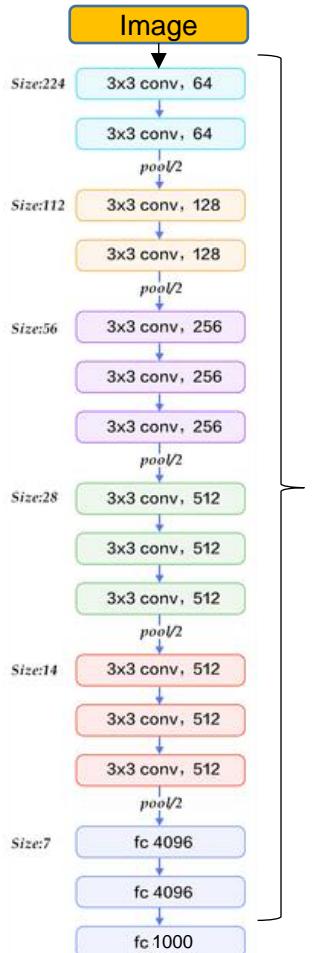
Other Variants...

- NASNet (CVPR 2018)

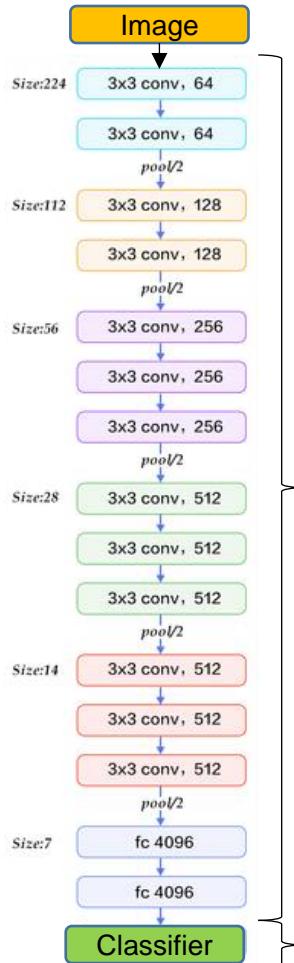


Transfer Learning

Transfer Learning



Transfer
Weights



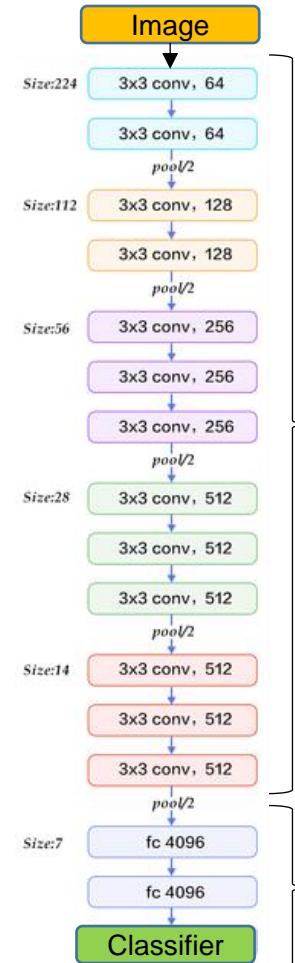
Freeze

...or

Reinitialize and Train

VGG trained on imangenet

Your small dataset



Freeze

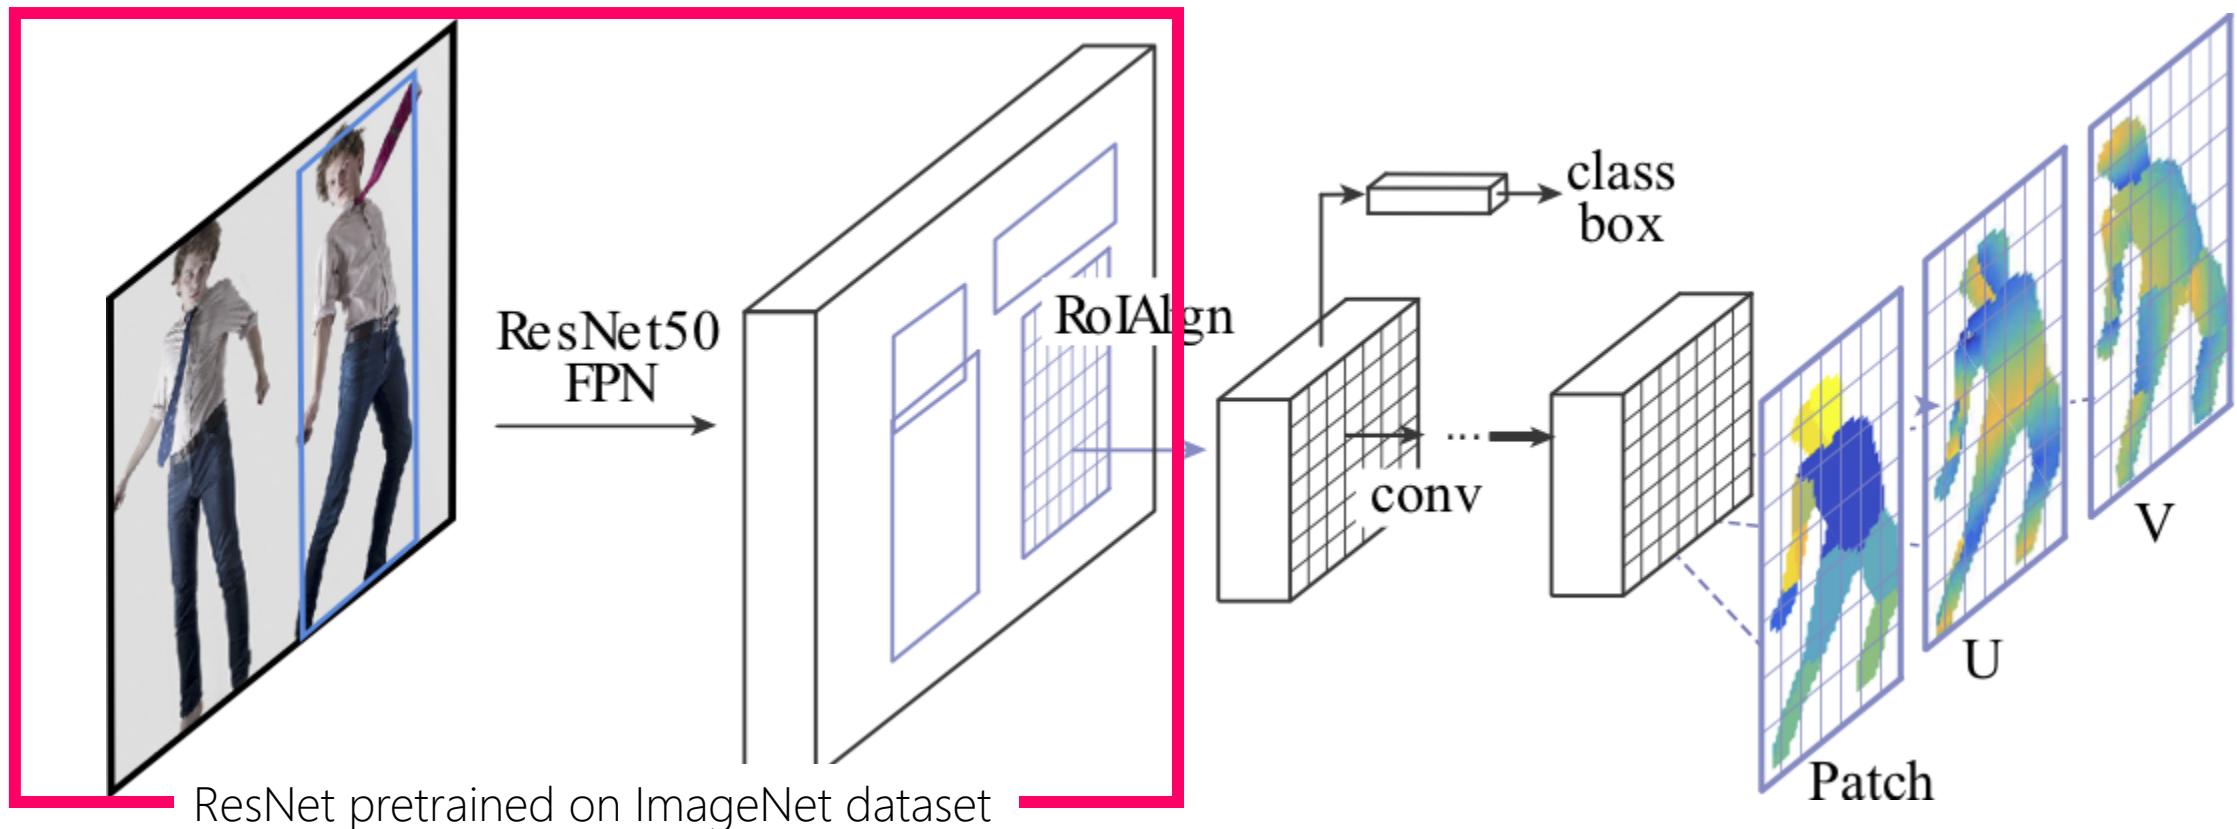
Reinitialize and Train

Transfer Learning

- You may not need to train a large network on your own ☺
 - Save time
 - Save money (to buy an expensive computer)
- You may not need a large dataset ☺

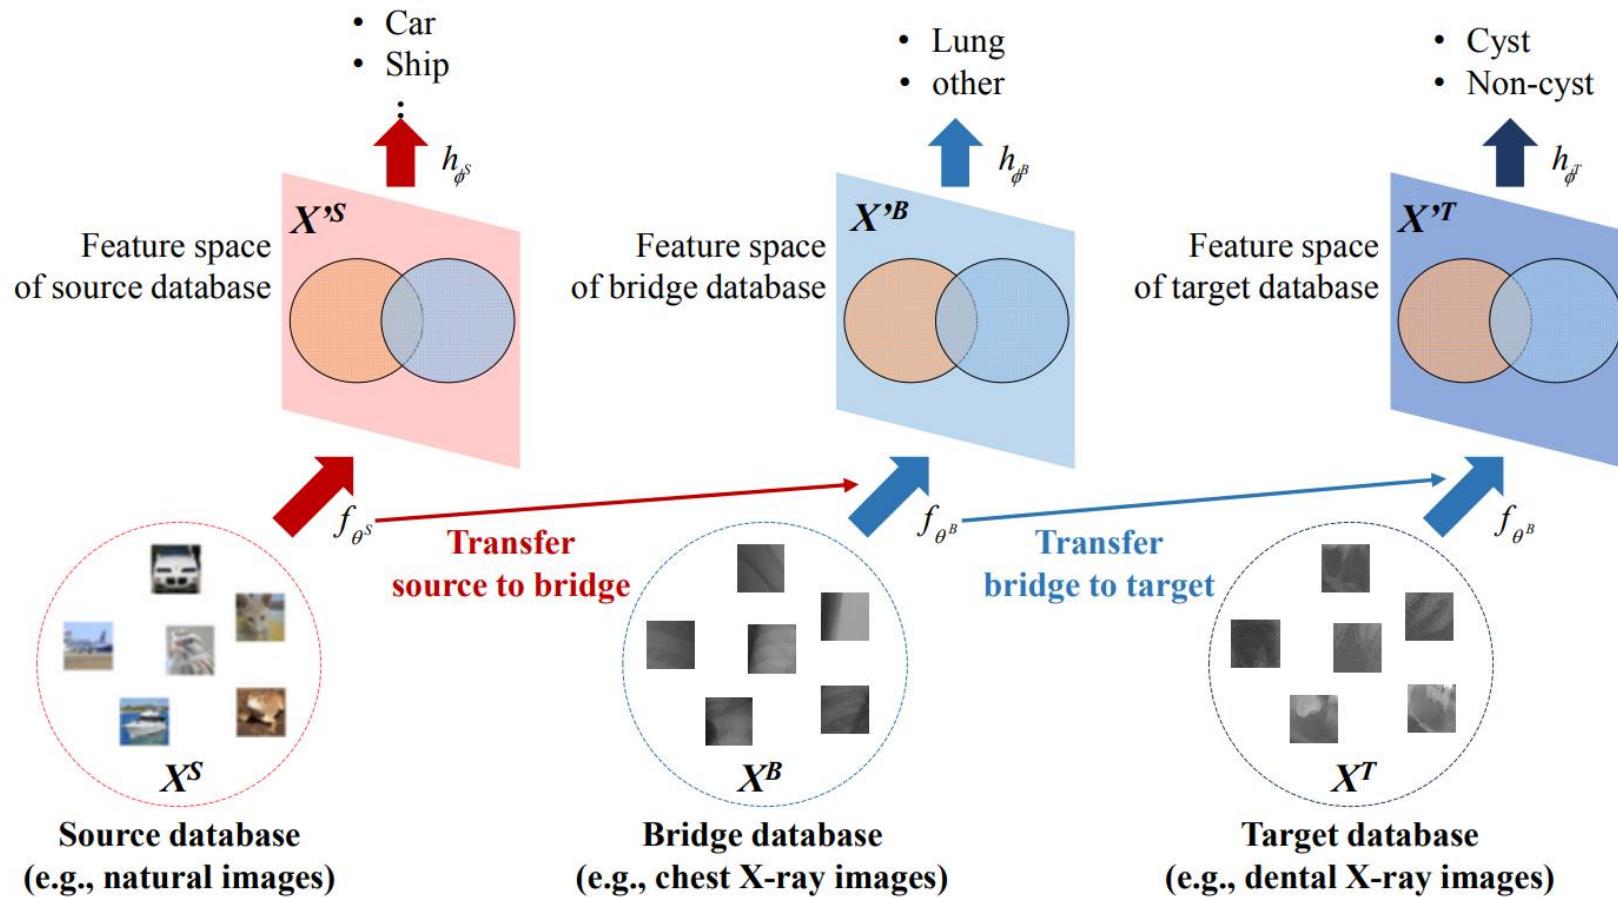
Transfer Learning is VERY Common

- DensePose (Facebook Research, 2018)



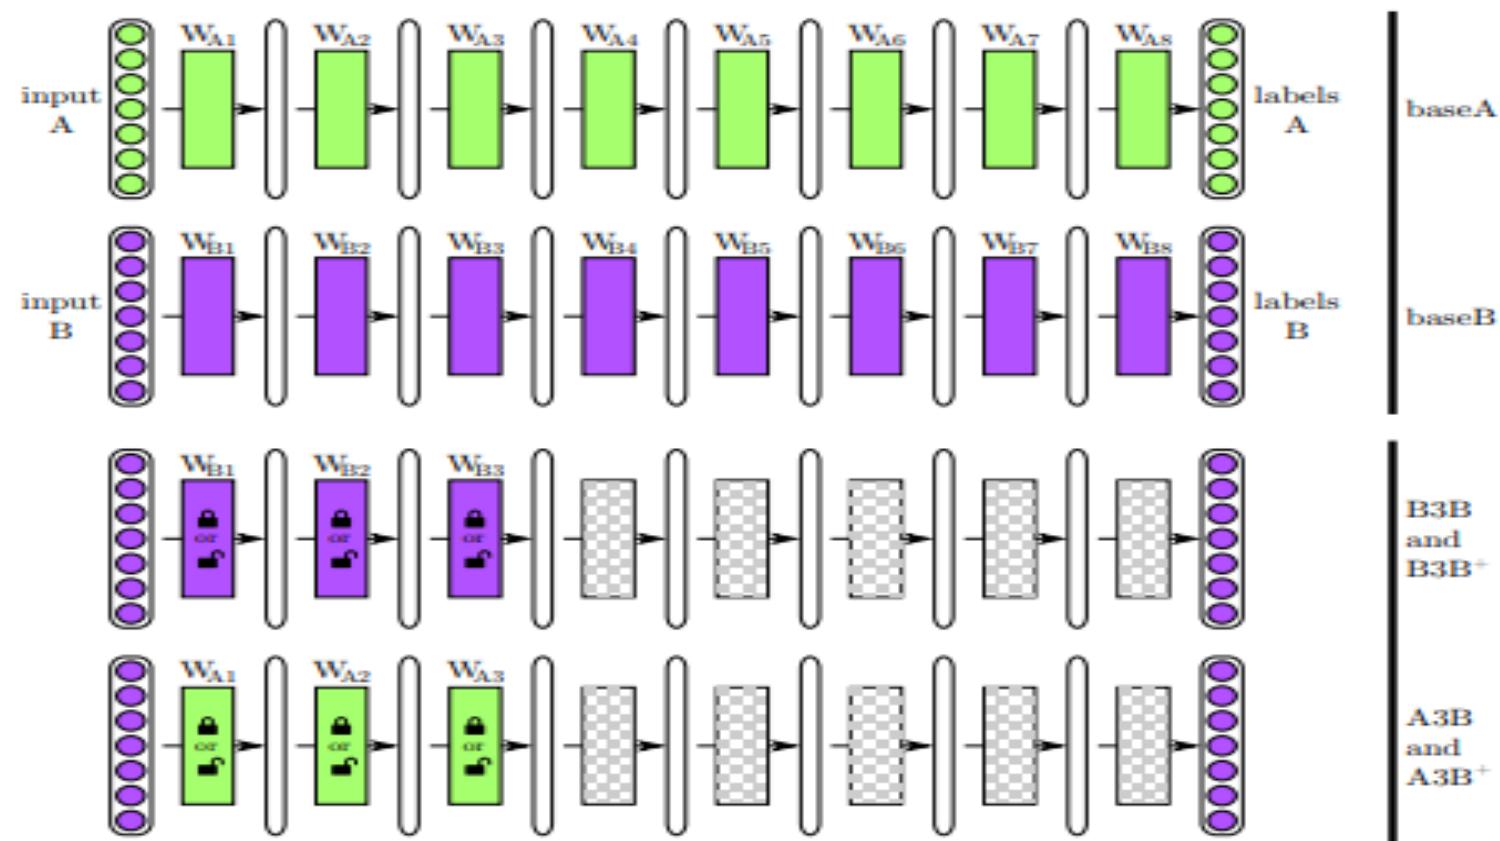
Transfer Learning is VERY Common

- Modality Bridging (Kim, Choi, & Ro 2017)



Transferability of the Layers

- Yosinski et al. (2014)



Transferability of the Layers

- Yosinski et al. (2014)
 - First few layers are more general/universal, later layers are more problem-specific.
 - Transfer + fine-tuning improves generalization

